Unravel: A Deterministic Pre-Analysis Layer for LLM-Assisted Software Debugging

Sambhav Jain

Independent Research

https://github.com/EruditeCoder108/unravelai eruditespartan@gmail.com

March 2026

Every LLM debugging session ends at the crash site. The bug rarely lives there. We present **Unravel**, a debugging engine built on the thesis that LLM-assisted debugging must be reframed as *constraint satisfaction*, not inference. Before any LLM sees the code, a deterministic AST pre-analysis layer (web-tree-sitter, WebAssembly) extracts every variable mutation chain, closure capture, async boundary, floating promise, and cross-file call graph as verified structural facts injected as ground truth into the prompt. The model cannot hallucinate about what does not exist in these facts. It cannot guess. It must trace. These constraints are enforced architecturally: a post-generation Claim Verifier deterministically rejects any output fabricating file names, line numbers, or variable references, independently of what the model was instructed. A two-pass LLM router selects diagnostically relevant files from large repositories. An 8-phase structured reasoning pipeline enforces hypothesis generation followed by evidence-based elimination. Seven anti-sycophancy rules, embedded as hard prompt constraints rather than suggestions, include formally defined guards against *Proximate Fixation* (attributing failure to the nearest code to the crash site rather than the state corruption origin) and the *Name-Behavior Fallacy* (trusting identifier semantics over actual execution paths). A Fix Completeness Verifier uses the already-built call graph to detect uncovered downstream callers in proposed patches at zero added latency. Real-world validation is provided by two merged pull requests in major open-source repositories: tldraw (PR #8161, 8,000+ file TypeScript monorepo) and cal.com (PR #28296), where Unravel produced structured diagnoses used to construct patches accepted by maintainers, with the causal mechanism credited by the issue reporter by name. All experiments were conducted using Gemini 2.5 Flash on the free API tier.

Keywords: automated debugging, multi-file defect localization, static analysis, AST pre-analysis, constraint-based code reasoning, hallucination mitigation, proximate fixation, hypothesis elimination, claim verification, anti-sycophancy, program dependence graph

1. Introduction

LLM-based debugging tools fail systematically on a specific class of bug: defects where the observable symptom and the causal origin are separated by multiple function calls, file boundaries, or async execution steps. This is not a matter of model scale. Claude Opus 4.6, Gemini 3.1 Pro, and GPT-5.4 all exhibit the same failure mode:

they attribute the defect to the most proximate, structurally familiar code rather than tracing state backward through the actual mutation sequence [3, 4].

The root cause of this failure is architectural. The model receives a *description* of what the code should do and a natural language account of the symptom. The structural facts it needs (which variables are mutated at which lines, which closures capture stale state, which callers will be affected by a proposed fix) are present in the code but are not surfaced in a machine-verifiable form. Without these facts as explicit constraints, the model is free to confabulate plausible-sounding evidence for whichever hypothesis is most consistent with its training distribution [2].

On SWE-bench (2,294 real GitHub issues), the best single-model baseline achieved a 1.96% resolution rate [1]. Agentic improvements have raised resolution to 40–55% on selected categories, but the architecture of failure is unchanged: 491 of 497 successfully resolved SWE-bench issues required modifying only a single file [8]. Multi-file bugs, which constitute the majority of non-trivial production defects, remain largely unresolved.

We present **Unravel**, which addresses this problem through a four-layer architecture. The central thesis is that LLM-assisted debugging is most productively understood not as *inference* (pattern-matching from symptom to likely code) but as *constraint satisfaction*: facts extracted deterministically from the AST constrain the hypothesis space, and the LLM's role is to find an explanation consistent with those constraints rather than one consistent merely with its training distribution. Layer 0 implements the constraint extraction; Layer 3's Claim Verifier enforces adherence post-generation. The result is a system that reduces structured entity hallucination to zero on verified fields, enables correct diagnosis on multi-file bugs where unaugmented models fail, and has produced structured diagnoses used to construct pull requests reviewed and merged in two major open-source production repositories.

Contributions:

1. A **four-layer architecture** (deterministic AST pre-analysis, pre-reasoning guards, graph-frontier file routing, 8-phase structured LLM pipeline, and multi-component verification) that reframes debugging as constraint satisfaction over verified AST facts, enabling correct diagnosis on complex multi-file bugs.
2. **Seven Anti-Sycophancy Rules:** formally specified guardrails targeting the two most prevalent reasoning failures: Proximate Fixation (tendency to attribute failure to the nearest code) and the Name-Behavior Fallacy (assuming execution semantics from identifier names).

3. **Symptom Contradiction Module:** a pre-LLM heuristic pass that detects mismatches between the user's symptom description and AST evidence before any model call is made, preventing the LLM from accepting a false premise as its starting constraint.
4. **CFG Branch Annotation:** every mutation in the AST output is annotated as conditional or unconditional (`[CONDITIONAL]` tag). Without this, the LLM treats branch-gated mutations as unconditional evidence, causing correct hypotheses to be incorrectly eliminated.
5. **Hypothesis Elimination Scoring:** the `eliminatedBy` schema field requires every hypothesis elimination in Phase 5 to cite the exact AST-verified file and line that contradicts it. A hypothesis without a code citation is flagged as unverified reasoning, making confident-wrong outputs structurally distinguishable from evidence-backed ones.
6. **Fix Completeness Verifier:** a zero-latency cross-file check using the already-built call graph to identify callers of modified functions omitted from a proposed fix.
7. **Layer Boundary Detector and External Fix Target Verdict:** a formally specified solvability check that identifies when the root cause originates upstream of the provided codebase (`LAYER_BOUNDARY`), and a complementary verdict (`EXTERNAL_FIX_TARGET`) that preserves the full diagnosis when the fix target resides in a different repository.
8. **Real-world validation:** Unravel produced structured diagnoses used to construct pull requests reviewed and merged by maintainers of tldraw (8,000+ file TypeScript monorepo) and cal.com (one of the largest open-source Next.js applications), providing external validation that diagnoses matched the repositories' intended behavior.

We additionally propose **Reasoning Integrity** as a future evaluation metric, distinguishing correctly-reached answers from correctly-reasoned answers using deterministic AST fact verification, as an open problem for the field.

2. Related Work

2.1 LLM Failure Modes in Debugging

Hallucination in code reasoning. Liu et al. [2] demonstrate that LLMs given a buggy prefix code context perform *worse* than zero-context baselines on code completion, fabricating references and behaviors consistent with the buggy context rather

than detecting the defect. Research on hallucination-associated neurons [16] provides a mechanistic explanation: specific neuron clusters activate during generation of factually unsupported content, producing statistically consistent but structurally false claims. Unravel's design addresses this at the prompt construction level: verified AST facts constrain the plausible-sounding hypothesis space before the model's token generation begins.

Proximate fixation. LLMs systematically attribute failures to the crash site rather than the distant corrupted state origin. Arora et al. [3] document a representative case: a model recommending kernel restarts for a bug caused by a single incorrect string literal. The Debugging Decay Index [4] demonstrates exponential effectiveness decay across successive LLM debugging attempts, confirming that the model anchors to its initial (usually proximate) hypothesis rather than revising it.

Cross-file reasoning degradation. Out of 497 successfully resolved SWE-bench issues, 491 required modifying only a single file [8]. Multi-SWE-bench confirms universal performance degradation on cross-file issues across eight languages [7].

Anchoring bias. Log-probability analyses confirm that initial hypothesis tokens occupy disproportionately high-weight positions in the context window, creating a feedback loop that prevents hypothesis revision [5]. RLHF-aligned models compound this with sycophantic responding: a user-proposed incorrect hypothesis receives supporting evidence rather than contradiction.

2.2 Program Graph Representations

Def-use chains and Data Flow Graphs model variable definitions and uses across function boundaries, enabling detection of state corruption paths [11, 10]. Unravel's `extractMutationChains()` is a practical deterministic implementation of def-use chain analysis with line-level attribution.

Program Dependence Graphs (PDGs) unify control and data dependencies, enabling backward program slicing from a failure point to all statements that transitively influenced corrupted state [9]. Unravel's Phase 5 DIAGNOSE step enforces this backward tracing discipline explicitly.

Call Graphs map cross-file function invocations and are essential for identifying cascading fix requirements [6]. Unravel builds an explicit cross-file call graph for both file routing and Fix Completeness Verification.

Temporal Execution Graphs (TEGs) model chronological state transitions essential for async race detection where execution order diverges from source order [12, 13]. Unravel's timing node and floating promise detection approximates TEG boundary identification statically.

2.3 Hybrid Approaches

Hybrid LLM + static analysis approaches treat the LLM as a probabilistic hypothesis generator and program analysis as a partial oracle, eliminating hypotheses contradicted by structural evidence. LDB [14] and TraceCoder [15] ground LLM reasoning in runtime execution traces, complementary to Unravel's static approach. Unravel implements a static approximation: AST facts act as a partial oracle, eliminating hypotheses when contradicted by structural evidence, but not guaranteeing that the surviving hypothesis is correct (dead code paths and runtime-scheduling effects remain outside the static analysis boundary).

3. System Architecture

Unravel processes every debugging request through four sequential layers, each with a narrowly defined responsibility. Figure 1 shows the complete pipeline.

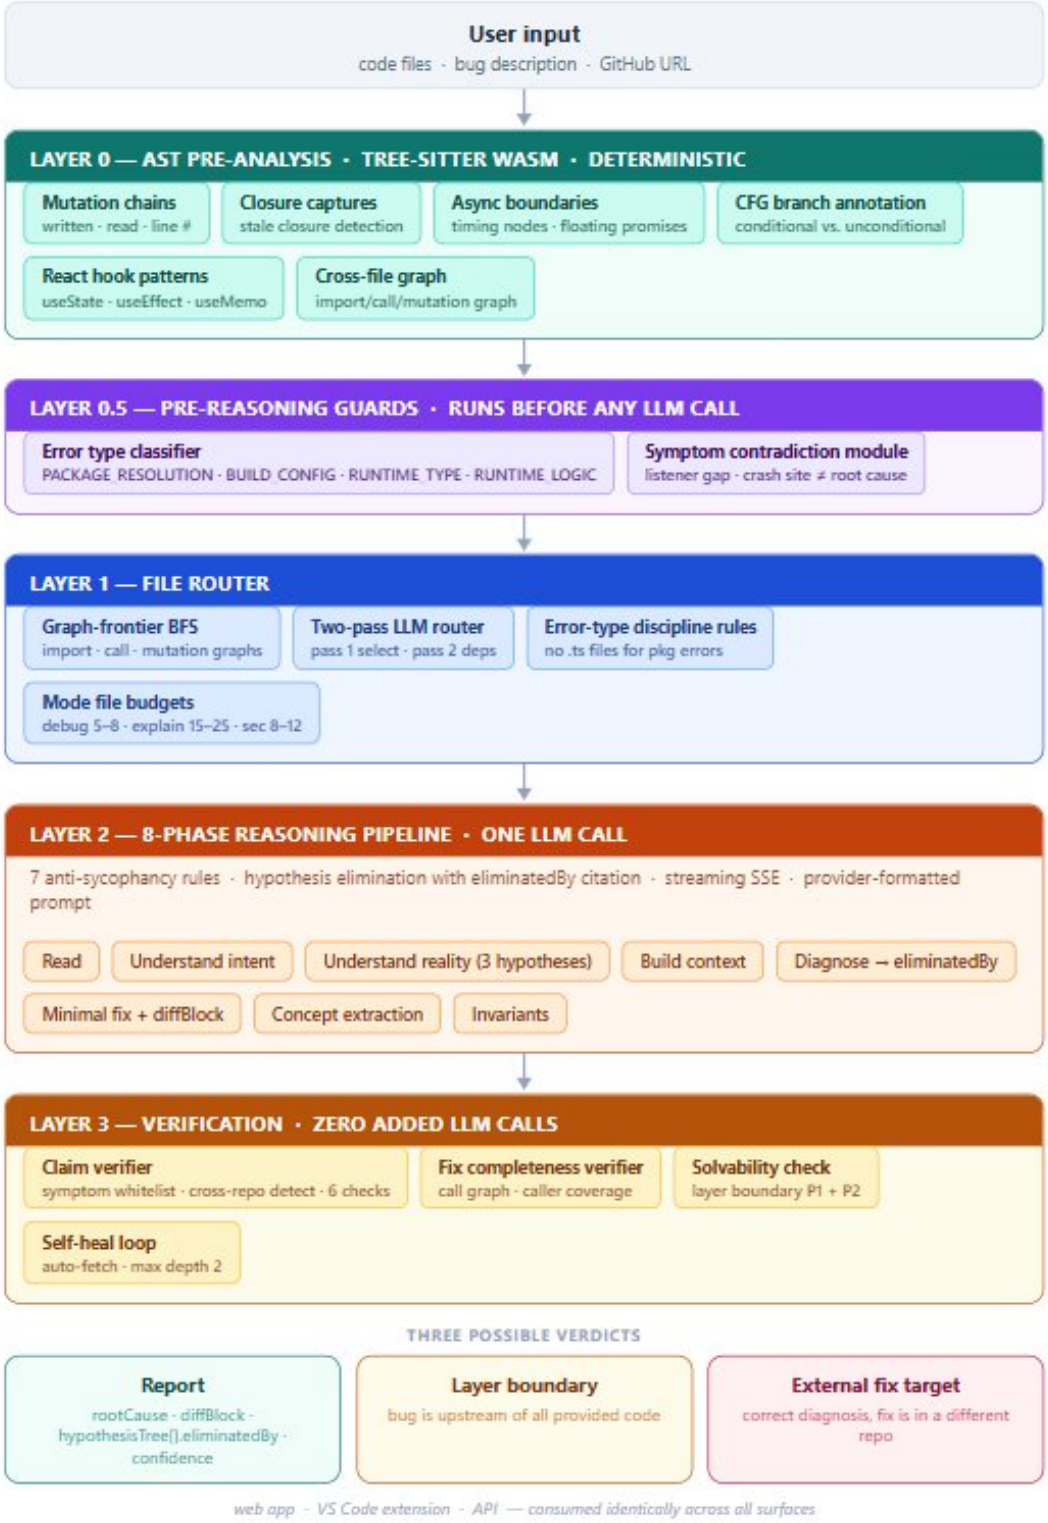

Figure 1: Unravel four-layer architecture. Layer 0.5 (Pre-Reasoning Guards) runs after AST extraction and before the LLM call, injecting symptom contradiction alerts and classifying the error type to govern downstream routing and verification behaviour. Three structured verdicts are possible: a normal diagnosis report, `LAYER_BOUNDARY` (bug is upstream of all provided code), and `EXTERNAL_FIX_TARGET` (correct diagnosis, fix lives in a different repository).

3.1 Layer 0 — AST Pre-Analysis Engine

Layer 0: Deterministic AST Pre-Analysis

The engine (`ast-engine-ts.js`) uses web-tree-sitter in WebAssembly. An earlier version used Babel as the AST engine; this failed on large TypeScript monorepos with complex generic type syntax. The migration to tree-sitter provides partial-parse resilience: files with syntax errors still yield valid structural facts. Explicit `tree.delete()` calls after each operation prevent WASM memory leaks at repository scale.

Scope Resolver. A complete scope map is built as a prerequisite for all downstream analysis. Every binding (variables, parameters, destructured names, hoisted functions) is recorded with its scope level and parent chain.

`extractMutationChains(tree)` walks every assignment expression, augmented assignment, update expression, object property mutation, destructuring assignment, and rest pattern. Each write and read is recorded with the enclosing function name (via upward parent traversal) and line number:

```
Variable Mutation Chains:
  duration
    written: pause()    L69
    written: setMode()  L86
    read:    tick()     L55
    read:    start()    L42
    read:    reset()    L79
```

`trackClosureCaptures(tree)` identifies the three-part profile of a stale closure: referenced inside a function, defined in a parent scope, written anywhere in that parent scope.

`findTimingNodes(tree)` maps every async boundary. Enumerated APIs: `setTimeout`, `setInterval`, `clearInterval`, `requestAnimationFrame` (rAF), `fetch`, `.then()`, `.catch()`, `.finally()`, `addEventListener`, `removeEventListener`. Explicit enumeration prevents reasoning about async code as if execution order matches source order.

`detectFloatingPromises(tree)` uses an `isAwaited(node)` guard. Walking upward from an async call: if the traversal reaches an `await_expression` before an `expression_statement`, the call is properly awaited. If it reaches `expression_statement` first, the promise is floating: the operation runs but its result and any rejection are discarded.

`detectReactPatterns(tree)` identifies three React anti-patterns: `setState` inside a timer without the state variable in the dependency array; `useEffect` with side effects but no cleanup return; missing dependency array entries in `useMemo`/`useCallback`.

Cross-file resolution (`ast-project.js`). After per-file analysis, a cross-file pass

builds the module map, resolves symbol origins across import/export boundaries, expands mutation chains across file boundaries, and emits risk signals for cross-file mutation, async state races, and stale closures.

Symptom Contradiction Module. Before any LLM call, a heuristic pass compares the user's symptom description against AST evidence. Two checks run in debug mode:

- **Listener Gap:** If the symptom contains phrases like "not firing" or "event not triggered" but AST confirms `addEventListener` is wired for the relevant event type, an alert is injected: the listener exists, the bug is in the handler logic, not the binding.
- **Crash Site $\neq$ Root Cause:** If the symptom names a specific function and AST shows that function makes *no writes* (reads state only), it is flagged as the crash site rather than the root cause. The model is instructed to trace backward from it, not diagnose it directly.

Alerts are injected as `SYMPTOM CONTRADICTION ALERTS` into the ground truth block. They never block analysis; they bias the model toward correct framing before token generation begins.

The combined AST output is injected as a structured ground-truth block framed explicitly as instructions the model is not permitted to contradict:

```
VERIFIED GROUND TRUTH -- confirmed by static analysis.
Instructions: do not contradict any fact below.
============================================================
Variable Mutation Chains:
  duration
    written: pause() L69    <- potential root cause
    read:    tick() L55, start() L42, reset() L79

Async Boundaries:
  setInterval -> tick() [L57]
  addEventListener("visibilitychange") -> handler() [L110]

Closure Captures:
  tick() captures: duration, remaining, interval
  STALE CANDIDATE: tick() reads duration from outer scope
```

3.2 Layer 0.5 — Pre-Reasoning Guards

Layer 0.5: Pre-Reasoning Guards

Two deterministic checks execute after AST extraction and before any LLM call. Their purpose is to prevent the model from accepting a false premise as its starting constraint. Neither check blocks analysis; both inject corrective alerts into the ground truth block so the model is challenged *before* its first token is generated.

Error Type Classifier. A zero-LLM-cost regex classifier assigns every analysis request to one of four categories: `PACKAGE_RESOLUTION`, `BUILD_CONFIG`, `RUNTIME_TYPE`, or `RUNTIME_LOGIC`. This classification propagates to three downstream systems: (1) Layer 1's Pass 2 discipline rules (e.g. `PACKAGE_RESOLUTION` errors are explicitly prohibited from requesting `.ts`/`.js` source files, since package resolution bugs are never fixed in application source code); (2) the solvability check in Layer 3, which short-circuits the `LAYER_BOUNDARY` classifier for configuration-fixable error types; and (3) the claim verifier, whose cross-repo detection behaviour is conditioned on whether the error is structural or runtime.

Symptom Contradiction Module. Described in §3.1. Two checks run in debug mode: the Listener Gap check (user says "not firing" but AST confirms `addEventListener` is wired) and the Crash Site $\neq$ Root Cause check (user names a function that makes no state writes). Alerts are injected as `SYMPTOM CONTRADICTION ALERTS` into the ground truth block. The model is explicitly instructed to investigate the contradiction before accepting the user's framing at face value.

3.3 Layer 1 — File Router

Layer 1: File Router

Unravel supports three input modes. For **paste/upload** mode, a graph-frontier BFS router traverses import and call graph edges from symptom-adjacent files. For **GitHub URL** mode, a dedicated two-pass fetch pipeline operates as follows.

3.3.1 GitHub File-Fetch Pipeline

When given a GitHub repository or issue URL, the following pipeline executes before any AST analysis:

Figure 2: GitHub file-fetch pipeline (Layer 1). Step 3 classifies the error type deterministically at zero LLM cost; this governs Pass 2 discipline rules (e.g. `PACKAGE_RESOLUTION` errors are never routed to `.ts`/`.js` source files). The self-heal gate distinguishes cross-repository references (emits `EXTERNAL_FIX_TARGET`) from genuinely missing files (auto-fetches, max depth 2).

Mode-specific file budgets (upload/BFS mode):

Mode	File Range	Selection Logic
Debug	5–8 files	Max call-graph proximity to symptom
Explain	15–25 files	Breadth-first coverage
Security	8–12 files	Attack-surface paths

Per-file truncation injects explicit `[TRUNCATED by Unravel]` markers. Every result records `_provenance.routerStrategy` (`"graph-frontier"`, `"llm-heuristic"`, or `"all-files"`), `engineVersion`, and `astVersion` for auditability.

3.4 Layer 2 — The 8-Phase Reasoning Pipeline

Layer 2: Core Engine

One LLM call is made per analysis, but the system prompt enforces traversal through eight phases in strict order. Phases 1–4 are shared across all three modes (Debug, Explain, Security). Phases 5–8 are mode-specific. Phase 5 (DIAGNOSE) requires every hypothesis elimination to populate `eliminatedBy` with the exact AST-verified file and line that contradicts it; a hypothesis without a citation is flagged as unverified reasoning. Phase 6 (MINIMAL FIX) produces a `diffBlock` field containing a unified diff (`-- file:line / - old / + new`) in addition to prose explanation.

Phase 6 defaults to the smallest possible surgical change. An *architectural note* is added — never instead of the surgical fix, always in addition to it — in any of three cases: (A) the patch suppresses the visible error but does not eliminate the underlying incorrect state (a null-check around a crash site when null should never reach there); (B) the patch technically works but violates framework best practices or creates maintenance debt (adding deps to a `useCallback` when the hook should be inlined into `useEffect`; adding a lock around shared state when the real fix is to eliminate shared ownership); or (C) the bug exists because of wrong data ownership, shared mutable state across async boundaries, or a missing abstraction layer that will cause the same class of bug to recur. `RACE_CONDITION` bugs involving shared async state explicitly qualify for Case C. The surgical patch remains actionable in all three cases; the architectural note documents the structural property that was violated.

The prompt is formatted differently per provider: Claude receives XML tags (`<instructions>`, `<rules>`, `<phase>`) for higher instruction-following fidelity; Gemini receives Markdown headers; OpenAI-compatible models receive section delimiters. The reasoning protocol is identical across providers.

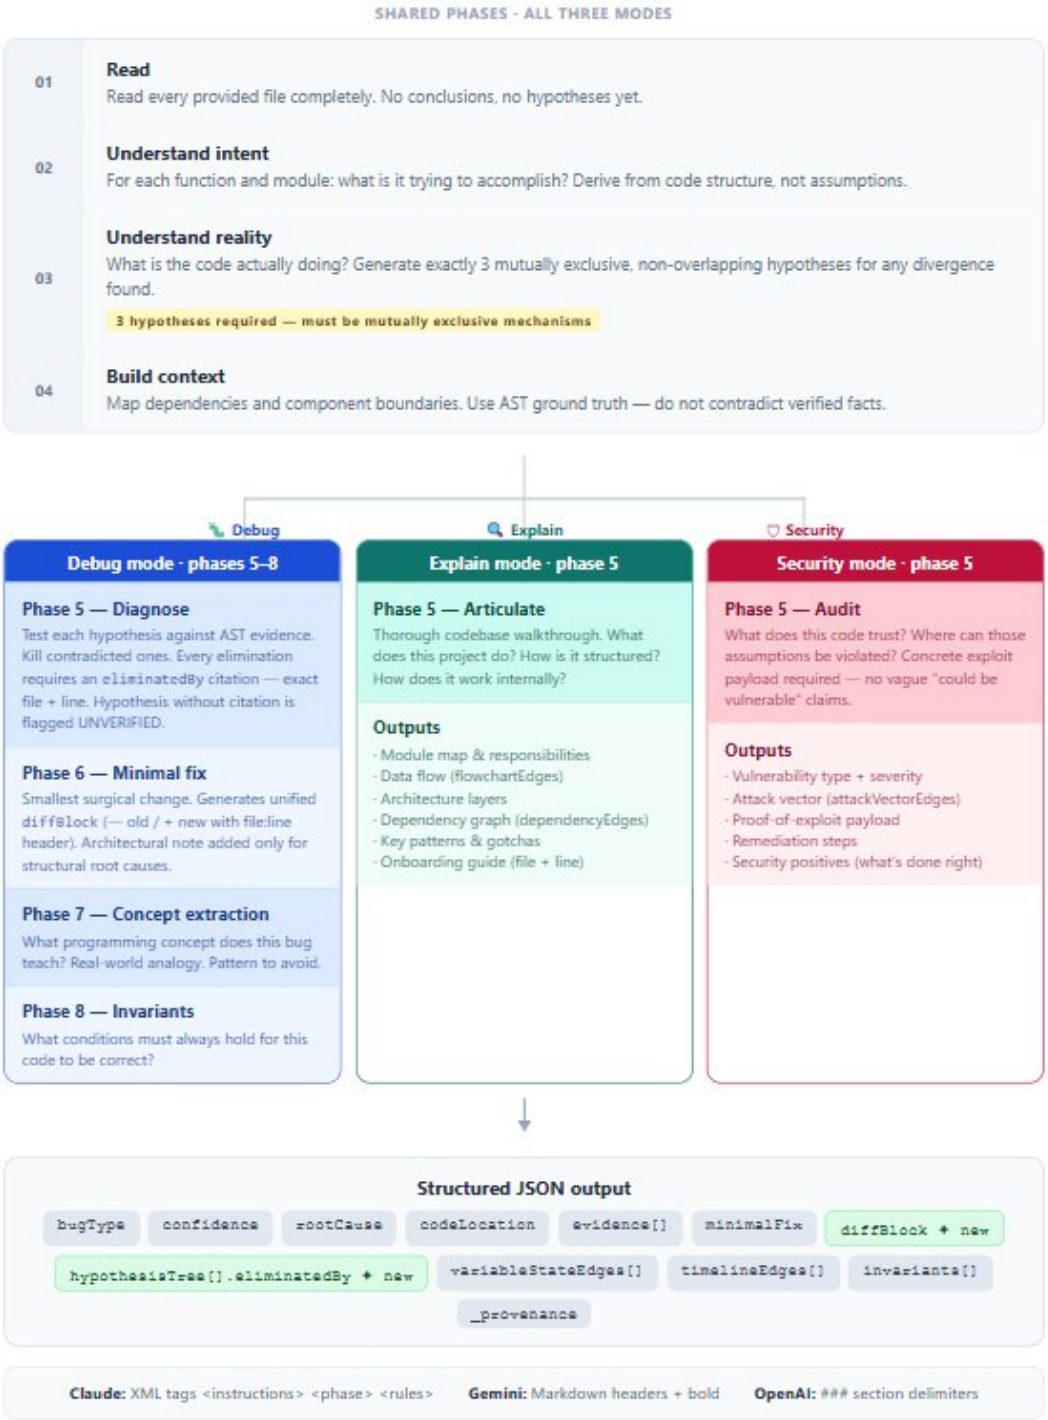

Figure 3: 8-phase reasoning pipeline. Phases 1–4 are common to all three modes. Phase 5 branches by mode; Debug mode continues through Phase 8. The `eliminatedBy` field (new) requires each hypothesis elimination in Phase 5 to cite the exact AST-verified file and line that contradicts it; a hypothesis without a citation is flagged as unverified reasoning. The `diffBlock` field (new) in Phase 6 produces a unified diff (`-- file:line / - old / + new`) for every fix.

Output schema. Every analysis produces typed JSON consumed identically by the web app, VS Code extension, and future clients:

```
{
  "bugType":      "STATE_MUTATION",
  "confidence":   0.92,
  "rootCause":    "Variable ‘duration‘ overwritten in pause() at
     line 69. All
                   subsequent reads in tick(), start(), reset()
                      receive the
                   corrupted value.",
  "codeLocation": { "file": "script.js", "line": 69 },
  "evidence": [
    "AST confirms: duration written at pause() L69",
    "AST confirms: duration read at tick() L55 -- receives
       corrupted value"
  ],
  "diffBlock": "--- script.js L69\n-    duration = remaining\n+
     lastActiveRemaining = remaining",
  "hypotheses": [
    { "id": "H1", "text": "Stale closure in tick() captures
       duration
                            at initialization",
      "status": "eliminated",
      "eliminatedBy": "script.js L55: tick() resolves duration
         from outer
                        scope at call time, not capture time",
      "reason": "Scope resolver confirms tick() resolves duration
         from
                 outer scope at call time, not capture time" },
    { "id": "H2", "text": "duration mutated in pause() before
       reset()",
      "status": "survived",
      "eliminatedBy": null,
      "reason": "Mutation chain confirms write at L69; reset()
         reads at L79" }
  ],
  "minimalFix": "Remove ‘duration = remaining‘ from pause().
     duration must
                 remain constant for the session.",
  "invariants": ["duration must not change during active or paused
     session"],
  "_provenance": {
    "engineVersion": "3.3",
    "astVersion": "2.2",
    "routerStrategy": "graph-frontier",
    "model": "gemini-2.5-flash"
  }
}
```

3.5 Layer 3 — Verification and Integrity

Layer 3: Verification & Integrity

Four components execute synchronously after each LLM response, before any output is surfaced. The Claim Verifier is the enforcement mechanism for the verified structural facts injected in Layer 0: the LLM is not made infallible by prompt instruction, but any output that contradicts injected facts is rejected deterministically here. Three structured verdicts are possible: a normal diagnosis **Report**, a `LAYER_BOUNDARY` verdict when the root cause is upstream of all provided code, and an `EXTERNAL_FIX_TARGET` verdict when the diagnosis is correct but the fix must be applied in a different repository.

Claim Verifier. Six deterministic cross-checks against source files: (1) root cause file exists in provided inputs (hard rejection on failure); (2) cited line number is within file bounds; (3) evidence array filename references exist; (4) `codeLocation` file-line pairs validated via paired extraction; (5) security vulnerability file exists; (6) variable names in `evidence` present in the AST mutation map with fuzzy matching for property access patterns. Violations above threshold set `needsMoreInfo: true` and trigger the self-healing loop.

Fix Completeness Verifier. Uses the call graph from Layer 0 at zero additional LLM calls and zero latency. For each function in `minimalFix`, all callers are identified. If a caller file is absent from the fix, a warning is injected into the `uncertainties` field and confidence is penalized by 0.15. React component files (PascalCase filenames, `.jsx`/`.tsx` extensions) are excluded: fixing a custom hook without modifying its consumers is encapsulation working as intended. The check fires only on signature-breaking changes (removed parameters visible in `diffBlock`); backward-compatible changes such as adding a guard or early return never trigger it, since callers require no updates when the function's interface is unchanged.

Claim Verifier precision improvements. Five false-positive patterns were identified and resolved across UDB-20 grading sessions. (1) *Class property matching (B-05):* TypeScript class properties are stored in the AST as `this.propName`; the verifier now registers the bare name (`propName`) in addition to the prefixed form so that AI-reported `currentConfig` correctly matches `this.currentConfig`. (2) *Array subscript stripping (B-05):* variables reported as `serverLikeCounts[]` are now normalized by stripping the bracket suffix before lookup. (3) *Fix Completeness scope (B-05):* the completeness check previously fired on any internal change; it now fires exclusively on diffs containing removed lines with function parameter signatures — backward-compatible additions (guards, early returns) do not trigger it. (4) *Claim-side `this.` stripping (B-06):* the verifier previously stripped `this.` only when building `knownVars`; it now also strips it from the AI's claimed variable name at match time, so `this.isReady` in

the output correctly matches `isReady` in the AST. (5) *Empty-mutations skip (B-07):* when `astRaw.mutations` is empty (no top-level variable mutations in the scanned file set), Check 4 is skipped entirely — previously, an empty `knownVars` caused every claim to fail, generating false warnings on stable service registries.

AST engine additions (B-10). Three detection capabilities were added to `ast-engine-ts.js` during B-10 grading and are now active on all future runs.

(1) `react_unstable_dep`: Detects when a `useEffect` dependency is created unstably inside a component body — `new Service()`, `[]`, `{}`, or an arrow function — without being wrapped in `useMemo` or `useCallback`. This class of bug causes the effect to re-run on every render as the object reference changes, accumulating subscriptions or side effects over time.

(2) `detectListenerParity()`: Compares each `addEventListener` / `removeEventListener` pair and emits a structured annotation. Mismatches in the `capture` flag (which the DOM specification uses as part of listener identity) are flagged as breaking. Mismatches in `passive` or `once` are flagged as informational style warnings only, with the annotation explicitly stating: *"passive/once omission is harmless (spec: only capture affects identity)."* This annotation directly prevented the pre-fix misdiagnosis on B-10 (§6.3.1).

(3) `isWrappedInMemoHook()`: Before flagging an object or function as unstably created, the engine checks whether it is wrapped in `useMemo` or `useCallback`. If it is, no instability flag is emitted. This prevents false positives in codebases that correctly stabilize dependencies with memoization.

3.6 Layer Boundary Detector

Definition: Layer Boundary Bug

A *layer boundary bug* is a defect whose causal origin resides in a software or system layer outside the scope of the provided codebase. Attempting to fix such bugs from within the application layer produces patches that are incorrect by construction: they address a symptom of an upstream condition rather than the condition itself.

This insight arose from a concrete failure: when analyzing a VS Code keyboard issue where `cmd+shift+]` failed on a Mac with the EurKEY layout, Unravel correctly traced the root cause to the OS translating the physical `]` key into `Digit9` before any VS Code code runs. Two completely different physical actions arrive looking identical at the application boundary. The fix does not exist in the VS Code codebase, because the distinguishing information was lost one layer upstream.

Formal definition. A diagnosis is classified as `LAYER_BOUNDARY` when both predicates hold:

```
P1 (deterministic):
  forall file_ref in extract_file_refs(rootCause):
    file_ref NOT IN provided_codebase
    AND NOT (verification.rootCauseRejected)   // hallucination !=
       boundary

P2 (heuristic):
  exists keyword in UPSTREAM_KEYWORDS:
    keyword IN (rootCause OR evidence OR symptom)
```

P1 is the hard gate: if the root cause mentions any provided file, the check does not fire regardless of keyword matches. P2 prevents false positives from bugs that simply omit file names from the root cause text. The keyword set covers OS event routing, browser event internals, Electron native bindings, external API layers, and keyboard layout translation.

Confidence scoring. The classifier assigns:

$$\text{confidence} = \min\big(0.95,\ 0.70 + 0.05 \times |\text{matched_keywords}| + 0.10 \times \mathbf{1}[\text{rootCause cites no files}]\big)$$

Case study: tldraw CLI (PR #8161). The primary CLI bug (wrong install directory) was diagnosed as a `DATA_FLOW` defect in the provided code. A secondary report (files created after user cancellation) also triggered this detector: `ensureDirectoryEmpty` calls `process.exit(1)` on cancellation, a behavior originating in npm's scaffolding infrastructure upstream of the CLI code. Unravel correctly separated these two diagnoses rather than generating a single incorrect fix.

Self-Healing Context Loop. Triggered by two signals detected after claim verification: (A) the model's `minimalFix` text contains phrases indicating it knows the fix target is an unseen file (e.g. "implementation is not provided"); or (B) `codeLocation` references a filename not present in the input set while the root cause cites a file that is present. In GitHub URL mode, the missing file is automatically fetched from the stored repository tree using exact path matching with fuzzy filename fallback. Depth is limited to two iterations: the first reliably recovers direct missing dependencies; the second resolves indirect imports discovered after the first fetch expands the visible interface boundary. Beyond two iterations, additional file fetches expand context superlinearly while rarely improving diagnostic resolution in practice.

External Fix Target Verdict. A third structured verdict (`EXTERNAL_FIX_TARGET`) is emitted when the claim verifier detects that the model's fix target belongs to an external dependency package rather than the scanned codebase. This is distinguished from hallucination by checking whether the cited file's package prefix matches a name present in the project's dependency declarations. When triggered, the self-heal loop

is skipped entirely (the file is in a different repository and cannot be fetched from the local tree). The full diagnosis is preserved alongside `targetRepository`, `targetFile`, and `suggestedAction` fields, giving the developer the complete root cause and a precise pointer to where the fix must be applied.

3.7 Output Presets and Diagnostic Invariants

Unravel exposes three output presets (Quick Fix, Developer, Full Report) and a user experience level selector (Beginner, Vibe Coder, Basic, Developer). A key design constraint governs both: *presets and levels control presentation, not diagnostic depth.*

Preset-invariant diagnostic core. The 8-phase reasoning pipeline executes identically across all presets. Critically, the output fields that enforce hypothesis elimination — `hypothesisTree`, `eliminatedBy`, `evidence[]`, `proximate_crash_site`, and `diffBlock` — are bundled into the `rootCause` schema group and cannot be removed by any preset. This is an architectural invariant: a Quick Fix run and a Full Report run receive the same system prompt, the same 8-phase reasoning instructions, and the same diagnostic output requirements. The only difference is that Full Report additionally requests presentational sections (execution timeline, concept extraction, invariants, variable state visualization). An engine that produces a weaker root cause diagnosis under Quick Fix mode than Full Report mode is exhibiting a prompt construction defect, not expected behavior. The UDB-20 mode parity policy (§6.1.1) explicitly grades for this.

Level-scope constraint. The user experience level (Beginner through Developer) governs the *explanation style* of text fields — the warmth of the analogy, the jargon density of `whyFixWorks`, the complexity of `conceptExtraction`. It does not govern the diagnostic fields. The level instructions are explicitly prefixed with a diagnostic precision mandate: evidence citations, line numbers, mutation chains, and hypothesis eliminations are produced at full technical accuracy regardless of audience. The model is instructed to *do expert-level tracing first, then translate the explanation* — not to simplify the tracing itself. A beginner user receives the same root cause, the same evidence, and the same fix as an intermediate developer; they receive it explained differently.

Language-scope constraint. The output language selector (English, Hinglish, Hindi) governs natural language production only. The English instruction was revised from a simplification mandate ("explain like a 15-year-old") to a professional neutral baseline that inherits depth from the level instruction. Indian cultural analogies are gated to Hinglish and Hindi settings and do not appear in English output.

4. Anti-Sycophancy Rules

Every LLM debugging tool defaults to agreement. If the symptom description names a component, the model focuses on it. If a user proposes an incorrect hypothesis, supporting evidence is generated rather than a contradiction [5]. Unravel embeds seven rules hardcoded into every prompt: structural constraints on what the model is permitted to output, not suggestions.

Formal Failure Mode Definitions

Proximate Fixation: The tendency of an LLM to attribute failure to the nearest structurally familiar code to the crash or error site, rather than tracing state backward to the distant origin of corruption. The model stops at the first plausible explanation rather than eliminating alternatives.

Name-Behavior Fallacy: The tendency to assume that the semantic behavior of a function or variable matches its identifier name, treating a function named `cleanup()` as if it performs cleanup, or `isAuthenticated` as if it guarantees authentication, without verifying the actual execution path.

#	Constraint	Failure Mode Targeted
1	If the code is correct, say "No bug found." Do not invent defects.	Sycophantic RLHF alignment
2	If the user's description contradicts the code, state the contradiction explicitly.	Confirmation bias
3	If uncertain, say "Cannot confirm without runtime execution."	Fabricated confidence
4	Every bug claim requires an exact line number and code fragment.	Evidence fabrication
5	Never describe code behavior not verifiable from the provided files.	Hallucinated behavior
6	The crash site is the symptom. Trace state backward through mutation chains to where it was first corrupted.	Proximate Fixation
7	A variable named `isPaused` does not guarantee the code pauses. Verify execution chains, not naming conventions.	Name-Behavior Fallacy

Rules 6 and 7 address the two most pervasive failure patterns identified in our analysis of LLM debugging behavior. Rule 6 directly targets proximate fixation, the most cited failure mode in the literature. Rule 7 targets a subtler failure: LLMs assign semantic certainty to identifiers (`cleanup()`, `isAuthenticated`, `handleSave`) without verifying the actual execution path, missing bugs that look correct on a surface reading precisely because they have semantically correct names.

5. Bug Taxonomy

Every diagnosis is classified against a 12-category formal taxonomy that is part of the output schema. Classification enables cross-run pattern analysis and drives the concept extraction phase.

Category	Description
`STATE_MUTATION`	Variable intended to be constant is modified unexpectedly
`STALE_CLOSURE`	Function captures outdated variable value
`RACE_CONDITION`	Multiple async operations conflict on shared state
`TEMPORAL_LOGIC`	Timing assumptions break: drift, wrong timestamps, sequence violations
`EVENT_LIFECYCLE`	Missing cleanup, double-binds, wrong event registration order
`TYPE_COERCION`	Implicit conversion causes unexpected computation
`ENV_DEPENDENCY`	Code behaves differently across environments
`ASYNC_ORDERING`	Operations execute in wrong sequence
`DATA_FLOW`	Data passes incorrectly between components or files
`UI_LOGIC`	Visual behavior does not match intended state
`MEMORY_LEAK`	Resources not released, accumulate over time
`INFINITE_LOOP`	Recursive or cyclic behavior creates runaway execution

6. Evaluation

The primary evidence for Unravel's diagnostic correctness is real-world validation: diagnoses that were implemented as merged pull requests in major open-source repositories, where maintainer acceptance provides an external correctness signal. We additionally provide: (a) a comparative evaluation against unaugmented frontier models on two hand-constructed test cases, and (b) UDB-20, a reproducible 20-bug benchmark evaluated on three axes designed to directly measure the failure modes Unravel targets. Results for ten completed bugs (B-01–B-10) show Unravel scoring 59/60 (98.3%) vs. baseline 52/60 (86.7%), a delta of +7 points across all three axes — with the largest gap on Causal Flow Reconstruction (CFR), the multi-hop causal chain explanation axis (20/20 vs. 14/20).

UDB-11 (the original 11-bug internal benchmark) is retained as historical context for the Claim Verifier design but is not presented as a primary benchmark result. Its limitations — self-referential construction, single-axis scoring, LLM-only grading — are superseded by UDB-20's design. The UDB-11 data appears in §6.5 for completeness.

6.1 Evaluation Methodology

6.1.1 UDB-20 Three-Axis Rubric

UDB-20 scores each diagnosis on three independent axes, each scored 0–2, for a maximum of 6 points per bug and 120 points across all 20 bugs. The three axes are designed to independently measure the specific failure modes Unravel's architecture addresses.

Axis	Definition
RCA Root Cause Accuracy	2: Correct file + correct mechanism + correct lines (within ±3). 1: Correct file, wrong mechanism — or correct mechanism, wrong file. 0: Wrong file, wrong mechanism, or refusal to diagnose.
PFR Proximate Fixation Resistance	2: Explicitly names the planted trap and eliminates it with code evidence. 1: Finds correct root cause but never explicitly rebuts the trap. 0: Fixates on the trap or spends >50% of diagnosis on the wrong location.
CFR Causal Flow Reconstruction	2: Full chain — root cause → intermediate state change → symptom, with file/line evidence at each hop. 1: Partial chain — missing intermediate steps, or no evidence for each hop. 0: No causal chain. States fix without explaining mechanism.

The PFR axis is central to Unravel's thesis. Every bug in UDB-20 contains a *planted proximate trap*: a plausible-looking location that a naive model will fixate on, typically the crash site or the file named in the symptom report. PFR measures whether the engine explicitly eliminates the trap with evidence rather than passively arriving at the correct location. An engine that always scores PFR = 1 (passively correct) is susceptible to confident-wrong outputs on harder bugs where the correct location is less obvious.

The CFR axis captures a failure mode complementary to PFR: an engine can identify the correct root cause file and mechanism, then fail to explain *how* that cause produces the observed symptom. A CFR = 0 answer is effectively unactionable — the developer cannot verify the diagnosis or trust the fix without independently retracing the causal chain.

Hallucination penalty. If an engine cites a line number that does not exist in the provided files, or claims a variable appears at a location where it does not, the entire axis in which that hallucinated evidence was used is capped at 0. This is the structural HR rule: hallucinated evidence does not partially degrade a score, it invalidates it.

Mode parity policy. All three output presets (Quick Fix, Developer, Full Report) must produce equal RCA, PFR, and CFR quality. Mode differences are noted separately and flagged as an engine defect, not a valid score variance. This policy reflects the design principle that presets control *output volume*, not *diagnostic depth*: the 8-phase reasoning pipeline and hypothesis elimination run identically regardless of preset.

Difficulty tiers.

Tier	Bugs	Description
Easy	B-01 to B-07	Single-file root cause. Trap in same feature area.
Medium	B-08 to B-14	Cross-file root cause. Trap is a different module.
Hard	B-15 to B-20	Multi-hop causal chain. Trap structurally similar to root cause.

Ground truth and verifiability. Each bug in UDB-20 ships with: (a) `bug.test.*` — a test that fails before the fix, confirming the bug is real; (b) `fixed.test.*` — a test that passes after the ground truth fix, confirming the fix is correct; (c) `ground-truth.md` — exact root cause file, lines, mechanism, proximate trap, and correct fix. Ground truth is locked before any model sees the code. Any researcher can independently run `pnpm test` to verify both the bug and the fix. This is external validity by construction rather than author assertion.

Grading. All axes are graded by the author against `ground-truth.md`. Results are archived in `validation/results/b-XX-<name>/grade.md` alongside raw engine outputs. The full 3-axis rubric and scoring policies are published in `validation/grading/README.md`.

Implementation note. All experiments were conducted using the **Gemini 2.5 Flash free API tier**. The structural grounding approach achieves the reported diagnostic performance without frontier model scale: the AST layer does not require a more powerful model; it requires a more structured context.

Additional metric definition.

- **Hallucination Rate (HR):** Percentage of diagnostic entity claims (file names, line numbers, variable names in the evidence block) that reference entities not present in the provided source files, as detected by the Claim Verifier's six deterministic checks. Formally: HR = |failed Claim Verifier checks|/|total entity claims assessed|.

Scope note. HR measures ***structured entity hallucination*** *only: fabricated file names, line numbers, and variable references in the output schema. It does not capture reasoning hallucination (confabulated causal chains) or behavioral hallucination (unverifiable semantic claims). These remain open problems addressed by the proposed Reasoning Integrity metric (§7.3).*

6.2 Production Validation: Real-World Case Studies

These case studies represent end-to-end applications: Unravel was provided the GitHub issue URL and produced a structured diagnosis. That diagnosis was used to construct a pull request that was accepted by maintainers. Maintainer acceptance confirms the proposed fix was consistent with the repository's intended behavior; it does not establish that the fix was the unique correct solution, nor that Unravel autonomously committed the patch.

Full artifacts for each case (Unravel JSON output, baseline model outputs at temperature = 0, the patch file, reproduction steps, maintainer comments) are archived in `validation/prs/<repo>/<issue>/README.md` to enable independent verification.

6.2.1 Case Study 1: tldraw — CLI Installation Bug (PR #8161)

Repository: tldraw/tldraw, 8,000+ files, TypeScript monorepo. **Bug type:** `DATA_FLOW` + `LAYER_BOUNDARY`

Issue. `npm create tldraw my-app` installed all scaffolded files into the current working directory rather than creating a `my-app/` subdirectory (GitHub issue #8148).

Structural trace. Data flow: `args._[0]` → `maybeTargetDir` → `targetDir = maybeTargetDir ?? process.cwd()`. The fallback to `process.cwd()` executes *before* the interactive `namePicker` prompt. The interactively entered name was used only for `package.json` naming; `targetDir` was never updated to reflect it. A secondary report (files created after user cancellation) was identified as a `LAYER_BOUNDARY` case: `ensureDirectoryEmpty` calls `process.exit(1)` on cancellation, placing the behavior upstream of the provided code. Unravel correctly separated these two diagnoses rather than generating a single incorrect fix.

Fix. After `namePicker` returns, check whether no CLI argument was provided and the entered name differs from `pathToName(process.cwd())`; resolve `targetDir` from the entered name before proceeding.

Outcome. The Unravel diagnosis was used to construct PR #8161. The fix was accepted by maintainers and the issue was closed.

6.2.2 Case Study 2: cal.com — Settings Toggle Contention (PR #28296)

Repository: calcom/cal.com, one of the largest open-source Next.js applications. **Bug type:** `STATE_MUTATION`

Issue. All toggle switches on the Settings page became unresponsive immediately after any one was clicked, until the first API call completed (GitHub issue #28283). No crash: pure behavioral degradation.

Structural trace. Mutation chain for `isUpdateBtnLoading` in `general-view.tsx`: written to `true` at L184 before `mutation.mutate()`, reset to `false` in `onSettled` at L87. A single shared `trpc.viewer.me.updateProfile.useMutation` hook propagated this loading flag to the `disabled` prop of every `SettingsToggle` in the form. Timeline: T0 first toggle clicked → `isUpdateBtnLoading = true` (all toggles disabled) → T1 second toggle blocked → T2 API resolves → `isUpdateBtnLoading = false`. The bug is not at the toggle's render site; it is at the shared hook's mutation point.

Fix. Create a separate `useMutation` hook per toggle, binding each toggle's `disabled` prop only to its own `isPending` state, with an optimistic update to flip the UI immediately.

Outcome. The diagnosis was used to construct PR #28296. The fix was accepted by maintainers. The issue reporter credited the Unravel diagnosis by name in the thread.

6.2.3 Aggregate Metrics

#	Repository	Bug Type	RCA	HR	Outcome
1	tldraw/tldraw #8148	`DATA_FLOW` + `LAYER_BOUNDARY`	✓	0%	Fix accepted, issue closed
2	calcom/cal.com #28283	`STATE_MUTATION`	✓	0%	Fix accepted, credited by reporter

Validation rate: 2/2 (100%). Both diagnoses used to construct patches that were accepted by maintainers, updated as additional cases are validated.

6.3 UDB-20 Results (B-01 to B-10)

B-01 through B-10 have been fully constructed, run, and graded against locked ground truth. Baseline is Claude Sonnet 4.6 at temperature = 0, no AST pre-analysis. B-10 was run twice: once before and once after the `detectListenerParity()` AST improvement; both runs are reported (§6.3.1).

			Unravel (Gemini 2.5 Flash + AST)				Baseline (Claude Sonnet 4.6)			
Bug	**Tier**	**Category**	**RCA**	**PFR**	**CFR**	**/6**	**RCA**	**PFR**	**CFR**	**/6**
B-01 *invisible-update*	Easy	STATE_MUTATION	2	2	2	**6**	2	2	1	5
B-02 *phantom-preference*	Hard	STATE_MUTATION	2	2	2	**6**	2	2	1	5
B-03 *loyal-closure*	Medium	STALE_CLOSURE	2	2	2	**6**	2	2	1	5
B-04 *deaf-plugin*	Hard	STALE_CLOSURE	2	2	2	**6**	2	2	1	5
B-05 *fickle-like*	Medium	ASYNC_ORDERING	1	2	2	5	2	2	1	5
B-06 *silent-await*	Easy	ASYNC_ORDERING	2	2	2	**6**	2	1	2	5
B-07 *ghost-ref*	Medium	ASYNC_ORDERING	2	2	2	**6**	2	2	2	**6**
B-08 *twin-effect*	Medium	EVENT_LIFECYCLE	2	2	2	**6**	2	2	2	**6**
B-09 *trusted-input*	Hard	TYPE_COERCION	2	2	2	**6**	2	2	2	**6**
B-10 *orphan-listener*	Hard	EVENT_LIFECYCLE	2	2	2	**6**[†]	1	2	1	4
Total /60			**19**	**20**	**20**	**59**	19	19	14	**52**

[†]B-10 score reflects the post-fix run after `detectListenerParity()` was added to the AST engine. Pre-fix Unravel scored 4/6 on B-10 (wrong mechanism). Both runs are reported in §6.3.1.

Table 1: UDB-20 results (B-01–B-10). RCA, PFR, CFR each scored 0–2; per-bug max 6, ten-bug max 60. Unravel: 59/60 (98.3%). Baseline: 52/60 (86.7%). Delta: +7. Grade sheets: `validation/results/b-XX-<n>/grade.md`.

Axis	**Unravel /20**	**Baseline /20**	**Δ**	**Pattern**
RCA	19	17	+2	Unravel wrong file on B-05 (right mechanism); baseline wrong
PFR	20	19	+1	Baseline half-credit on B-06: trap acknowledged but not fo
CFR	20	14	+6	Unravel full multi-hop chains on all 10; baseline dropped hops
Total	**59**	**52**	**+7**	

Table 2: Per-axis aggregate across B-01–B-10. CFR is the dominant differentiator.

Key findings from B-01–B-10.
CFR is the dominant differentiator. *Unravel scored CFR = 2 on all ten bugs (20/20); the baseline scored 14/20. The gap is not in identifying the root cause file — both engines locate the correct file on most bugs — but in tracing the complete causal chain from root mutation to observed symptom. On B-04 (four-file chain: Application → PluginRegistry → NotificationPlugin → handler closure → stale config read), Unravel traced all four hops with file and line evidence; the baseline omitted the PluginRegistry step. This pattern repeated on every multi-hop bug in B-01–B-05 and on B-10.*
PFR held for both engines across all difficulty tiers. *Both engines scored PFR ≥ 1 on every bug. The planted proximate traps were recognized and rebuffed in nearly every case, with the single exception of B-06 where the baseline acknowledged the trap but did not formally eliminate it (PFR = 1 vs Unravel's PFR = 2).*
B-05 and B-10 are the two non-perfect Unravel scores, and they are instructive in opposite ways. *B-05: Unravel found the correct mechanism (unconditional overwrite via missing version guard) but diagnosed the wrong file (`optimisticStore.ts` vs ground truth `likeService.ts`). The baseline found the correct file. This is the case where caller-side design intent was not visible in the AST. B-10: the pre-fix Unravel diagnosed the correct file but anchored on a visually asymmetric but mechanically irrelevant detail — the missing `{passive: true}` options in `removeEventListener`. The baseline correctly dismissed this. One 47-word AST annotation (§6.3.1) eliminated the false anchor entirely.*
B-09 fix quality: the baseline wins. *Both engines scored 6/6 on all three axes for B-09, but the fix quality differed. Unravel used `parseFloat(String(total).replace(non-numeric chars, ")) || 0` — the `|| 0` fallback silently converts completely invalid input to zero rather than rejecting it, which can produce misleading validation outcomes. The baseline used `parseFloat(total)` with an explicit `isNaN` guard returning HTTP 400, and replaced `as unknown as CartSummary` with `satisfies CartSummary` for compile-time safety. The rubric scores diagnosis, not fix quality; this distinction is noted for completeness.*

6.3.1 B-10: Two Runs and What One AST Annotation Changed

B-10 (*orphan-listener*) is the only bug in UDB-20 where the AST engine was modified between runs. Both runs are reported in full; only the post-fix score appears in the aggregate table, and the pre-fix score is explicitly noted.

Pre-fix run (Unravel 4/6, Baseline 4/6). The ground truth root cause is accumulated scroll listeners from `PageView` components not being properly unmounted during SPA navigation — each navigation adds a listener without the previous one being removed, producing exactly N listeners after N navigations. Unravel (pre-fix) scored RCA = 1: it identified the correct file (`useScrollAnalytics.ts`) but diagnosed the wrong mechanism, fixating on the syntactic asymmetry between `addEventListener('scroll',`

`fn, {passive: true})` and `removeEventListener('scroll', fn)` (no options object). This is a *visual proximate fixation* failure: the code *looked* asymmetric, so the model treated the asymmetry as evidence before checking whether it was mechanically significant. Per the DOM specification, `passive` is not part of the listener identity key for removal; only `capture` is. The proposed fix (extracting the options object into a `const`) was a no-op.

The baseline (Claude Sonnet 4.6) also scored 4/6 on B-10, but for structurally different reasons: it correctly dismissed the `passive` mismatch, correctly hypothesized concurrent mounted `PageView` instances as the accumulation mechanism, but was unable to confirm the hypothesis from the three provided files and explicitly flagged its diagnosis as inconclusive. Correct reasoning, insufficient context.

Post-fix run (Unravel 6/6). One function was added to `ast-engine-ts.js`: `detectListenerParity()` which compares each `addEventListener`/`removeEventListener` pair and outputs a structured annotation. When the options objects differ only in `passive` or `once` (not `capture`), the annotation explicitly states: *"passive/once omission is harmless (spec: only capture affects identity)."* This annotation is 47 words. It appears in the verified AST ground-truth block before the LLM sees the code.

On the post-fix run, hypothesis H1 (`passive` mismatch) was eliminated in Phase 5 immediately, citing the AST annotation verbatim. Hypothesis H3 (multiple `PageView` instances not unmounting) survived elimination and was confirmed as the root cause. The full causal chain — mount A, navigate, router fails to unmount A, mount B, both `onScroll` closures fire, duplicate `track()` calls — was traced correctly with `isBugPoint: true` placed at the missed unmount edge.

The lesson is direct: the LLM did not hallucinate about `passive` because it was careless or underpowered — it hallucinated because the visual evidence pointed to a plausible-looking mismatch and no authoritative counter-fact was available. Once the AST provided the counter-fact as a verified ground-truth annotation, the hallucination did not occur. Deterministic structural knowledge does not merely help the LLM reason better; it eliminates entire categories of false hypothesis before they can anchor.

6.4 Comparative Evaluation

To contextualize Unravel's output, we conducted a qualitative comparison across two hand-constructed test cases, evaluated against Claude 4.6, ChatGPT 5.3, and Gemini 3.1 Pro. All models received identical code context at temperature = 0. These results are illustrative rather than statistically conclusive due to sample size; they characterize output categories, not model intelligence rankings.

6.4.1 Test 1: The Heisenbug (Single File)

A race condition in a dashboard initializer where two async operations mutate shared rendering state independently. The Heisenbug property: adding `console.log` for debugging changes microtask scheduling just enough to eliminate the bug, making it unreproducible under observation.

Table A: Core Diagnostic Correctness

Criterion	Claude	GPT-5.3	Gemini	Unravel
Bug category (`RACE_CONDITION`)	✓	✓	✓	✓
Heisenbug identified	✓	×	✓	✓
Microtask scheduling explained	✓	×	✓	✓
Spread non-atomicity caught	Approx	✓	✓	✓
Correct fix (`Promise.all`)	✓	✓	✓	✓
Error handling in fix	×	×	✓	✓

Table B: Extended Output Capabilities

Output capability	Claude	GPT-5.3	Gemini	Unravel
Developer debugging path predicted	×	×	×	✓
Variable state lifecycle	×	×	×	✓
Invariants documented	×	×	×	✓(4)
Execution timeline	Basic	Basic	×	Full
Competing hypotheses	×	×	×	✓
Structured JSON output	×	×	×	✓

6.4.2 Test 2: The Phantom Accumulator (5 Files)

A 4-way emergent bug across 5 files: a memoization reference bug, a WeakRef GC footgun, a microtask async ordering issue, and a Heisenbug observation effect combined.

Table A: Core Diagnostic Correctness

Criterion	Claude	GPT-5.3	Gemini	Unravel
Root cause correct	✓	✓	✓	✓
Mutation sites + exact lines	✓	✓	✓	✓(L22, L28, L32)
`resetBoard` non-contributing	✓	×	×	✓
Minimal fix (not full rewrite)	✓	×	✓	✓
Timestamped failure trace	Basic	×	×	Full (0s–10.5s)
Causal loop traced (8 steps)	✓	Partial	×	✓

Table B: Extended Output Capabilities

Output capability	Claude	GPT-5.3	Gemini	Unravel
Variable lifecycle tracker	×	×	×	✓(2 vars)
Invariants documented	×	×	×	✓(8)
Competing hypotheses	×	×	×	✓(3)
Structured reproduction steps	×	×	×	✓
AI fix prompt	×	×	×	✓
Structured JSON output	×	×	×	✓

Key finding. *On core diagnostic correctness (Table A), all four tools identified the root cause on both test cases, confirming that frontier models are sufficient for tractable bugs when given adequate context. The divergence is in Table B: Unravel additionally produces output categories (invariants, variable lifecycle traces, hypothesis enumerations, timestamped causal chains, reproduction steps, structured JSON) not observed in any baseline output. These categories are not produced by prompting; they are structural outputs of the 8-phase pipeline.*

Limitation. These are two hand-selected test cases. They are not randomly sampled and are not representative of a distribution. Statistical conclusions cannot be drawn.

6.5 Historical Context: UDB-11

UDB-11 was the original 11-bug internal benchmark used during Claim Verifier design. It is retained here as historical context, not as a primary benchmark result. Its design limitations — self-referential construction, single-axis scoring, and author selection bias — are superseded by UDB-20.

Metric	Baseline (no AST)	Unravel (full pipeline)
RCA	91% (10/11)	100% (11/11)
HR	1.5% (1/66)	0.0% (0/66)

The baseline HR of 1.5% reflects a fabricated line number on one diagnosis (cited line exceeded file length). The Claim Verifier catches this deterministically regardless of model instruction compliance.

UDB-11's primary contribution was motivating the Claim Verifier design and demonstrating that AST pre-analysis produces a measurable signal even on an 11-bug sample. The real-world merged PRs (§6.2) constitute the primary correctness evidence.

7. Discussion

7.1 Structural Understanding as a General Primitive

A system that traces the complete structural behavior of a codebase (which variables are mutated at which lines, which functions propagate data across file boundaries, where async execution diverges from lexical order, what invariants the code implicitly relies on) has acquired knowledge beyond what debugging requires. Debugging is one application of structural understanding.

Code explanation. Unravel's Explain mode uses the same Layer 0 facts to produce entry point hierarchies, data flow maps, and dependency graphs. Structural understanding is the hardest part of developer onboarding; it is automated deterministically.

Security auditing. Security mode uses the call graph and mutation chains to trace data from user-controlled input to sensitive operations. Injection vulnerabilities and authentication bypasses are structural properties of the code: they appear in the AST whether or not the model is explicitly told to look for them.

Feature implementation. If Unravel knows which variables hold which invariants and which callers will be affected by a change, it can reason about *where to add a feature* without introducing structural violations. Current agentic coding tools generate code without structural knowledge of the codebase they are generating into. A code generation agent that has not verified which invariants must be preserved will, predictably, violate some of them.

7.2 Structural Grounding vs. Context Window Scaling

A predictable objection: frontier models with 1M-token context windows can ingest entire repositories, eliminating the need for selective routing. This addresses context size but not the structural problem:

Hypothesis elimination. Long-context models still anchor to the symptom description. Injecting 500K tokens of source code does not change the fact that the model pattern-matches from the symptom to the nearest familiar-looking code [5]. The AST fact that `duration` is written in `pause()` at L69 *eliminates* the stale-closure hypothesis that the model would otherwise anchor on. No amount of context makes this elimination happen automatically: it requires the verified structural constraint.

Ground truth injection. The statement "duration is written at L69, verified by static analysis" is categorically different from asking the model to infer that mutation from reading 80 lines of code. The former is a deterministic constraint injected before token generation begins; the latter is a probabilistic inference over the training distribution.

Cost. Unravel achieves the reported results on the Gemini 2.5 Flash free tier, demon-

strating that structural grounding extracts strong diagnostic performance without requiring frontier model scale.

7.3 Reasoning Integrity: A Proposed Future Metric

The Claim Verifier checks whether *answers* are grounded in real code. It does not verify whether the *reasoning path* is grounded. Three failure modes are distinct:

Type	Description	Caught by current system
1	Wrong answer, fabricated evidence	Yes: Claim Verifier
2	Correct answer, fabricated reasoning	No
3	Mostly-correct reasoning, one unjustified inferential leap	No

Type 2 is the dangerous case: the fix works, the developer trusts the diagnosis, but the reasoning was confabulated. The next similar bug is misdiagnosed using the same flawed logic.

We propose **Reasoning Integrity (RI)** as a metric: for each hypothesis elimination in Phase 5, extract the factual claim and verify it against the AST scope map and mutation chains. RI is the percentage of elimination steps in which the model's factual claims are AST-verified. This enables characterizations such as: *"Model X: RCA 100%, RI 78%"*, distinguishing models that get correct answers by correct reasoning from those that get correct answers by confabulation. The infrastructure is already resident in Unravel: the AST facts and the Phase 5 `hypothesisTree` output are both in memory simultaneously; the cross-check is not yet automated.

7.4 Limitations

Language coverage. AST pre-analysis is currently limited to JavaScript, TypeScript, JSX, and TSX. Python, Go, Rust, and other languages receive the full 8-phase LLM pipeline but without structural ground-truth injection. tree-sitter grammars exist for all common languages; extension is structurally straightforward.

Static analysis scope. Without runtime execution, structural analysis cannot verify which code paths are actually taken under specific inputs. A mutation present in the AST may be on a dead code path. A race condition that manifests only under specific scheduling may not be statically identifiable.

Floating promise precision. The `isAwaited` guard correctly handles most cases but can produce false positives in complex chained callback expressions where upward parent traversal encounters ambiguous scope boundaries.

Benchmark scale and methodology. Ten of UDB-20's 20 bugs have been graded

at time of writing. The results (59/60 Unravel vs. 52/60 baseline) provide directional signal; the remaining ten bugs (B-11–B-20) are required for full statistical coverage. All 20 bugs are constructed by the same author who designed the system; grading is human against pre-locked ground truth, not independently replicated. The planted proximate traps were designed by the author and not independently validated as genuinely misleading before embedding. B-10 was re-run after an AST improvement; only the post-fix score appears in the aggregate, with the pre-fix score reported separately (S6.3.1). A fully independent replication would substantially strengthen all quantitative claims.

Failure case: scheduling-dependent race condition. In one analysis involving a race condition where two concurrent async writes converged on the same state object, AST analysis correctly identified both mutation paths and flagged the pattern as `RACE_CONDITION`. However, static analysis could not determine which of the two execution orderings was realized at runtime: both were structurally possible. Unravel returned an ambiguous diagnosis with two surviving hypotheses and confidence 0.61, noting that disambiguation required runtime instrumentation. This is an expected limitation of static-only analysis.

7.5 Benchmark Design Evolution and the Evaluation Validity Problem

The evaluation strategy for Unravel has evolved through three distinct stages, each exposing a limitation in the previous approach. Documenting this evolution is itself useful: the failure modes encountered generalize to any system whose claimed advantage is specifically about complex, multi-file, real-world debugging.

Stage 1 — Large but transparent synthetic bugs. The earliest internal testing used a suite of deliberately constructed bugs, each designed to be architecturally complex: cross-file state corruption, async boundary violations, layered causation. The intent was to stress-test the system. The outcome was unexpected: both Unravel and unaugmented frontier models solved nearly all of them reliably, and the performance delta was small. Post-analysis identified the reason. The bugs, despite their architectural sophistication, were contained in small, focused codebases specifically written to isolate the defect. With minimal surrounding code, the root cause was effectively visible in the files provided — there was nowhere else to look. Unravel's structural grounding still produced a measurable advantage, but the signal was weak because the alternative hypothesis space was small. A system built specifically for large codebases does not reveal its edge when tested on small ones.

Stage 2 — Real GitHub issues. The limitations of Stage 1 motivated applying Unravel to real open-source GitHub issues: tldraw issue #8148 (8,000+ file TypeScript monorepo) and cal.com issue #28283 (one of the largest open-source Next.js codebases). On these cases Unravel performed qualitatively differently. The file routing

problem, the cross-file causal chain reconstruction, and the hypothesis elimination under high-noise conditions all became non-trivial. The structured diagnosis it produced was used to construct patches accepted by maintainers. This is the external validation signal the synthetic bugs could not provide.

However, GitHub issues introduce a different evaluation problem. Obtaining a PR acceptance for each additional case requires submitting patches to real repositories and waiting for maintainer review — a process that may take days to weeks per case, is subject to factors entirely outside diagnostic quality (maintainer availability, competing PRs, project roadmap), and cannot be run in controlled comparative conditions. Two cases provide proof of concept; they do not provide a repeatable, ablation-ready benchmark.

Stage 3 — UDB-20: constructed bugs with realistic structural properties. UDB-20 represents an attempt to occupy the space between Stages 1 and 2: bugs constructed by the author but designed to exhibit the structural properties that make real GitHub issues hard. Each bug has a deliberately planted proximate trap, a multi-file causal chain, and a ground truth locked before any model runs. The test suite verifies both the bug and the fix independently via `pnpm test`. The three-axis rubric (RCA, PFR, CFR) is designed to directly measure the failure modes Unravel addresses rather than binary pass/fail.

The results across B-01–B-10 (Unravel 59/60 vs. baseline 52/60, delta +7) confirm a consistent advantage on CFR — the causal chain reconstruction axis (20/20 vs. 14/20) — while showing near-parity on RCA (19/20 vs. 17/20) and PFR (20/20 vs. 19/20). The author's assessment is that this delta, while real, understates Unravel's advantage on genuinely large-scale real-world bugs. UDB-20's bugs, despite deliberate effort to make them structurally realistic, remain smaller and more self-contained than a typical multi-file issue in a production codebase with tens of thousands of lines of surrounding context. The hypothesis space is still narrower than in the wild. The structural grounding layer contributes most when the number of plausible files and causal paths is large; at smaller scales, both augmented and unaugmented models can arrive at the correct answer by different routes.

The correct interpretation of current results. The three tiers of evidence in this paper should be read as measuring different things. The GitHub PR case studies (§6.2) measure whether Unravel produces correct, actionable diagnoses on real production bugs at the scale and complexity of codebases with thousands of files — they do not support comparative claims. UDB-20 (§6.3) measures a repeatable performance delta against a controlled baseline — it does not yet reflect the full difficulty distribution of real-world bugs. The expected pattern, consistent with both theoretical motivation and preliminary results, is that Unravel's advantage grows with bug complexity, codebase size, and causal chain length. Validating this prediction rigorously requires either

a larger constructed benchmark with provably realistic complexity, or a systematic real-world validation program that pairs identical bug inputs to both Unravel and a baseline and evaluates the resulting diagnoses against the ground truth established by maintainer-accepted fixes. Both remain open work.

8. Future Work

8.1 Heavy Mode: Four-Agent Parallel Analysis. For repositories too large for a single context window, a four-agent parallel system distributes the codebase: Agents 1, 2, and 3 each receive a distinct non-overlapping partition and run the full 8-phase pipeline independently. Agent 4 (Synthesis) receives all three JSON outputs, resolves agreement and contradiction, and issues a final merged diagnosis weighted by per-agent confidence. This is a parallelism strategy, not a debate: agents do not see each other's work during analysis.

8.2 Three-Agent Sequential Divergent System. Activated when single-agent confidence falls below a tuned threshold. Agent 2 receives Agent 1's summary as a negative prompt: it must look in a completely different direction and cannot reproduce Agent 1's hypothesis. Agent 3 receives both summaries as negative prompts. If all three agents converge independently, confidence in the shared hypothesis is elevated. If they diverge, all three diagnoses are surfaced: the developer receives the full uncertainty picture rather than a false consensus.

8.3 Reasoning Integrity Verifier. Automated implementation of the RI metric described in §7.3: for each hypothesis elimination in Phase 5, extract the factual claim, cross-check it against the AST scope map and mutation chains, and flag reasoning steps that are factually unverifiable.

8.4 WebContainers Runtime Execution. Run reproduction steps in an in-browser sandboxed Node.js environment. Capture concrete variable values at the crash site and inject them as additional verified facts alongside AST ground truth. This converts "AST indicates this mutation could cause the bug" into "runtime confirms value was 1490, not 1500."

8.5 Multi-Language AST Support. tree-sitter grammars exist for Python, Go, Rust, Java, C, and C++. The scope resolver, mutation extractor, and closure detector are language-agnostic in design; extension requires grammar-specific node type mapping.

8.6 Semantic Grounding Layer. Embedding-based similarity for symptom-to-known-bug matching (inject top-3 confirmed prior bugs as grounding context before the LLM call), confidence calibration via response-consistency clustering across multiple temperature-0.3 runs, and hypothesis deduplication by collapsing semantically

identical candidates before Phase 4.

8.7 Long-Term Vision. The structural intelligence layer that makes debugging reliable (deterministic pre-analysis, call graph routing, hypothesis elimination, verified output) is the same layer that makes all AI code interaction reliable. The natural evolution is toward an IDE-level environment in which structural analysis precedes every AI interaction with the codebase: feature additions, refactoring, security review, onboarding. Structural pre-analysis is that layer.

9. Conclusion

We presented Unravel, a debugging engine built on the thesis that deterministic structural pre-analysis is a necessary layer, not an optimization, without which LLM reasoning over code is architecturally unreliable. The four-layer architecture implements the generate-verify-eliminate loop that hypothesis-driven debugging research identifies as the primary mitigation for anchoring bias and proximate fixation.

The meaningful empirical record is tldraw PR #8161 and cal.com PR #28296: real production bugs in repositories with thousands of files, diagnosed correctly using Gemini 2.5 Flash on the free API tier, with the exact causal mechanism confirmed by maintainer acceptance. UDB-20 (§6.1.1) provides a reproducible three-axis evaluation framework — RCA, Proximate Fixation Resistance (PFR), and Causal Flow Reconstruction (CFR) — designed to directly measure the failure modes Unravel targets. Results from the first ten bugs (59/60 Unravel vs. 52/60 baseline, delta +7) show Unravel's advantage concentrated on CFR (20/20 vs. 14/20): complete multi-hop causal chains traced on all ten bugs versus the baseline dropping hops on six. B-10 additionally demonstrated that a 47-word deterministic AST annotation eliminated an entire false hypothesis category — converting a 4/6 pre-fix run to 6/6 — without any change to the prompt or model. The remaining ten bugs (B-11–B-20) are forthcoming.

Larger models will continue to improve at pattern matching over code. The structural reasoning gap (knowing which variables are mutated at which lines, which callers break if a function changes, what invariants a fix must preserve) does not close through scale alone. It requires verified structural facts about the specific code being analyzed. AST pre-analysis surfaces exactly those facts. The model does not become more capable. It finally has something verified to reason about.

References

[1] Jimenez, C.E., et al. (2024). "SWE-BENCH: Can Language Models Resolve Real-World GitHub Issues?" *ICLR 2024*. https://proceedings.iclr.cc/paper_files/paper/2024/file/edac78c3e300629acfe6cbe9ca88fb84-Paper-Conference.pdf

[2] Liu, et al. (2023). "Large Language Models of Code Fail at Completing Code with Potential Bugs." *NeurIPS 2023*. https://proceedings.neurips.cc/paper_files/paper/2023/file/819cebb05f993840e8a52d7564c5c282-Paper-Conference.pdf

[3] Arora, M., et al. (2026). "Cognitive Biases in LLM-Assisted Software Development." *arXiv:2601.08045*. https://arxiv.org/abs/2601.08045

[4] "The Debugging Decay Index." (2026). *arXiv:2506.18403*. https://arxiv.org/abs/2506.18403

[5] Zhang, et al. (2026). "Behavioral and Attributional Evidence of Anchoring Bias in LLMs." *arXiv:2511.05766*. https://arxiv.org/abs/2511.05766

[6] Cai, Z., et al. (2023). "Unleashing the Power of Type-Based Call Graph Construction." *USENIX Security 2023*. https://www.usenix.org/system/files/sec23winter-prepub-350-cai.pdf

[7] "Multi-SWE-bench: A Multilingual Benchmark for Software Engineering." *NeurIPS 2025*. https://neurips.cc/virtual/2025/poster/121692

[8] "Agentless: Demystifying LLM-Based Software Engineering Agents." *arXiv:2407.01489 / findings cited in arXiv:2509.13941*. https://arxiv.org/abs/2509.13941

[9] Ferrante, J., Ottenstein, K.J., Warren, J.D. (1987). "The Program Dependence Graph and Its Use in Optimization." *ACM TOPLAS*. https://web.eecs.umich.edu/~mahlke/courses/583f23/reading/ferrante_toplas_87.pdf

[10] Horwitz, S., Reps, T., Binkley, D. (1990). "Interprocedural Slicing Using Dependence Graphs." *ACM TOPLAS*.

[11] Aho, A.V., Lam, M.S., Sethi, R., Ullman, J.D. (2006). *Compilers: Principles, Techniques, and Tools*. Addison-Wesley.

[12] "Unveiling the Power of Intermediate Representations for Static Analysis." *arXiv:2405.12841*. https://arxiv.org/abs/2405.12841

[13] "When Threads Meet Events: Precise Static Race Detection with O2." O2Lab Technical Report. https://o2lab.github.io/p/o2.pdf

[14] "LDB: A Large Language Model Debugger via Verifying Runtime Execution Step by Step." *arXiv:2402.16906*. https://arxiv.org/abs/2402.16906

[15] "TraceCoder: A Trace-Driven Multi-Agent Framework for Automated Debugging." *arXiv:2602.06875*. https://arxiv.org/abs/2602.06875

[16] Chen, J., et al. (2024). "H-Neurons: On the Existence, Impact, and Origin of Hallucination-Associated Neurons in LLMs." *arXiv:2512.01797*. https://arxiv.org/abs/2512.01797

[17] tldraw PR #8161 (merged). https://github.com/tldraw/tldraw/pull/8161

[18] cal.com PR #28296 (merged). https://github.com/calcom/cal.com/pull/28296

[19] Unravel source code (BSL-1.1 license). https://github.com/EruditeCoder108/

unravelai

[20] Unravel web application (live). https://vibeunravel.netlify.app

---

Submitted to arXiv cs.SE. March 2026. *Source: https://github.com/EruditeCoder108/unravelai*