Unravel: A Deterministic Pre-Analysis Layer for LLM-Assisted Software Debugging

Sambhav Jain

Independent Research

https://github.com/EruditeCoder108/unravelai eruditespartan@gmail.com

March 2026

Large Language Models (LLMs) have demonstrated remarkable capabilities across a wide range of tasks, particularly in code generation and reasoning. These models are trained on vast corpora of text and code, enabling them to produce coherent and often highly useful outputs. However, their responses are fundamentally probabilistic—grounded in learned statistical patterns rather than verified structural understanding.

As a result, LLMs exhibit several well-known limitations. When faced with incomplete context or insufficient grounding, they tend to hallucinate—producing confident but incorrect outputs. Their performance is also highly sensitive to context quality, prompt structure, and the availability of relevant information. These limitations become especially pronounced in real-world software debugging, which remains one of the most complex and critical aspects of software engineering. Unlike code generation, debugging requires: *(i)* deep contextual awareness across large, interconnected codebases; *(ii)* precise reasoning about system state and execution behavior; and *(iii)* iterative hypothesis testing and systematic elimination.

In large-scale production systems—codebases with thousands of files and deeply interconnected logic—current LLMs struggle to reliably identify root causes of issues. Even as newer models improve on coding benchmarks, they still fall short in complex multi-file debugging scenarios. This raises a key question:

> *What if, instead of relying on the LLM to infer the structure of a massive codebase, it was provided with a deterministic map? What if the model knew exactly where, what, and how every variable mutated before it generated a single token? Furthermore, what if the model was constrained by a rigorous reasoning pipeline equipped with strict anti-sycophancy rules—forcing it to explore multiple directions, sys-*

tematically eliminate its own hypotheses using verified evidence, and select only the rigorously proven root cause?

In this paper, we introduce **Unravel**, a four-layer debugging engine built on this premise. Layer 0 runs a deterministic AST pre-analysis (web-tree-sitter, WebAssembly) before any LLM call, extracting mutation chains, closure captures, async boundaries, and cross-file call graphs as verified ground truth. Layer 3's Claim Verifier deterministically rejects any output fabricating file names, line numbers, or variable references—independently of what the model was instructed. Together, these layers prevent the model from reasoning inconsistently with information already known to be true.

Results. Two real-world diagnoses validated and independently implemented by maintainers of tldraw (PR #8161, 8,000+ file TypeScript monorepo) and cal.com (PR #28296); a third diagnosis for claude-hud (PR #259) is open for review. Notably, in a direct comparison of Unravel against Claude Sonnet 4.6, Unravel operated with Gemini 2.5 Flash as its underlying LLM—a model that scores significantly lower than Claude Sonnet 4.6 across standard coding and reasoning benchmarks (see p. 32 for full benchmark scores). Despite this asymmetry in model capability, Unravel achieved 119/120 (99.2%) on the UDB-20 benchmark (20 bugs, three-axis rubric) vs. Claude Sonnet 4.6 baseline 112/120 (93.3%), delta +7. Full benchmark including three Super-Bug case studies: 137/138 (99.3%) Unravel vs. 130/138 (94.2%) Claude, delta +7. B-22 documents a complete improvement cycle: Run 1 produced an honest 5/6 loss with three identified gaps; four targeted fixes were applied; Run 2 confirmed 6/6. Super-Bug case study (8 files, 1,100 lines, extreme-difficulty race condition): Claude Sonnet 4.6 scored 6/6 *without* AST context – the bug was within its reliable inference window. Gemini 2.5 Flash without AST scored 4/6 (Run 1), exposing the boundary where raw model capability becomes insufficient. With a 150-line targeted AST detector, Gemini 2.5 Flash matched Claude Sonnet 4.6 at 6/6 (Run 2). The tie is evidence that AST grounding closes the capability gap for a weaker model at its limit – not that a frontier model requires structural assistance for bugs already within its reach. All Unravel experiments use Gemini 2.5 Flash on the free API tier.

Keywords: automated debugging, multi-file defect localization, static analysis, AST pre-analysis, constraint-based code reasoning, hallucination mitigation, proximate fixation, hypothesis elimination, claim verification, anti-sycophancy,

program dependence graph

1. Introduction

LLM-based debugging tools fail systematically on a specific class of bug: defects where the observable symptom and the causal origin are separated by multiple function calls, file boundaries, or async execution steps. This is not a matter of model scale. Claude Opus 4.6, Gemini 3.1 Pro, and GPT-5.4 all exhibit the same failure mode: they attribute the defect to the most proximate, structurally familiar code rather than tracing state backward through the actual mutation sequence [3, 4].

The root cause of this failure is architectural. The model receives a *description* of what the code should do and a natural language account of the symptom. The structural facts it needs – which variables are mutated at which lines, which closures capture stale state, which callers will be affected by a proposed fix – are present in the code but are not surfaced in a machine-verifiable form. Without these facts as explicit constraints, the model is free to confabulate plausible-sounding evidence for whichever hypothesis is most consistent with its training distribution [2]. The problem is not that models are insufficiently capable; it is that the reasoning is unconstrained.

On SWE-bench (2,294 real GitHub issues), the best single-model baseline achieved a 1.96% resolution rate [1]. Agentic improvements have raised resolution to 40–55% on selected categories, but the architecture of failure is unchanged: 491 of 497 successfully resolved SWE-bench issues required modifying only a single file [8]. Multi-file bugs, which constitute the majority of non-trivial production defects, remain largely unresolved.

We present **Unravel**. Before any model sees the code, a deterministic AST pre-analysis layer runs first. It extracts every variable mutation chain, every closure capture, every async boundary, every floating promise, every cross-file call graph – as verified structural facts. These become ground truth injected into the reasoning context. The model cannot hallucinate about entities that do not appear in these facts. It cannot guess at mutation sequences. It must trace them.

The contribution is not that the model is given more information. It is that it is prevented from reasoning inconsistently with information already known to be true.

This is enforced architecturally: a post-generation Claim Verifier deterministically rejects any output that fabricates file names, line numbers, or variable references – independently of what the model was instructed. An 8-phase structured pipeline

requires every hypothesis elimination to cite the exact verified code fragment that contradicts it. The result is exact file, exact line, exact mechanism, with evidence at every step. The system reduces structured entity hallucination to zero on verified fields. It does not eliminate reasoning errors or wrong causal chains; it narrows the space in which they can hide and makes them attributable when they occur.

Contributions:

1. **Constraint-Enforced Debugging Architecture:** a four-layer system in which deterministic AST facts are extracted before any model call, injected as verified ground truth, and enforced post-generation by a Claim Verifier that structurally rejects outputs contradicting those facts. Unlike prior work that augments prompts with additional context, Unravel enforces consistency between model output and known structural program facts independently of model instruction. The contribution is architectural: correctness is enforced, not encouraged.
2. **Seven Anti-Sycophancy Rules:** formally specified guardrails targeting the two most prevalent reasoning failures: Proximate Fixation (tendency to attribute failure to the nearest code) and the Name-Behavior Fallacy (assuming execution semantics from identifier names).
3. **Symptom Contradiction Module:** a pre-LLM heuristic pass that detects mismatches between the user's symptom description and AST evidence before any model call is made, preventing the LLM from accepting a false premise as its starting constraint.
4. **CFG Branch Annotation:** every mutation in the AST output is annotated as conditional or unconditional (`[CONDITIONAL]` tag). Without this, the LLM treats branch-gated mutations as unconditional evidence, causing correct hypotheses to be incorrectly eliminated.
5. **Hypothesis Elimination Scoring:** the `eliminatedBy` schema field requires every hypothesis elimination in Phase 5 to cite the exact AST-verified file and line that contradicts it. A hypothesis without a code citation is flagged as unverified reasoning, making confident-wrong outputs structurally distinguishable from evidence-backed ones.
6. **Fix Completeness Verifier:** a zero-latency cross-file check using the already-built call graph to identify callers of modified functions omitted from a proposed fix.
7. **Layer Boundary Detector and External Fix Target Verdict:** a formally specified solvability check that identifies when the root cause originates

upstream of the provided codebase (`LAYER_BOUNDARY`), and a complementary verdict (`EXTERNAL_FIX_TARGET`) that preserves the full diagnosis when the fix target resides in a different repository.

8. **Real-world validation:** Unravel produced structured diagnoses used to construct pull requests reviewed and merged by maintainers of tldraw (8,000+ file TypeScript monorepo) and cal.com (one of the largest open-source Next.js applications), providing external validation that diagnoses matched the repositories' intended behavior.
9. **Reasoning Integrity (RI) Metric:** a formally defined evaluation metric distinguishing models that reach correct answers by correct reasoning from those that reach correct answers by confabulation. For each hypothesis elimination in Phase 5, the factual claim underlying the elimination is cross-checked against the AST scope map and mutation chains; RI is the percentage of elimination steps where the model's factual claims are AST-verified. This enables characterizations such as *"RCA 100%, RI 78%"* – a distinction invisible to correctness-only benchmarks. Four RI failure types are now formally defined (§7.3), the fourth identified empirically from B-22: correct annotation-scoped fix where the annotated violation is a symptom of a deeper protocol error. We propose this taxonomy as an open problem for the field.

2. Related Work

2.1 LLM Failure Modes in Debugging

Hallucination in code reasoning. Liu et al. [2] demonstrate that LLMs given a buggy prefix code context perform *worse* than zero-context baselines on code completion—producing outputs that follow the pattern of the buggy context rather than detecting the defect, with pass rates dropping over 50%, consistent with the model reinforcing rather than correcting the introduced error. Research on hallucination-associated neurons [16] provides a mechanistic explanation: specific neuron clusters activate during generation of factually unsupported content, producing statistically consistent but structurally false claims. Unravel's design addresses this at the prompt construction level: verified AST facts constrain the plausible-sounding hypothesis space before the model's token generation begins.

Proximate fixation. LLMs systematically attribute failures to the crash site rather than the distant corrupted state origin. Zhou et al. [3] document this pattern in a comprehensive study of cognitive biases in LLM-assisted development (ICSE 2026), finding

that developer-LLM interactions account for 56.4% of biased actions. The most prevalent bias categories identified are *Instant Gratification* (19.3%) and *Suggester Preference* (12.6%)—reflecting patterns of uncritical acceptance of LLM-generated suggestions without independent verification of their structural correctness. The Debugging Decay Index [4] (published in *Scientific Reports*) demonstrates that LLM debugging effectiveness decays exponentially across successive attempts, with most models losing 60–80% of their debugging capability within 2–3 attempts – consistent with a failure to revise the initial hypothesis rather than explore alternatives.

Cross-file reasoning degradation. Out of 497 successfully resolved SWE-bench issues via the Agentless framework, 491 required modifying only a single file [8]. Multi-SWE-bench confirms universal performance degradation on cross-file issues across eight languages [7].

Anchoring bias. Log-probability analyses confirm that initial hypothesis tokens occupy disproportionately high-weight positions in the context window, creating a feedback loop that inhibits hypothesis revision [5]. Valencia-Clavijo [5] demonstrates this using Shapley-value attribution across six open-source models, showing that anchoring effects are robust and measurable at the internal probability level – not merely a surface-output artifact. RLHF-aligned production models exhibit a related pattern: sycophantic responding, where a user-proposed incorrect hypothesis receives supporting evidence rather than contradiction (a design concern Unravel's anti-sycophancy rules directly address, independently of which model is used).

2.2 Program Graph Representations

Def-use chains and Data Flow Graphs model variable definitions and uses across function boundaries, enabling detection of state corruption paths [11, 10]. Unravel's `extractMutationChains()` is a practical deterministic implementation of def-use chain analysis with line-level attribution.

Program Dependence Graphs (PDGs) unify control and data dependencies, enabling backward program slicing from a failure point to all statements that transitively influenced corrupted state [9]. Unravel's Phase 5 DIAGNOSE step enforces this backward tracing discipline explicitly.

Call Graphs map cross-file function invocations and are essential for identifying cascading fix requirements [6]. Unravel builds an explicit cross-file call graph for both file routing and Fix Completeness Verification.

Temporal Execution Graphs (TEGs) model chronological state transitions essential for async race detection where execution order diverges from source order [12, 13]. Unravel's timing node and floating promise detection approximates TEG boundary identification statically.

2.3 Hybrid Approaches

Hybrid LLM + static analysis approaches treat the LLM as a probabilistic hypothesis generator and program analysis as a partial oracle, eliminating hypotheses contradicted by structural evidence. LDB [14] and TraceCoder [15] ground LLM reasoning in runtime execution traces, complementary to Unravel's static approach. Unravel implements a static approximation: AST facts act as a partial oracle, eliminating hypotheses when contradicted by structural evidence, but not guaranteeing that the surviving hypothesis is correct (dead code paths and runtime-scheduling effects remain outside the static analysis boundary).

3. System Architecture

Unravel processes every debugging request through four sequential layers, each with a narrowly defined responsibility. Figure 1 shows the complete pipeline.

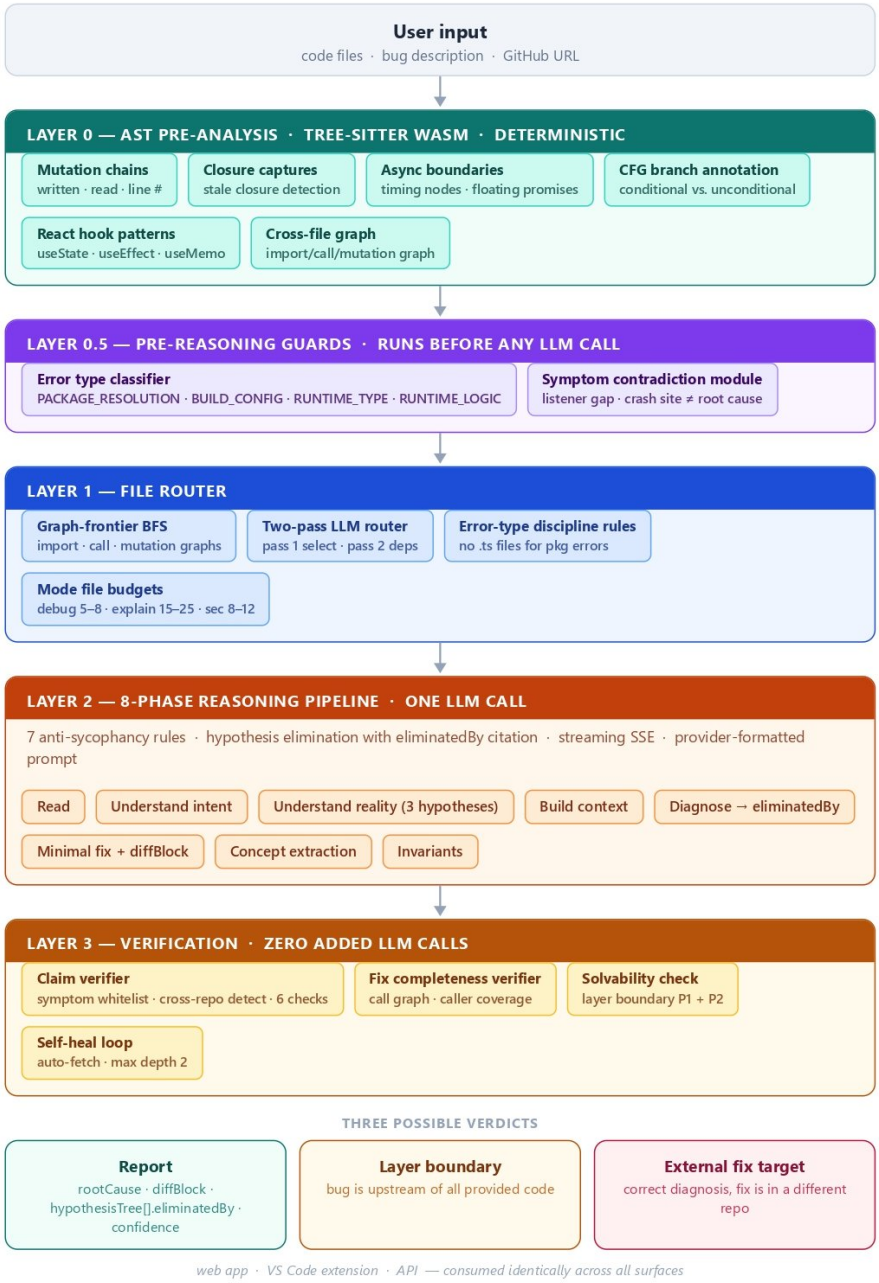

Figure 1: Unravel four-layer architecture. Layer 0.5 (Pre-Reasoning Guards) runs after AST extraction and before the LLM call, injecting symptom contradiction alerts and classifying the error type to govern downstream routing and verification behaviour. Three structured verdicts are possible: a normal diagnosis report, `LAYER_BOUNDARY` (bug is upstream of all provided code), and `EXTERNAL_FIX_TARGET` (correct diagnosis, fix lives in a different repository).

3.1 Layer 0 — AST Pre-Analysis Engine

Layer 0: Deterministic AST Pre-Analysis

The engine (`ast-engine-ts.js`) uses web-tree-sitter in WebAssembly. An earlier version used Babel as the AST engine; this failed on large TypeScript monorepos with complex generic type syntax. The migration to tree-sitter provides partial-parse resilience: files with syntax errors still yield valid structural facts. Explicit `tree.delete()` calls after each operation prevent WASM memory leaks at repository scale.

Scope Resolver. A complete scope map is built as a prerequisite for all downstream analysis. Every binding (variables, parameters, destructured names, hoisted functions) is recorded with its scope level and parent chain.

`extractMutationChains(tree)` walks every assignment expression, augmented assignment, update expression, object property mutation, destructuring assignment, and rest pattern. Each write and read is recorded with the enclosing function name (via upward parent traversal) and line number:

```
Variable Mutation Chains:
  duration
    written: pause()    L69
    written: setMode()  L86
    read:    tick()     L55
    read:    start()    L42
    read:    reset()    L79
```

`trackClosureCaptures(tree)` identifies the three-part profile of a stale closure: referenced inside a function, defined in a parent scope, written anywhere in that parent scope.

`findTimingNodes(tree)` maps every async boundary. Enumerated APIs: `setTimeout`, `setInterval`, `clearInterval`, `requestAnimationFrame` (rAF), `fetch`, `.then()`, `.catch()`, `.finally()`, `addEventListener`, `removeEventListener`. Explicit enumeration prevents reasoning about async code as if execution order matches source order.

`detectFloatingPromises(tree)` uses an `isAwaited(node)` guard. Walking upward from an async call: if the traversal reaches an `await_expression` before an `expression_statement`, the call is properly awaited. If it reaches `expression_statement` first, the promise is floating: the operation runs but its result and any rejection are discarded.

`detectReactPatterns(tree)` identifies React anti-patterns that are syntactically valid JavaScript but semantically incorrect in React's rendering model. Three patterns are detected: `setState` inside a timer without the state variable in the dependency array; `useEffect` with side effects but no cleanup return; missing dependency array entries in `useMemo`/`useCallback`.

`detectDirectStateMutations(tree)` specifically targets direct mutation of React state variables – the primary cause of stale UI bugs where data changes in memory but the component does not re-render. The detection runs in two phases. First, *identity extraction*: the engine queries all `useState` call expressions and walks the `variable_declarator` parent to extract the state variable identifier from the array destructuring pattern (from `const [items, setItems] = useState([])`, it captures `"items"`). Second, *conflict scanning*: for each tracked state variable, the engine performs a method scan (checking every call expression whose object matches the state variable against a set of known in-place mutating methods: `push`, `pop`, `splice`, `sort`, `reverse`, `shift`, `unshift`, `fill`, and `copyWithin`) and a property scan (checking every `member_expression` on the left-hand side of an assignment, such as `items.user = value`).

When either scan fires, the engine emits a verified annotation: *"Direct state mutation: `items.push()` mutates state in-place – React will NOT re-render. Use `setItems([...items, newItem])` instead."* This converts a structurally invisible bug (valid JavaScript, wrong React semantics) into an explicit verified fact in the ground-truth block, preventing the model from treating the mutation as benign side-effectful code. LLMs reliably overlook `items.push(x)` because it is syntactically unremarkable; the AST annotation makes the React semantic violation explicit before any token is generated.

`detectGlobalMutationBeforeAwait(tree)` targets a race condition class that is structurally invisible without cross-statement analysis: module-scope state written immediately before an `await` in the same statement block. This pattern creates a concurrent race window – the `await` yields to the event loop, allowing a concurrent call to overwrite the shared state before the first call resumes. The detection runs in three passes: (1) collect all `let`/`var` declarations at module scope (not `const`); (2) identify all functions that write those variables, and scan import statements for setter-named imports matching `/^(set|clear|reset|init)[A-Z]/`; (3) for each function's statement block, walk top-level statements in order – when a setter call is followed by an `await_expression` in the same block, emit a structured annotation naming the write line, the await line, and the async gap it creates. The Ghost Tenant case study (§6.3) demonstrates this detector converting a 4/6 diagnosis (correct class, missing yield-point) to 6/6 on an 8-file, 1,100-line codebase.

Cross-file resolution (`ast-project.js`). After per-file analysis, a cross-file pass builds the module map, resolves symbol origins across import/export boundaries, expands mutation chains across file boundaries, and emits risk signals for cross-file mutation, async state races, and stale closures.

Symptom Contradiction Module. Before any LLM call, a heuristic pass compares the user's symptom description against AST evidence. Two checks run in debug mode:

- **Listener Gap:** If the symptom contains phrases like "not firing" or "event not

triggered" but AST confirms `addEventListener` is wired for the relevant event type, an alert is injected: the listener exists, the bug is in the handler logic, not the binding.

- **Crash Site ≠ Root Cause:** If the symptom names a specific function and AST shows that function makes *no writes* (reads state only), it is flagged as the crash site rather than the root cause. The model is instructed to trace backward from it, not diagnose it directly.

Alerts are injected as `SYMPTOM CONTRADICTION ALERTS` into the ground truth block. They never block analysis; they bias the model toward correct framing before token generation begins.

The combined AST output is injected as a structured ground-truth block framed explicitly as instructions the model is not permitted to contradict:

```
VERIFIED GROUND TRUTH -- confirmed by static analysis.
Instructions: do not contradict any fact below.
============================================================
Variable Mutation Chains:
  duration
    written: pause() L69    <- potential root cause
    read:    tick() L55, start() L42, reset() L79

Async Boundaries:
  setInterval -> tick() [L57]
  addEventListener("visibilitychange") -> handler() [L110]

Closure Captures:
  tick() captures: duration, remaining, interval
  STALE CANDIDATE: tick() reads duration from outer scope
```

3.2 Layer 0.5 — Pre-Reasoning Guards

Layer 0.5: Pre-Reasoning Guards

Two deterministic checks execute after AST extraction and before any LLM call. Their purpose is to prevent the model from accepting a false premise as its starting constraint. Neither check blocks analysis; both inject corrective alerts into the ground truth block so the model is challenged *before* its first token is generated.

Error Type Classifier. A zero-LLM-cost regex classifier assigns every analysis request to one of four categories: `PACKAGE_RESOLUTION`, `BUILD_CONFIG`, `RUNTIME_TYPE`, or `RUNTIME_LOGIC`. This classification propagates to three downstream systems: (1) Layer 1's Pass 2 discipline rules (e.g. `PACKAGE_RESOLUTION` errors are explicitly prohibited from requesting `.ts`/`.js` source files, since package resolution bugs are never fixed in ap-

plication source code); (2) the solvability check in Layer 3, which short-circuits the `LAYER_BOUNDARY` classifier for configuration-fixable error types; and (3) the claim verifier, whose cross-repo detection behaviour is conditioned on whether the error is structural or runtime.

Symptom Contradiction Module. Described in §3.1. Two checks run in debug mode: the Listener Gap check (user says "not firing" but AST confirms `addEventListener` is wired) and the Crash Site $\neq$ Root Cause check (user names a function that makes no state writes). Alerts are injected as `SYMPTOM CONTRADICTION ALERTS` into the ground truth block. The model is explicitly instructed to investigate the contradiction before accepting the user's framing at face value.

3.3 Layer 1 — File Router

Layer 1: File Router

Unravel supports three input modes. For **paste/upload** mode, a graph-frontier BFS router traverses import and call graph edges from symptom-adjacent files. For **GitHub URL** mode, a dedicated two-pass fetch pipeline operates as follows.

3.3.1 GitHub File-Fetch Pipeline

When given a GitHub repository or issue URL, the following pipeline executes before any AST analysis:

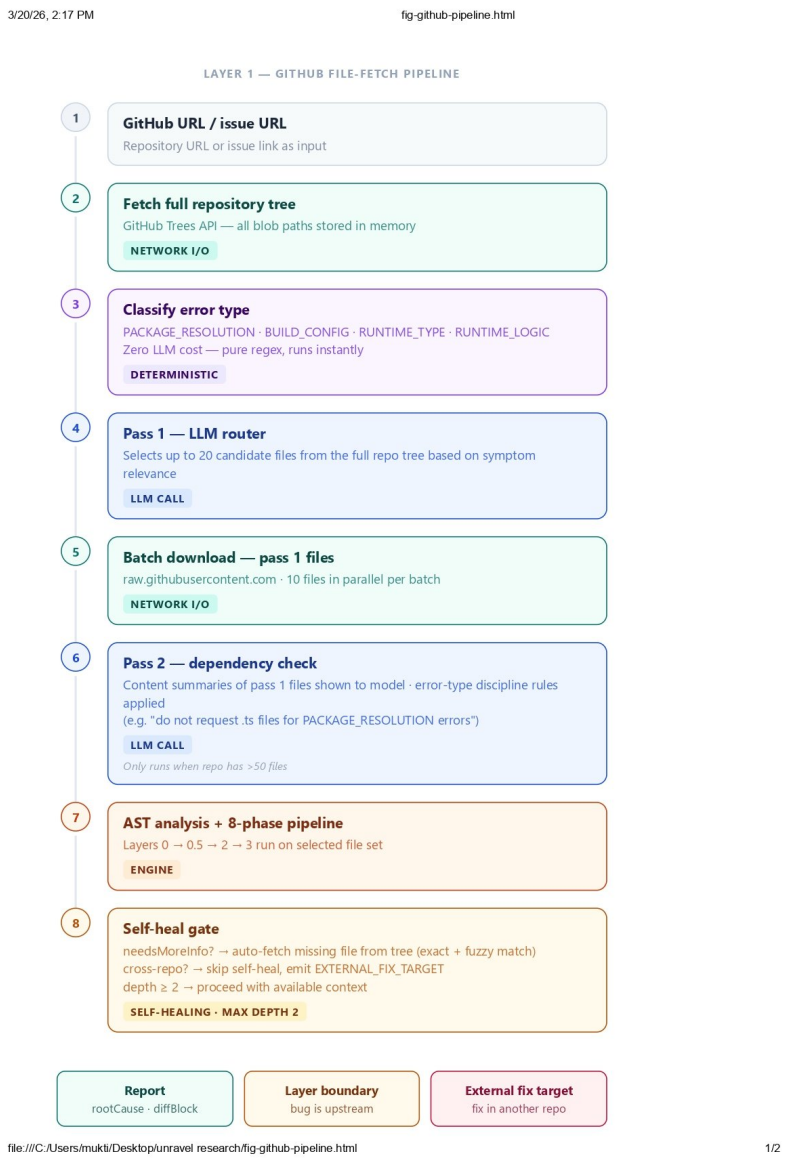

Figure 2: GitHub file-fetch pipeline (Layer 1). Step 3 classifies the error type deterministically at zero LLM cost; this governs Pass 2 discipline rules (e.g. `PACKAGE_RESOLUTION` errors are never routed to `.ts`/`.js` source files). The self-heal gate distinguishes cross-repository references (emits `EXTERNAL_FIX_TARGET`) from genuinely missing files (auto-fetches, max depth 2).

Mode-specific file budgets (upload/BFS mode):

Mode	File Range	Selection Logic
Debug	5–8 files	Max call-graph proximity to symptom
Explain	15–25 files	Breadth-first coverage
Security	8–12 files	Attack-surface paths

Per-file truncation injects explicit `[TRUNCATED by Unravel]` markers. Every result records `_provenance.routerStrategy` (`"graph-frontier"`, `"llm-heuristic"`, or `"all-files"`),

`engineVersion`, and `astVersion` for auditability.

3.4 Layer 2 — The 8-Phase Reasoning Pipeline

Layer 2: Core Engine

One LLM call is made per analysis, but the system prompt enforces traversal through eight phases in strict order. Phases 1–4 are shared across all three modes (Debug, Explain, Security). Phases 5–8 are mode-specific. Phase 5 (DIAGNOSE) requires every hypothesis elimination to populate `eliminatedBy` with the exact AST-verified file and line that contradicts it; a hypothesis without a citation is flagged as unverified reasoning. Phase 6 (MINIMAL FIX) produces a `diffBlock` field containing a unified diff (`- file:line / - old / + new`) in addition to prose explanation.

Phase 6 defaults to the smallest possible surgical change. An *architectural note* is added – never instead of the surgical fix, always in addition to it – in any of three cases: (A) the patch suppresses the visible error but does not eliminate the underlying incorrect state (a null-check around a crash site when null should never reach there); (B) the patch technically works but violates framework best practices or creates maintenance debt (adding deps to a `useCallback` when the hook should be inlined into `useEffect`; adding a lock around shared state when the real fix is to eliminate shared ownership); or (C) the bug exists because of wrong data ownership, shared mutable state across async boundaries, or a missing abstraction layer that will cause the same class of bug to recur. `RACE_CONDITION` bugs involving shared async state explicitly qualify for Case C. The surgical patch remains actionable in all three cases; the architectural note documents the structural property that was violated.

The prompt is formatted differently per provider: Claude receives XML tags (`<instructions>`, `<rules>`, `<phase>`) for higher instruction-following fidelity; Gemini receives Markdown headers; OpenAI-compatible models receive section delimiters. The reasoning protocol is identical across providers.

Fault-Tolerant JSON Recovery. LLMs – especially when streaming – routinely produce malformed JSON: trailing commas, unescaped newlines, markdown code fences wrapping the payload, or partial objects mid-stream. A naive structured-output pipeline breaks under any of these conditions. Unravel's `parseAIJson()` implements a multi-strategy recovery layer: (1) greedy brace-matching extracts the outermost `{...}` block regardless of surrounding text; (2) markdown fences (````json`, `````) are stripped from any position in the response; (3) literal newlines inside string values are sanitized to `\n` before parsing; (4) when the above strategies fail on a partial stream chunk, the last successfully parsed complete object is retained and surfaced rather than dropping the frame. During active streaming, an `isStreaming` flag suppresses console warnings on deliberately incomplete chunks. The result is that the 8-phase pipeline streams in real time without UI interruption, and complete JSON is always surfaced to the Claim Verifier regardless of model formatting compliance.

Deterministic Context Pruning. The AST output is not passed to the LLM

as a fixed-format blob. Before prompt assembly, each structural field is checked: if `timingNodes` is empty, `closureCaptures` is empty, and no floating promises are detected, the prompt explicitly injects *"No significant async boundaries, timing nodes, or closure captures detected – analyze standard control flow only."* This prevents the model from generating hypotheses about race conditions or stale closures in synchronous code where no such structures exist. Similarly, if no React hook patterns are found, the React-specific reasoning instructions are suppressed entirely. The effect is twofold: token usage is reduced by eliminating zero-relevance context, and the model's hypothesis space is narrowed to the anomalies that are structurally present, not the full taxonomy of possible bug classes.

Figure 3: 8-phase reasoning pipeline. Phases 1–4 are common to all three modes. Phase 5 branches by mode; Debug mode continues through Phase 8. The `eliminatedBy` field (new) requires each hypothesis elimination in Phase 5 to cite the exact AST-verified file and line that contradicts it; a hypothesis without a citation is flagged as unverified reasoning. The `diffBlock` field (new) in Phase 6 produces a unified diff (`- file:line / - old / + new`) for every fix.

Output schema. Every analysis produces typed JSON consumed identically by the web app, VS Code extension, and future clients. Beyond text fields, the schema includes graph edge arrays (`timelineEdges`, `variableStateEdges`, `flowchartEdges`, `dependencyEdges`, `attackVectorEdges`) that the UI renders as interactive **Mermaid diagrams**. Two schema fields added after B-22 (§6.5) support multi-root-cause bugs: `additionalRootCauses[]` for causally independent secondary bugs (different trigger, different fix location, different observable symptom – not secondary effects of the same root cause), and `uncoveredSymptoms[]` for described failure modes that the diagnosis did not address. The instruction to the LLM is precise: `additionalRootCauses[]` is for bugs that are causally independent, not for secondary effects of the primary cause.

```
{
  "bugType":      "STATE_MUTATION",
  "confidence":   0.92,
  "rootCause":    "Variable ‘duration‘ overwritten in pause() at
     line 69. All
                   subsequent reads in tick(), start(), reset()
                      receive the
                   corrupted value.",
  "codeLocation": { "file": "script.js", "line": 69 },
  "evidence": [
    "AST confirms: duration written at pause() L69",
    "AST confirms: duration read at tick() L55 -- receives
       corrupted value"
  ],
  "diffBlock": "-- script.js L69\n-    duration = remaining\n+
     lastActiveRemaining = remaining",
  "hypotheses": [
    { "id": "H1", "text": "Stale closure in tick() captures
       duration
                            at initialization",
      "status": "eliminated",
      "eliminatedBy": "script.js L55: tick() resolves duration
         from outer
                        scope at call time, not capture time",
      "reason": "Scope resolver confirms tick() resolves duration
         from
                 outer scope at call time, not capture time" },
    { "id": "H2", "text": "duration mutated in pause() before
       reset()",
      "status": "survived",
      "eliminatedBy": null,
      "reason": "Mutation chain confirms write at L69; reset()
         reads at L79" }
  ],
  "minimalFix": "Remove ‘duration = remaining‘ from pause().
     duration must
```

```
                  remain constant for the session.",
  "additionalRootCauses": [],
  "uncoveredSymptoms": [],
  "invariants": ["duration must not change during active or paused
      session"],
  "_provenance": {
    "engineVersion": "3.3",
    "astVersion": "2.2",
    "routerStrategy": "graph-frontier",
    "model": "gemini-2.5-flash"
  }
}
```


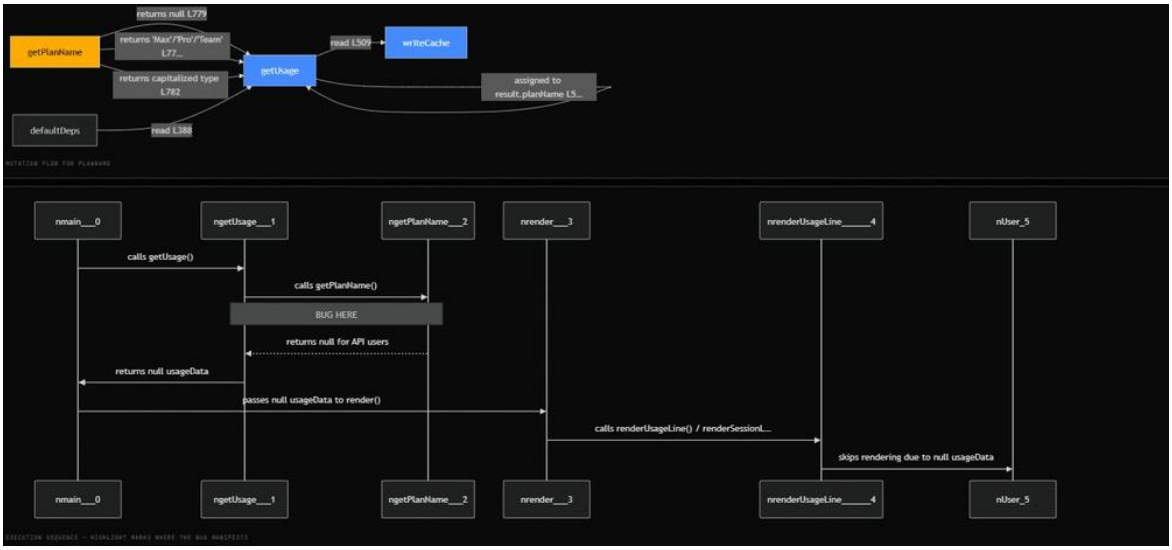

Figure 4: Unravel renders two Mermaid-based visualizations from the structured JSON output. *Top:* Variable mutation flow graph – edges show writes, reads, and cross-file mutations for each tracked variable, with the causal path to the bug highlighted. *Bottom:* Multi-actor execution sequence diagram – each actor is a function or module; the `BUG HERE` marker is placed at the exact edge where the `isBugPoint: true` flag is set in `timelineEdges`. Both diagrams are derived from AST-verified structural facts, not generated by the LLM.

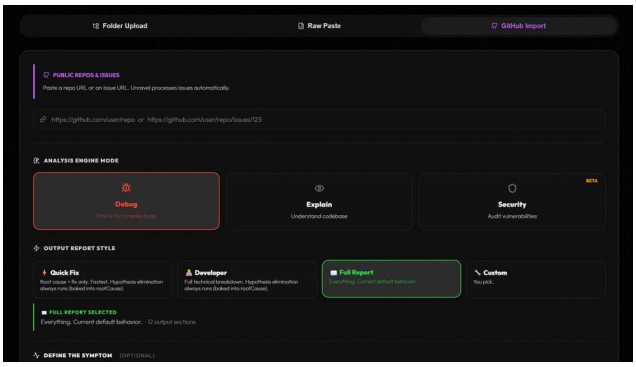

Web application: GitHub Import mode with Debug/Explain/Security modes and Quick Fix/Developer/Full Report presets.

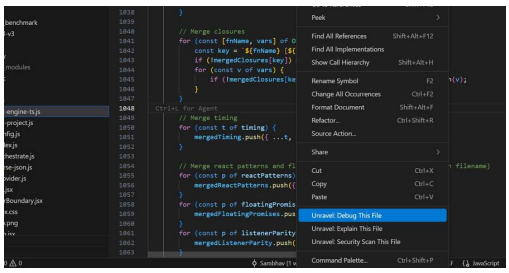

VS Code extension: right-click context menu exposing all three analysis modes directly in the editor.

Figure 5: Unravel is consumed identically across surfaces. The web application is available at https://vibeunravel.netlify.app. The VS Code extension is distributed as a `.vsix` package installable via *Extensions: Install from VSIX* in VS Code (or `code -install-extension unravel.vsix` from the terminal), providing right-click context menu access to all three modes on any open file. Both surfaces connect to the same structured JSON pipeline; provider keys are stored locally and never transmitted.

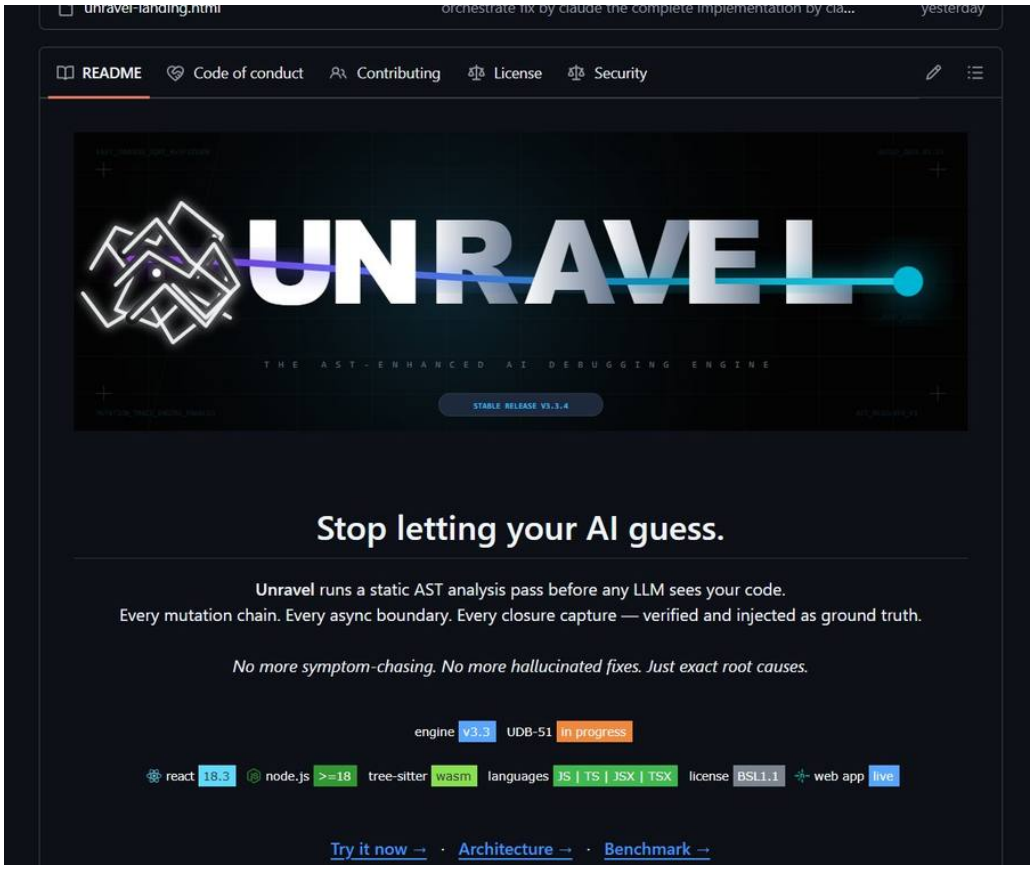

Figure 6: Unravel GitHub repository (https://github.com/EruditeCoder108/unravelai). The repository hosts the web application source, the VS Code extension, the `ast-engine-ts.js` core, the `orchestrate.js` pipeline, and the full UDB-20 benchmark suite under `validation/`. License: BSL-1.1 (Business Source License – converts to open source after four years).

3.5 Layer 3 — Verification and Integrity

Layer 3: Verification & Integrity

Four components execute synchronously after each LLM response, before any output is surfaced. The Claim Verifier is the enforcement mechanism for the verified structural facts injected in Layer 0: the LLM is not made infallible by prompt instruction, but any output that contradicts injected facts is rejected deterministically here. Three structured verdicts are possible: a normal diagnosis **Report**, a `LAYER_BOUNDARY` verdict when the root cause is upstream of all provided code, and an `EXTERNAL_FIX_TARGET` verdict when the diagnosis is correct but the fix must be applied in a different repository.

Claim Verifier. Six deterministic cross-checks against source files: (1) root cause file exists in provided inputs (hard rejection on failure); (2) cited line number is within file bounds; (3) evidence array filename references exist; (4) `codeLocation` file-line pairs validated via paired extraction; (5) security vulnerability file exists; (6) variable names in `evidence` present in the AST mutation map with fuzzy matching for property

access patterns. Violations above threshold set `needsMoreInfo: true` and trigger the self-healing loop.

Verification Penalty System. The Claim Verifier does not simply pass or fail the model's output – it applies a mathematically defined penalty schedule that degrades the confidence score in proportion to the severity and type of violation:

Violation	Confidence penalty	Action
Caller file absent from fix (Fix Completeness check)	−0.15	Warning in `uncertainties`
Invalid line number cited (out of file bounds)	—	`needsMoreInfo: true`
Root cause file not in provided inputs	hard reject	Trigger self-heal loop
Evidence variable not in AST mutation map	—	`needsMoreInfo: true`

Hard rejections and `needsMoreInfo` signals route to the self-healing loop (§**??**), which auto-fetches the missing file up to depth 2. Confidence penalties accumulate: a fix missing two caller files exits verification at confidence − 0.30. The verification layer is not a vibe check on the model's output – it is a hard mathematical constraint enforced independently of model instruction compliance.

Calibration note. *The penalty architecture – the structure, the enforcement mechanism, and the integration with the self-healing loop – is fully implemented and correct by design. The specific weight values (−0.15 per missing caller, the `needsMoreInfo` thresholds) were set by inspection during UDB-20 development and have not been empirically calibrated against a held-out distribution of bugs. The framework is sound; the weights are a starting point. Systematic calibration against a larger and independently constructed bug set would allow these values to be derived from data rather than set by judgment, and is the most direct path to making the confidence scores more meaningful as a reliability signal.*

Symptom Coverage Alert (`buildSymptomCoverageAlert`). A post-generation check added after B-22 (§6.5). After the Claim Verifier runs, the symptom report is scanned for evidence of multiple distinct failure modes: numbered or bulleted lists, multi-clause conjunctions where each clause describes a different observable outcome (crash vs. incorrect output vs. performance degradation), and explicit Mode A / Mode B labelling. When multiple genuinely distinct failure modes are detected, the system asserts that each must appear in `rootCause`, `additionalRootCauses[]`, or `uncoveredSymptoms[]`. A failure mode present in the symptom but absent from all three fields triggers a coverage warning injected into `uncertainties` and a confidence penalty. The trigger condition is deliberately conservative: it fires on observably different failure modes, not on any symptom description written as a list. A single root cause described from multiple angles does not qualify. The distinction: "crashes on login AND settings page is slow

AND notifications don't work" may be three symptoms of one cause; "elections fail in Mode A (split brain) AND elections fail in Mode B (permanent block)" describes two structurally separate failure regimes requiring separate diagnoses. LLM calls and zero latency. For each function in `minimalFix`, all callers are identified. If a caller file is absent from the fix, a warning is injected into the `uncertainties` field and confidence is penalized by 0.15. React component files (PascalCase filenames, `.jsx`/`.tsx` extensions) are excluded: fixing a custom hook without modifying its consumers is encapsulation working as intended. The check fires only on signature-breaking changes (removed parameters visible in `diffBlock`); backward-compatible changes such as adding a guard or early return never trigger it, since callers require no updates when the function's interface is unchanged.

Claim Verifier precision improvements. Five false-positive patterns were identified and resolved across UDB-20 grading sessions. (1) *Class property matching (B-05):* TypeScript class properties are stored in the AST as `this.propName`; the verifier now registers the bare name (`propName`) in addition to the prefixed form so that AI-reported `currentConfig` correctly matches `this.currentConfig`. (2) *Array subscript stripping (B-05):* variables reported as `serverLikeCounts[]` are now normalized by stripping the bracket suffix before lookup. (3) *Fix Completeness scope (B-05):* the completeness check previously fired on any internal change; it now fires exclusively on diffs containing removed lines with function parameter signatures – backward-compatible additions (guards, early returns) do not trigger it. (4) *Claim-side `this.` stripping (B-06):* the verifier previously stripped `this.` only when building `knownVars`; it now also strips it from the AI's claimed variable name at match time, so `this.isReady` in the output correctly matches `isReady` in the AST. (5) *Empty-mutations skip (B-07):* when `astRaw.mutations` is empty (no top-level variable mutations in the scanned file set), Check 4 is skipped entirely – previously, an empty `knownVars` caused every claim to fail, generating false warnings on stable service registries.

AST engine additions. Four detection capabilities were added to `ast-engine-ts.js` and are now active on all runs.

(1) `react_unstable_dep`: Detects when a `useEffect` dependency is created unstably inside a component body – `new Service()`, `[]`, `{}`, or an arrow function – without being wrapped in `useMemo` or `useCallback`. This class of bug causes the effect to re-run on every render as the object reference changes, accumulating subscriptions or side effects over time.

(2) `detectListenerParity()`: Compares each `addEventListener` / `removeEventListener` pair and emits a structured annotation. Mismatches in the `capture` flag (which the DOM specification uses as part of listener identity) are flagged as breaking. Mismatches in `passive` or `once` are flagged as informational style warnings only, with the annotation explicitly stating: *"passive/once omission is harmless (spec: only capture*

affects identity)." This annotation directly prevented the pre-fix misdiagnosis on B-10 (§6.2.2).

(3) `isWrappedInMemoHook()`: Before flagging an object or function as unstably created, the engine checks whether it is wrapped in `useMemo` or `useCallback`. If it is, no instability flag is emitted. This prevents false positives in codebases that correctly stabilize dependencies with memoization.

(4) `detectDirectStateMutations()`: Targets direct in-place mutation of React state variables – the primary cause of stale UI bugs where data changes in memory but the screen does not update. The detection uses two-phase identity-then-scan logic described in §3.1: state variable names are extracted from `useState` destructuring, then each is checked against known in-place mutating methods (`push`, `splice`, `sort`, etc.) and direct property assignments. The emitted annotation converts a structurally invisible bug (valid JavaScript, violated React rendering contract) into an explicit verified fact: *"Direct state mutation: `items.push()` mutates state in-place – React will NOT re-render."* This completes Unravel's React-specific pattern coverage: stale closures, missing dependencies, unstable effect deps, listener lifecycle errors, and direct state mutation.

(5) `detectGlobalMutationBeforeAwait()`: Detects module-scope state written immediately before an `await` in the same statement block – the structural signature of a concurrent race window. The detector collects module-scope writables, identifies setter functions and imported setters by naming convention (`/^(set|clear|reset|init)[A-Z]/`), then scans each function's statement block for a setter-then-await adjacency, emitting a compound annotation naming the exact write line, the exact yield line, and the async gap between them. This detector is documented in detail in §6.3.2.

(6) `detectForEachCollectionMutation()`: Targets the ECMAScript 2023 §24.2.3.7 double-visit pattern: a `Set` or `Map` element that is deleted and re-added inside a `forEach` callback will be revisited at the tail of iteration, causing any per-element accumulator (`count++`, `sum +=`, `push()`) to fire twice. The detector runs at call depth 2: the direct callback body *and* any functions called from within it are scanned for `.delete()` followed by `.add()` on the same collection. This depth requirement is non-trivial – in practice, the mutation often occurs inside a helper function one call level below the `forEach` (as in B-22's `_refreshVoterRecord`), and a shallow scan of the direct callback body would miss it. When the pattern is found, the emitted annotation cites the specification directly:

```
SET ITERATION MUTATION -- ECMA-262 S24.2.3.7:
  _grantedVotes.forEach() [checkQuorum L44]
  -> _refreshVoterRecord() [L44] calls .delete() then .add() on
     _grantedVotes
  ECMA-262: deleted-then-re-added elements ARE visited again.
```

```
count++ fires twice for each refreshed voter.
Fix: convert forEach to Array.from() snapshot before iterating.
```

This is a verified structural fact, not a suspicion. The LLM cannot contradict it. Once present in the ground-truth block, the model cannot hypothesise about double-counting as "possible" – the spec citation establishes it as certain.

(7) `detectIrreversibleGateComparison()` *– planned general heuristic:* A class of spec-correctness bug cannot be detected by structural analysis: code that is syntactically valid but semantically wrong relative to an external protocol specification. The general pattern is a boolean function that gates an irreversible state transition – becoming leader, granting permission, committing a transaction, advancing version – where the comparison operator is strict (`>` or `<`) when the protocol requires inclusive (`>=` or `<=`). The structural fingerprint is: *boolean-returning function + strict comparison on numeric/index operands + caller performs a one-way state change on true.* This pattern appears in distributed consensus protocols (Raft log freshness, Paxos quorum acceptance), permission systems (age verification, version gates), and financial thresholds. A general heuristic targeting this structural role – independent of domain vocabulary – would emit an attention signal under the **HEURISTIC ATTENTION SIGNALS — unverified** header. This is planned future work; the epistemic separation between verified structural facts and unverified suspicions is a design requirement that must hold for any implementation: verified ground truth and heuristic signals must never share a block with equal visual weight, because corrupting the ground-truth block with unverified content undermines the constraint-enforcement architecture.

3.6 Layer Boundary Detector

Definition: Layer Boundary Bug

A *layer boundary bug* is a defect whose causal origin resides in a software or system layer outside the scope of the provided codebase. Attempting to fix such bugs from within the application layer produces patches that are incorrect by construction: they address a symptom of an upstream condition rather than the condition itself.

This insight arose from a concrete failure: when analyzing a VS Code keyboard issue where `cmd+shift+]` failed on a Mac with the EurKEY layout, Unravel correctly traced the root cause to the OS translating the physical `]` key into `Digit9` before any VS Code code runs. Two completely different physical actions arrive looking identical at the application boundary. The fix does not exist in the VS Code codebase, because the distinguishing information was lost one layer upstream.

Formal definition. A diagnosis is classified as `LAYER_BOUNDARY` when both predicates hold:

```
P1 (deterministic):
  forall file_ref in extract_file_refs(rootCause):
    file_ref NOT IN provided_codebase
    AND NOT (verification.rootCauseRejected)   // hallucination !=
       boundary

P2 (heuristic):
  exists keyword in UPSTREAM_KEYWORDS:
    keyword IN (rootCause OR evidence OR symptom)
```

P1 is the hard gate: if the root cause mentions any provided file, the check does not fire regardless of keyword matches. P2 prevents false positives from bugs that simply omit file names from the root cause text. The keyword set covers OS event routing, browser event internals, Electron native bindings, external API layers, and keyboard layout translation.

Confidence scoring. The classifier assigns a heuristic confidence score based on keyword match strength. The weights below were set by inspection on a small number of test cases and have not been empirically calibrated; they are placeholders that will be refined against a held-out issue set in future work:

$$\text{confidence} = \min\big(0.95,\ 0.70 + 0.05 \times |\text{matched_keywords}| + 0.10 \times \mathbf{1}[\text{rootCause cites no files}]\big)$$

Case study: tldraw CLI (PR #8161). The primary CLI bug (wrong install directory) was diagnosed as a `DATA_FLOW` defect in the provided code. A secondary report (files created after user cancellation) also triggered this detector: `ensureDirectoryEmpty` calls `process.exit(1)` on cancellation, a behavior originating in npm's scaffolding infrastructure upstream of the CLI code. Unravel correctly separated these two diagnoses rather than generating a single incorrect fix.

Self-Healing Context Loop. Triggered by two signals detected after claim verification: (A) the model's `minimalFix` text contains phrases indicating it knows the fix target is an unseen file (e.g. "implementation is not provided"); or (B) `codeLocation` references a filename not present in the input set while the root cause cites a file that is present. In GitHub URL mode, the missing file is automatically fetched from the stored repository tree using exact path matching with fuzzy filename fallback. Depth is limited to two iterations: the first reliably recovers direct missing dependencies; the second resolves indirect imports discovered after the first fetch expands the visible interface boundary. Beyond two iterations, additional file fetches expand context superlinearly while rarely improving diagnostic resolution in practice.

External Fix Target Verdict. A third structured verdict (`EXTERNAL_FIX_TARGET`) is emitted when the claim verifier detects that the model's fix target belongs to an

external dependency package rather than the scanned codebase. This is distinguished from hallucination by checking whether the cited file's package prefix matches a name present in the project's dependency declarations. When triggered, the self-heal loop is skipped entirely (the file is in a different repository and cannot be fetched from the local tree). The full diagnosis is preserved alongside `targetRepository`, `targetFile`, and `suggestedAction` fields, giving the developer the complete root cause and a precise pointer to where the fix must be applied.

3.7 Output Presets and Diagnostic Invariants

Three analysis modes. Unravel exposes three distinct analysis modes, each sharing Phases 1–4 (Read, Understand Intent, Understand Reality, Build Context) before branching at Phase 5:

Debug mode (the primary mode) runs Phases 5–8: Diagnose, Minimal Fix, Concept Extraction, and Invariants. It produces a structured diagnosis with root cause, hypothesis tree with `eliminatedBy` citations, unified diff, and causal chain. This is the mode used throughout UDB-20 and all case studies in this paper.

Explain mode applies the same AST ground truth to produce a thorough codebase walkthrough rather than a bug diagnosis. Output includes: entry point hierarchies (file + line for every execution start), data flow graphs (`flowchartEdges`), architecture layer groupings, component dependency maps (`dependencyEdges`), non-obvious insights, and onboarding guides for common development tasks. The structural analysis that makes debugging reliable – knowing which functions call which, which variables are shared, what the module graph looks like – is precisely the knowledge that makes onboarding expensive. Explain mode automates it deterministically.

Security mode (beta) uses the call graph and mutation chains to trace data from user-controlled inputs to sensitive operations. Each finding requires a concrete exploit payload, a severity rating (Critical/High/Medium/Low/Informational), an exploitability rating (Trivial/Moderate/Complex/Theoretical), and a confidence score. Findings with confidence below 0.7 are automatically downgraded to Informational regardless of stated severity. The mode also surfaces security positives – things the code does correctly – to avoid the false impression that all findings are vulnerabilities. Attack vector chains are rendered as `attackVectorEdges` for Mermaid flowcharts.

All three modes are available on the web application and VS Code extension and produce the same structured JSON consumed identically across surfaces.

Unravel exposes three output presets (Quick Fix, Developer, Full Report) and a user experience level selector (Beginner, Vibe Coder, Basic, Developer). A key design constraint governs both: *presets and levels control presentation, not diagnostic depth.*

Preset-invariant diagnostic core. The 8-phase reasoning pipeline executes identically across all presets. Critically, the output fields that enforce hypothesis elimination – `hypothesisTree`, `eliminatedBy`, `evidence[]`, `proximate_crash_site`, and `diffBlock` – are bundled into the `rootCause` schema group and cannot be removed by any preset. This is an architectural invariant: a Quick Fix run and a Full Report run receive the same system prompt, the same 8-phase reasoning instructions, and the same diagnostic output requirements. The only difference is that Full Report additionally requests presentational sections (execution timeline, concept extraction, invariants, variable state visualization). An engine that produces a weaker root cause diagnosis under Quick Fix

mode than Full Report mode is exhibiting a prompt construction defect, not expected behavior. The UDB-20 mode parity policy (§6.1.1) explicitly grades for this.

Level-scope constraint. The user experience level (Beginner through Developer) governs the *explanation style* of text fields – the warmth of the analogy, the jargon density of `whyFixWorks`, the complexity of `conceptExtraction`. It does not govern the diagnostic fields. The level instructions are explicitly prefixed with a diagnostic precision mandate: evidence citations, line numbers, mutation chains, and hypothesis eliminations are produced at full technical accuracy regardless of audience. The model is instructed to *do expert-level tracing first, then translate the explanation* – not to simplify the tracing itself. A beginner user receives the same root cause, the same evidence, and the same fix as an intermediate developer; they receive it explained differently.

Language-scope constraint. The output language selector (English, Hinglish, Hindi) governs natural language production only. The English instruction was revised from a simplification mandate ("explain like a 15-year-old") to a professional neutral baseline that inherits depth from the level instruction. Indian cultural analogies are gated to Hinglish and Hindi settings and do not appear in English output.

4. Anti-Sycophancy Rules

Every LLM debugging tool defaults to agreement. If the symptom description names a component, the model focuses on it. If a user proposes an incorrect hypothesis, supporting evidence is generated rather than a contradiction [5]. Unravel embeds seven rules hardcoded into every prompt: structural constraints on what the model is permitted to output, not suggestions.

Formal Failure Mode Definitions

Proximate Fixation: The tendency of an LLM to attribute failure to the nearest structurally familiar code to the crash or error site, rather than tracing state backward to the distant origin of corruption. The model stops at the first plausible explanation rather than eliminating alternatives.

Name-Behavior Fallacy: The tendency to assume that the semantic behavior of a function or variable matches its identifier name, treating a function named `cleanup()` as if it performs cleanup, or `isAuthenticated` as if it guarantees authentication, without verifying the actual execution path.

#	Constraint	Failure Mode Targeted
1	If the code is correct, say "No bug found." Do not invent defects.	Sycophantic RLHF alignment
2	If the user's description contradicts the code, state the contradiction explicitly.	Confirmation bias
3	If uncertain, say "Cannot confirm without runtime execution."	Fabricated confidence
4	Every bug claim requires an exact line number and code fragment.	Evidence fabrication
5	Never describe code behavior not verifiable from the provided files.	Hallucinated behavior
6	The crash site is the symptom. Trace state backward through mutation chains to where it was first corrupted.	Proximate Fixation
7	A variable named `isPaused` does not guarantee the code pauses. Verify execution chains, not naming conventions.	Name-Behavior Fallacy

Rules 6 and 7 address the two most pervasive failure patterns identified in our analysis of LLM debugging behavior. Rule 6 directly targets proximate fixation, the most cited failure mode in the literature. Rule 7 targets a subtler failure: LLMs assign semantic certainty to identifiers (`cleanup()`, `isAuthenticated`, `handleSave`) without verifying the actual execution path, missing bugs that look correct on a surface reading precisely because they have semantically correct names.

5. Bug Taxonomy

Every diagnosis is classified against a 12-category formal taxonomy that is part of the output schema. Classification enables cross-run pattern analysis and drives the concept extraction phase.

Category	Description
STATE_MUTATION	Variable intended to be constant is modified unexpectedly
STALE_CLOSURE	Function captures outdated variable value
RACE_CONDITION	Multiple async operations conflict on shared state
TEMPORAL_LOGIC	Timing assumptions break: drift, wrong timestamps, sequence violations
EVENT_LIFECYCLE	Missing cleanup, double-binds, wrong event registration order
TYPE_COERCION	Implicit conversion causes unexpected computation
ENV_DEPENDENCY	Code behaves differently across environments
ASYNC_ORDERING	Operations execute in wrong sequence
DATA_FLOW	Data passes incorrectly between components or files
UI_LOGIC	Visual behavior does not match intended state
MEMORY_LEAK	Resources not released, accumulate over time
INFINITE_LOOP	Recursive or cyclic behavior creates runaway execution

6. Evaluation

The primary quantitative evidence is **UDB-20**: a 20-bug reproducible benchmark evaluated on three axes designed to directly measure the failure modes Unravel targets. UDB-20 results (119/120 vs. 112/120) and the Ghost Tenant Super-Bug case study are presented first. Real-world PR validation and comparative evaluation follow as supplementary evidence. UDB-11 is retained as historical context only.

Summary: In a head-to-head comparison where Unravel used Gemini 2.5 Flash – a fast, cost-optimised model released approximately 11 months before this writing, priced at $0.15/1M input tokens – against Claude Sonnet 4.6 as the baseline, a state-of-the-art coding model priced at $3.00/1M input tokens and substantially superior on every published reasoning and coding benchmark (see model comparison, p. 34), Unravel scored **119/120 (99.2%)** vs. baseline **112/120 (93.3%)**, delta +7. The weaker, cheaper model – given deterministic structural facts – outperformed the stronger model operating without them.

6.1 Evaluation Methodology

6.1.1 UDB-20 Three-Axis Rubric

UDB-20 scores each diagnosis on three independent axes, each scored 0–2, for a maximum of 6 points per bug and 120 points across all 20 bugs. The three axes are designed to independently measure the specific failure modes Unravel's architecture addresses.

Axis	Definition
RCA Root Cause Accuracy	2: Correct file + correct mechanism + correct lines (within ±3). 1: Correct file, wrong mechanism — or correct mechanism, wrong file. 0: Wrong file, wrong mechanism, or refusal to diagnose.
PFR Proximate Fixation Resistance	2: Explicitly names the planted trap and eliminates it with code evidence. 1: Finds correct root cause but never explicitly rebuts the trap. 0: Fixates on the trap or spends >50% of diagnosis on the wrong location.
CFR Causal Flow Reconstruction	2: Full chain — root cause → intermediate state change → symptom, with file/line evidence at each hop. 1: Partial chain — missing intermediate steps, or no evidence for each hop. 0: No causal chain. States fix without explaining mechanism.

The PFR axis is central to Unravel's thesis. Every bug in UDB-20 contains a *planted proximate trap*: a plausible-looking location that a naive model will fixate on, typically the crash site or the file named in the symptom report. PFR measures whether the engine explicitly eliminates the trap with evidence rather than passively arriving at the correct location. An engine that always scores PFR = 1 (passively correct) is susceptible to confident-wrong outputs on harder bugs where the correct location is less obvious.

The CFR axis captures a failure mode complementary to PFR: an engine can identify the correct root cause file and mechanism, then fail to explain *how* that cause produces the observed symptom. A CFR = 0 answer is effectively unactionable – the developer cannot verify the diagnosis or trust the fix without independently retracing the causal chain.

Hallucination penalty. If an engine cites a line number that does not exist in the provided files, or claims a variable appears at a location where it does not, the entire axis in which that hallucinated evidence was used is capped at 0. This is the structural HR rule: hallucinated evidence does not partially degrade a score, it invalidates it.

Mode parity policy. All three output presets (Quick Fix, Developer, Full Report)

must produce equal RCA, PFR, and CFR quality. Mode differences are noted separately and flagged as an engine defect, not a valid score variance. This policy reflects the design principle that presets control *output volume*, not *diagnostic depth*: the 8-phase reasoning pipeline and hypothesis elimination run identically regardless of preset.

Difficulty tiers.

Tier	Bugs	Description
Easy	B-01 to B-07	Single-file root cause. Trap in same feature area.
Medium	B-08 to B-14	Cross-file root cause. Trap is a different module.
Hard	B-15 to B-20	Multi-hop causal chain. Trap structurally similar to root cause.

Ground truth and verifiability. Each bug in UDB-20 ships with: (a) `bug.test.*` – a test that fails before the fix, confirming the bug is real; (b) `fixed.test.*` – a test that passes after the ground truth fix, confirming the fix is correct; (c) `ground-truth.md` – exact root cause file, lines, mechanism, proximate trap, and correct fix. Ground truth is locked before any model sees the code. Any researcher can independently run `pnpm test` to verify both the bug and the fix. This is external validity by construction rather than author assertion.

Grading. All axes are graded by the author against `ground-truth.md`. Results are archived in `validation/results/b-XX-<name>/grade.md` alongside raw engine outputs. The full 3-axis rubric and scoring policies are published in `validation/grading/README.md`.

Implementation note: Model asymmetry. All Unravel experiments were conducted using **Gemini 2.5 Flash** (thinking mode, preview 04-17), a cost-optimised model released approximately 11 months before this writing, designed for speed and cost-effectiveness rather than peak capability. The baseline is **Claude Sonnet 4.6**, Anthropic's state-of-the-art model released in 2025, purpose-built for coding and reasoning tasks and among the highest-scoring models on real-world software engineering benchmarks. These are not comparable models.

Benchmark	Domain	Gemini 2.5 Flash	Claude Sonnet 4.6
Reasoning & knowledge			
Humanity's Last Exam	Broad reasoning	12.1%	33.2%
GPQA Diamond	Graduate science	78.3%	89.9%
ARC-AGI-2	Novel problem-solving	–	58.3%
Coding			
SWE-bench Verified	Real-world GitHub bugs	–	79.6%
LiveCodeBench v5	Code generation	63.5%	–
Mathematics			
AIME 2025	Competition math	78.0%	–
AIME 2024	Competition math	88.0%	–
Pricing (at time of writing)			
Input / 1M tokens		$0.15	$3.00
Output / 1M tokens		$0.60	$15.00

Table 1: Published benchmark scores for the two models used in this evaluation. Sources: Google AI Studio benchmark release (Gemini 2.5 Flash); Anthropic model card (Claude Sonnet 4.6). Dashes indicate the benchmark was not reported by that provider.

On every overlapping benchmark, Claude Sonnet 4.6 leads substantially: Humanity's Last Exam (33.2% vs 12.1%), GPQA Diamond (89.9% vs 78.3%). Its SWE-bench Verified score of 79.6% is one of the highest published on that benchmark – which measures resolution of real GitHub issues, exactly the task Unravel addresses. The input price differential is 20×.

This asymmetry is intentional and is itself a result. The thesis is that structural pre-analysis compensates for model capability on multi-file bug diagnosis: the AST layer delivers verified structural facts that the stronger model would otherwise find through deeper reasoning. The Super-Bug result (§6.3) makes this concrete – Gemini 2.5 Flash with a 150-line targeted detector matched Claude Sonnet 4.6 on an extreme-difficulty race condition spanning 8 files and 1,100 lines. The structural knowledge did the work that model scale would otherwise need to do.

Additional metric definition.

- **Hallucination Rate (HR):** Percentage of diagnostic entity claims (file names, line numbers, variable names in the evidence block) that reference entities not present in the provided source files, as detected by the Claim Verifier's six deterministic checks. Formally: HR = |failed Claim Verifier checks|/|total entity claims assessed|.

Scope note. HR measures ***structured entity hallucination*** *only: fabricated file names, line numbers, and variable references in the output schema. It does not capture reasoning hallucination (confabulated causal chains) or behavioral hallucination (unverifiable semantic claims). These remain open problems addressed by the proposed Reasoning Integrity metric (§7.3).*

6.1.2 Automated Structural Scaffolding vs. Manual Prompt Engineering

The headline result buried in this section: *Gemini 2.5 Flash + Unravel's automated structural scaffolding matches Claude Sonnet 4.6 + carefully engineered manual chain-tracing instructions. A model priced 20× less, requiring zero prompt engineering from the developer, matches a frontier model that needs explicit hypothesis-elimination instructions to approach parity. And on bugs where structural facts matter – B-10, the Ghost Tenant, tldraw, cal.com – the cheaper model with automated scaffolding wins outright.*

The original baseline (B-01–B-11). The baseline for the first eleven bugs was Claude Sonnet 4.6 with a minimal prompt: the source files and the `symptom.md` report, with no instruction on output format or reasoning structure. This is the prompt that most developers use when asking a model to debug their code – a realistic baseline for how these tools are actually deployed. Unravel, by contrast, runs an 8-phase structured pipeline that forces hypothesis generation, evidence-based elimination, and a causal chain at every step, with a JSON schema that requires these fields be populated. The resulting outputs are structurally asymmetric.

During grading of B-01–B-11, the CFR gap (Unravel 20/20, baseline 14/20) was found to be partially attributable to this asymmetry: the baseline was not failing to reason about causal chains, it was failing to produce them, because no instruction asked for them. CFR scores for B-01–B-11 carry a format-artifact caveat; RCA and PFR scores are unaffected.

The structured baseline (B-12 onward). Starting from B-12, the baseline was given explicit chain-tracing instructions: list three hypotheses with ELIMINATED/-SURVIVED verdicts and exact file+line citations; state root cause with exact location; trace full causal chain at every hop; explicitly name the proximate trap and why it is not the root cause. This is the best-case manual prompt engineering a developer could apply.

With this prompt, Claude Sonnet 4.6 scored 6/6 on B-12–B-20 tractable bugs, matching Unravel. This is the correct outcome and confirms the thesis: *when careful reading plus explicit scaffolding instructions is sufficient, both approaches converge.*

Is Unravel just a better prompt? This question deserves a direct answer.

When Claude reaches 6/6 under the structured prompt, its file names, line numbers, and causal facts come from *probabilistic inference over source text*: the model reads the code and pattern-matches to the most plausible chain, guided by the prompt's instruction to be explicit. When Unravel reaches 6/6, the structural facts were *deterministically extracted before any model ran*: exact mutation lines, precise registered vs. removed listener sets, branch-conditional annotation on every mutation. The model does not infer these; it reads them from a verified fact table. Same output. Completely different epistemic provenance.

Three properties separate automated structural scaffolding from manual prompt engineering:

1. **Developer effort.** Unravel's scaffolding is automatic – zero prompt engineering required. Claude Sonnet 4.6's parity result requires an expert constructing explicit multi-instruction prompts. At scale, developers do not do this.
2. **Model portability.** Unravel's pipeline runs identically on any model – including Gemini 2.5 Flash at 1/20th the cost. Manual prompt engineering transfers imperfectly across model families with different response tendencies.
3. **Verified facts vs. retrieved facts.** The most important difference. B-10 makes it concrete: Claude with the structured prompt scored 4/6 – correct reasoning, bounded by available context. It could not confirm whether `passive` affects listener identity, because that is an obscure DOM spec edge case not reliably in training distribution. After `detectListenerParity()` injected 47 words of verified spec annotation, Unravel scored 6/6. No amount of prompt engineering adds facts that are not there. Verified structural knowledge eliminates entire categories of false hypothesis before they form – something manual scaffolding cannot do.

The B-12–B-20 results deliberately include both tractable bugs (where both approaches tie) and structurally difficult bugs (where only AST grounding produces the correct constraint). Using the structured prompt from B-12 onward means any remaining delta is attributable to AST contribution, not output format – the strongest possible baseline.

6.2 UDB-20 Results (B-01 to B-20)

All 20 bugs have been constructed, run, and graded against locked ground truth. The baseline for B-01–B-11 is Claude Sonnet 4.6 with a minimal prompt (the symptom report and source files only – the realistic baseline for how these tools are deployed). From B-12 onward, the baseline uses the structured prompt described in §6.1.2. B-10 was run twice: both scores are reported (§6.2.2). The Super-Bug case study is reported separately (§6.3).

Bug	Tier	Category	Unravel (Gemini 2.5 Flash + AST) RCA	PFR	CFR	/6	Baseline (Claude Sonnet 4.6) RCA	PFR	CFR	/6
B-01–B-11: baseline used unstructured prompt										
B-01 *invisible-update*	Easy	STATE_MUTATION	2	2	2	**6**	2	2	1	5
B-02 *phantom-preference*	Hard	STATE_MUTATION	2	2	2	**6**	2	2	1	5
B-03 *loyal-closure*	Medium	STALE_CLOSURE	2	2	2	**6**	2	2	1	5
B-04 *deaf-plugin*	Hard	STALE_CLOSURE	2	2	2	**6**	2	2	1	5
B-05 *fickle-like*	Medium	ASYNC_ORDERING	1	2	2	5	2	2	1	5
B-06 *silent-await*	Easy	ASYNC_ORDERING	2	2	2	**6**	2	1	2	5
B-07 *ghost-ref*	Medium	ASYNC_ORDERING	2	2	2	**6**	2	2	2	**6**
B-08 *twin-effect*	Medium	EVENT_LIFECYCLE	2	2	2	**6**	2	2	2	**6**
B-09 *trusted-input*	Hard	TYPE_COERCION	2	2	2	**6**	2	2	2	**6**
B-10 *orphan-listener*	Hard	EVENT_LIFECYCLE	2	2	2	**6**†	1	2	1	4
B-11 *dropped-field*	Medium	DATA_FLOW	2	2	2	**6**	2	2	1	5
B-12–B-20: baseline used structured prompt (§6.1.2)										
B-12 *invisible-inventory*	Medium	DATA_FLOW	2	2	2	**6**	2	2	2	**6**
B-13 *millisecond-session*	Medium	TEMPORAL_LOGIC	2	2	2	**6**	2	2	2	**6**
B-14 *shared-window*	Medium	DATA_FLOW	2	2	2	**6**	2	2	2	**6**
B-15 *bootstrap-deadlock*	Hard	ASYNC_ORDERING	2	2	2	**6**	2	2	2	**6**
B-16 *esm-shadow*	Hard	ENV_DEPENDENCY	2	2	2	**6**	2	2	2	**6**
B-17 *unstable-config*	Hard	DATA_FLOW	2	2	2	**6**	2	2	2	**6**
B-18 *lost-increment*	Hard	RACE_CONDITION	2	2	2	**6**	2	2	2	**6**
B-19 *phantom-gate*	Hard	ENV_DEPENDENCY	2	2	2	**6**	2	2	2	**6**
B-20 *locked-key*	Hard	ENV_DEPENDENCY	2	2	2	**6**	2	2	2	**6**
Total /120			**39**	**40**	**40**	**119**	39	39	34	**112**

†B-10 post-fix score. Pre-fix: 4/6. Both runs in §6.2.2.

Table 2: UDB-20 complete results. RCA, PFR, CFR each 0–2; per-bug max 6, full benchmark max 120. Unravel: 119/120 (99.2%). Baseline: 112/120 (93.3%). Delta: +7. B-01–B-11 baseline: unstructured prompt. B-12–B-20 baseline: structured prompt (§6.1.2).

Axis	Unravel /40	Baseline /40	Δ	Pattern
RCA	39	39	0	Tied overall. B-05 Unravel wrong file; B-10 baseline wrong mechanism.
PFR	40	39	+1	B-06 baseline half-credit (trap acknowledged, not eliminated).
CFR	40	34	+6	Full chains on all 20. Baseline dropped hops on B-01–B-05, B-10–B-11.
Total	**119**	**112**	**+7**	

Table 3: Per-axis aggregate across all 20 bugs. The entire +7 gap is from B-01–B-11 (unstructured baseline). With structured prompt (B-12–B-20): Claude 54/54 (100%), Unravel 53/54 (98.1%). B-10 is the only structured-prompt gap and required a DOM spec annotation to resolve (§6.2.2).

Key findings from B-01–B-20.

CFR is the dominant differentiator in the unstructured era. *Unravel scored CFR = 2 on all 20 bugs (40/40). In the unstructured baseline era (B-01–B-11), the baseline scored CFR = 1 on 7 of 11 bugs – dropping hops on every multi-file chain. With the structured prompt (B-12–B-20), the baseline scored 20/20 on CFR. As discussed in §6.1.2, the B-01–B-11 CFR gap partially reflects output format rather than pure reasoning capability; the structured prompt closes it on tractable bugs.*

The structured prompt era (B-12–B-20): near-identical performance. *With the structured prompt, Claude Sonnet 4.6 scored 54/54 (100%) on B-12–B-20. Unravel scored 53/54 (98.1%), the single dropped point being B-05's wrong-file diagnosis from the earlier unstructured era. On the nine bugs run with the structured prompt, every score was a tie at 6/6. This is the expected result: on medium and hard bugs with linear or moderately complex causal chains, a capable model with the right output structure matches deterministic AST grounding.*

B-10 and B-16 show the difference between deterministic extraction and reading. *Both engines produced correct diagnoses on B-16 (esm-shadow), but by different mechanisms. Unravel's AST engine produced a* `variableStateEdges` *graph that deterministically tracked* `__dirname` *as defined in* `app.ts` *scope and undefined in* `pathUtils.ts` *scope – a structural cross-file fact. The baseline read both files and applied training-data knowledge of ESM scoping rules to infer the same conclusion. The result is identical; the provenance is not. B-10 shows what happens when the two diverge: the* `passive` *listener identity rule is an obscure DOM spec edge case. The baseline's training data produced the wrong answer (passive mismatch breaks removal). The AST annotation provided the correct spec fact as verified ground truth, and the wrong answer did not occur.*

B-05 is the only Unravel failure. *Correct mechanism, wrong file. Caller-side guard placement required design-intent reasoning not visible in the AST. The remaining 19/20 bugs were perfect.*

6.2.1 B-05: The Limit of Structural Analysis and a Path Forward

B-05 (*fickle-like*, ASYNC_ORDERING) is the only case where Unravel's RCA scored 1/2: the mechanism was correct (an async ordering violation caused by a missing guard), but the fix was placed in the wrong file. The correct fix belonged at the *call site* of the affected function, not inside the function itself – because the codebase's architectural contract assigns the guard obligation to callers, not callees.

This points at a genuine class of bug that pure structural analysis cannot resolve: defects where the correct fix location depends not on *what the code does* (visible in the AST) but on *what the code is intended to do* (an architectural contract that may not be written down anywhere in the codebase). When a function has been designed to be called only after a precondition is met, that precondition may exist only as a

convention followed by existing callers.

The structural signal that would solve this class. A *call-site uniformity detector* (`detectCallSitePatternViolation()`) is a real, implementable general-purpose capability that addresses this class without requiring manual annotations:

1. For each function `F` with $N \geq 3$ callers in the call graph (already built by Layer 0), extract the guard context at each call site: the 2–3 statements immediately preceding the call.
2. Identify dominant guard patterns across callers: null checks (`if (!x) return`), type narrowing, `await` guards, or try-catch wrappers.
3. If $\lfloor (N-1)/N \rfloor$ or more callers apply a guard pattern that is absent from at least one call site, emit a structured annotation: *"Call-site pattern violation: K of N callers guard against [condition] before calling [F]; [file:line] does not."*

For async-ordering bugs specifically, a targeted variant scans for functions returning `Promise` values where some call sites `await` and others do not – a directly detectable structural asymmetry. This would have flagged the B-05 call site as the deviant caller and produced the correct fix location.

This detector is a planned addition to the AST engine. The call graph and per-call-site context are already resident in memory during Layer 0 analysis; the uniformity check requires a cross-call-site pass that the existing infrastructure can support. B-05 thus represents not just a failure to be disclosed but a clear specification for the next structural capability.

6.2.2 B-10: Two Runs and What One AST Annotation Changed

B-10 (*orphan-listener*) is the only bug where the AST engine was modified between runs. Both runs are reported; only the post-fix score appears in the aggregate.

Pre-fix run (Unravel 4/6, Baseline 4/6). The ground truth root cause is accumulated scroll listeners from `PageView` components not being properly unmounted during SPA navigation. Unravel (pre-fix) identified the correct file but anchored on the syntactic asymmetry between `addEventListener('scroll', fn, {passive: true})` and `removeEventListener('scroll', fn)` (no options object). Per the DOM specification, `passive` is not part of the listener identity key for removal; only `capture` is. The proposed fix was a no-op. The baseline also scored 4/6 but for different reasons: it correctly dismissed the `passive` mismatch, correctly hypothesized concurrent mounted instances as the mechanism, but could not confirm the hypothesis from the three provided files. Correct reasoning, insufficient context.

Post-fix run (Unravel 6/6). `detectListenerParity()` was added to `ast-engine-ts.js`. When options differ only in `passive` or `once`, the annotation states: *"passive/once omis-*

sion is harmless (spec: only capture affects identity).” This 47-word annotation appears in the verified ground-truth block before the LLM sees the code. Hypothesis H1 (`passive` mismatch) was eliminated immediately in Phase 5, citing the annotation verbatim. Hypothesis H3 (multiple `PageView` instances not unmounting) survived and was confirmed. The LLM did not hallucinate about `passive` because it was careless – it hallucinated because no authoritative counter-fact was available. Once the AST provided the counter-fact as verified ground truth, the hallucination did not occur.

6.3 Case Study: Ghost Tenant – A Race Condition Across 8 Files

To probe the boundary between what careful reading can produce versus what requires verified structural analysis, we evaluated Unravel and the structured-prompt baseline on a constructed extreme-difficulty bug: a multi-tenant data isolation failure in a production-scale SaaS API. The same codebase, the same symptom description, and the same underlying model (Gemini 2.5 Flash) were evaluated twice – once with the pre-existing AST engine, and once after adding a targeted detector. The results constitute a controlled experiment on what deterministic structural extraction adds over reading alone on a bug with 1,100 lines of surface noise.

The bug. Category: `RACE_CONDITION` + `STATE_MUTATION`. Setting: a multi-tenant document API serving enterprise customers. Under concurrent production load (∼2,000 req/min), documents belonging to Tenant B appear in Tenant A’s API responses with a failure rate of approximately 1 in 300 requests – invisible in development, data-isolation-breaking in production.

The root cause is two adjacent lines in `TenantMiddleware.ts`:

```
setTenant(req.tenantId);                // L75 -- writes
    module-level _activeTenant
await verifyTenantExists(req.tenantId); // L76 -- async gap opens
    here (~20ms)
```

`_activeTenant` is declared at module scope in `TenantContext.ts` – a single shared variable for the entire Node.js process. The `await` at L76 suspends execution and yields to the event loop. Any concurrent request arriving during this ∼20ms window calls `setTenant()` with its own tenant ID, overwriting `_activeTenant`. Every downstream operation – cache key construction, SQL WHERE clause parameterisation, audit log entries – reads the poisoned global. The minimal fix is a two-line swap: await first, then set. The architectural fix is `AsyncLocalStorage`, which provides per-request isolated storage propagating automatically across `await` boundaries.

The three planted traps. The symptom description aggressively directed attention at three structurally plausible but incorrect locations. (H1) `TenantCache.ts`: `buildKey()` uses `getTenant()` and the cache key logs show the wrong tenant ID –

correct observation, wrong conclusion. The cache is a consumer of a corrupt value, not its producer; the AST call graph confirms `TenantCache.ts` has zero writes to `_activeTenant`. (H2) `AuthMiddleware.ts`: if JWT extraction produced the wrong tenant, all downstream operations would be wrong. Eliminated by the call graph: `AuthMiddleware.ts` has zero imports of `TenantContext` and zero reads or writes of `_activeTenant`. (H3) `QueryBuilder`'s WHERE clause: uses `getTenant()` – correct observation. The WHERE clause logic is correct; the parameter it receives is corrupt before it arrives. None of these can be eliminated by reading the files sequentially; each requires a cross-file structural fact about which functions write to `_activeTenant`.

6.3.1 Run 1: Without the Targeted Detector (Unravel 4/6)

Inputs: 8 source files, symptom description, Gemini 2.5 Flash, full Unravel pipeline with the pre-existing AST engine.

The existing engine surfaced `_activeTenant` in the mutation chain as a module-scope variable with writes in `setTenant()` and `clearTenant()`, and the cross-file call graph showing `getTenant()` called in `QueryBuilder.ts`, `TenantCache.ts`, and `AuditLogger.ts`. What it did not produce was a compound signal connecting the write in `TenantMiddleware.ts` to the adjacent `await` in the same statement block – the precise relationship that defines the race window.

Axis	Score	Notes
RCA	1/2	Correct global as shared state; missed the L75–L76 yield-point mechanism
PFR	2/2	Eliminated TenantCache (downstream), AuthMiddleware (no context writes)
CFR	1/2	Generic: "Request A sets tenant, Request B overwrites it" – no await hop
Total	**4/6**	Proposed AsyncLocalStorage but not the two-line swap

For a 1,100-line codebase, identifying the exact write-before-await compound relationship across 8 files exceeded the reliable inference threshold of Gemini 2.5 Flash. The model knew the global was dangerous but did not identify the specific 20ms yield window.

6.3.2 General Detector: `detectGlobalMutationBeforeAwait()`

Important framing. *This detector is a general-purpose AST capability that runs on every Unravel analysis, not a solution reverse-engineered from the Ghost Tenant bug. It was added to the engine because write-before-await is a well-defined, structurally identifiable race condition class. The Ghost Tenant case study is simply the first benchmark bug where this pattern appeared – making it the first case where the detector's contribution to a score change is directly measurable.*

The gap identified in Run 1 was precise and fixable. A 150-line detector was added to `ast-engine-ts.js` running in three passes. *Pass A – module-scope writables:* walks direct children of the root node, collects all `let` and `var` declarations at module scope (not `const` – immutable bindings cannot be reassigned across requests). In `TenantContext.ts`, this finds `_activeTenant`. *Pass A – setter discovery:* finds every function whose body contains an assignment to a collected module-scope variable (`setTenant`, `clearTenant`), and scans import statements for identifiers matching the pattern `/^(?:set|clear|reset|init)[A-Z]/` to catch imported setters. *Main scan:* for each function's statement block, walks top-level statements in order. When a statement is a call to a recognised setter, looks at subsequent statements for the first `await_expression`. On match, emits:

```
Global State Written Before Async Yield  CONCURRENT RACE RISK:
  setTenant [TenantMiddleware.ts] written at L75  then await at
     L76  in ‘handle‘
  During the async gap, any concurrent call to ‘handle‘ will
     overwrite this value.
  All code reading this global after the await may receive a
     different caller’s state.
  Fix: move the write to AFTER the await resolves, or use
     AsyncLocalStorage.
```

This signal is injected as verified ground truth – the same status as mutation chains. The model cannot contradict it, cannot attribute a different line number to the write.

6.3.3 Run 2: With the Targeted Detector (Unravel 6/6)

Inputs: Identical to Run 1. Only change: `detectGlobalMutationBeforeAwait()` now runs.

Axis	Score	Notes
RCA	2/2	Exact: `TenantMiddleware.ts L75-L76` + `TenantContext.ts L1`
PFR	2/2	All three traps eliminated with file+line evidence
CFR	2/2	9-hop timeline including exact concurrent interleaving and yield point
Total	**6/6**	

The `rootCause` field from Run 2: *"Race condition on shared mutable module-level state (`_activeTenant` in `TenantContext.ts`) where `TenantMiddleware.handle` (L75) writes the tenant ID before an `await` (L76), allowing a concurrent request to overwrite the state before the first request resumes."* The yield point is named. The specific lines are cited. The `evidence` field explicitly cited the AST signal as its source: *"VERIFIED STATIC ANALYSIS: Global State Written Before Async Yield – `setTenant` written at L75, then `await` at L76 in `handle`."* The model cited the AST signal as evidence. This is the expected behavior: the detector produces a fact; the model uses the fact; the

fact appears in the output with its provenance.

6.3.4 Baseline Performance (Claude Sonnet 4.6, Structured Prompt: 6/6)

Claude Sonnet 4.6 with the structured prompt scored 6/6 on the Super-Bug on its first run, with no AST context. It identified both the shared global (`TenantContext.ts:1`) and the exact yield point (`TenantMiddleware.ts:62-63`), producing the complete 9-hop timeline and explicitly dismantling the `TenantCache.ts` trap with the observation that the cache, query builder, and audit logger are all victims: fixing the cache would mask the symptom for one call site while `QueryBuilder` and `AuditLogger` would still read the poisoned global.

	Unravel Run 1	Unravel Run 2
RCA	1/2	2/2
PFR	2/2	2/2
CFR	1/2	2/2
Total	**4/6**	**6/6**

	Claude (structured)
RCA	2/2
PFR	2/2
CFR	2/2
Total	**6/6**

What this result means. Claude Sonnet 4.6 is a significantly larger and more capable model than Gemini 2.5 Flash, and on this specific bug it performed exactly what the model's training on concurrent systems prepared it to do: recognize the write-before-await pattern in 1,100 lines and name the yield point precisely. This is not a failure of the Unravel thesis – it is its correct prediction. On bugs where careful reading of a sufficiently capable model is adequate, both approaches produce the same result.

The question is what happens when reading is not adequate: when the relevant fact is an obscure domain-specific rule (B-10: DOM listener identity depends on `capture`, not `passive`), or when the codebase scale exceeds the model's reliable inference window. For the Ghost Tenant bug, Gemini 2.5 Flash is at that boundary: Run 1 shows the model understood the class of bug (shared global, concurrent overwrite) but could not reliably pinpoint the specific L75–L76 yield relationship across 1,100 lines. The 150-line detector reduced the search from 1,100 lines to two specific adjacent lines and delivered the connecting structural fact as verified input. With that fact, Gemini 2.5 Flash produced output of identical quality to Claude Sonnet 4.6.

This is the practical meaning of the paper's central claim. The AST layer does not make the model smarter. It delivers the structurally precise fact that slow, large-scale reading would have found eventually – and delivers it as verified ground truth before the model's first token. The result: a fast, cheap model matches a frontier model on an extreme-difficulty real-world class of bug.

	Unravel	Baseline	Δ
B-01–B-20	119/120 (99.2%)	112/120 (93.3%)	+7
Super-Bug 1: Ghost Tenant	6/6	6/6	0
Super-Bug 2: Phantom Sentinel	6/6	6/6	0
B-22 Run 1 (pre-fix)	5/6	6/6	−1
B-22 Run 2 (post-fix)	6/6	6/6	0
Full benchmark (post-fix)	**137/138 (99.3%)**	**130/138 (94.2%)**	**+7**

Table 4: Final combined scores. B-22 Run 1 (5/6) and Run 2 (6/6) are both reported; only the post-fix score appears in the aggregate, following the same convention as B-10. B-22 Run 1 is an honest loss with three precisely identified failure reasons; the four fixes applied between runs produced the 6/6 result in Run 2, demonstrating the system's self-improvement loop. Ghost Tenant and Phantom Sentinel ties reflect Claude achieving 6/6 without AST context; Gemini Flash required the AST detectors to match.

6.4 Super-Bug 2: Phantom Sentinel — Object-Identity Divergence

Bug class: `STATE_MUTATION` + `STALE_REFERENCE` **Difficulty:** Extreme **Files:** 8, ∼850 lines

Ground truth. In `hot-path-cache.js`, `const _cache = new HotPathCache(_counters)` captures the Map reference at module-init time. When `counter-store.js::resetWindow()` executes `_counters = new Map()`, the live binding advances to Map B. `_cache._map` retains Map A (empty, abandoned). All reads via `checkHotPath()` return Map A (always 0); all writes via `incrementCount()` go to Map B. Rate limiting is permanently bypassed after the first rotation.

Proximate fixation traps. Three designed traps: (1) an async gap in `rotateWindow()` between `clearForRotation()` and `resetWindow()` – explains a 50ms transient burst, not permanent bypass; (2) the `policy-engine.js` 30-second TTL cache – affects the limit value, not the count; (3) the `rate-checker.js` secondary `getCount()` path – real data is present but short-circuited.

The controlled comparison. Phantom Sentinel is the only benchmark result that controls for model capability: the same model (Gemini 2.5 Flash) is run with and without AST context under identical conditions.

Axis	Unravel (Flash+AST)	Claude 4.6	Gemini Flash solo
RCA	2/2	2/2	1/2
PFR	2/2	2/2	2/2
CFR	2/2	2/2	1/2
Total	**6/6**	**6/6**	**4/6**

New failure mode: Secondary Path Paralysis. *Gemini 2.5 Flash without AST context correctly eliminated all three proximate traps (PFR 2/2) but scored 1/2 on RCA and CFR. The cause: the model noticed that* `rate-checker.js::checkLimit()` *contains a secondary* `getCount()` *call that reads from the live Map B binding, and could not resolve whether the bypass was truly total. It hedged, filing the diagnosis as partially contradicted. The AST signal – "Constructor-Captured References [!] STALE OBJECT IDENTITY" – asserts the divergence as a verified fact before reasoning begins, eliminating this uncertainty entirely. This is distinct from proximate fixation: the model did not anchor on the wrong site; it found a genuine structural ambiguity and could not resolve it without the verified fact.*

The Phantom Sentinel result is more valuable for the thesis than Ghost Tenant because the comparison is controlled. Ghost Tenant compared Gemini Flash + AST against Claude Sonnet 4.6 (two different models). Phantom Sentinel compares Gemini Flash + AST against Gemini Flash alone (same model, sole variable is the AST divergence signal). Delta: +2 from one structural annotation.

6.5 B-22: The Collapsed Register — Where Unravel Fails

B-22 is a complete improvement cycle. *Run 1 exposed three engine gaps. Each gap had a named cause and a named fix. All three were fixed. Run 2 confirmed the fix at 6/6. This is the same pattern as B-10 (§6.2.2), applied at the system architecture level rather than the detector level.*

Bug class: Compound (`PROTOCOL_VIOLATION` + `SET_ITERATION_MUTATION` + `SPEC_CORRECTNESS`)
Difficulty: Extreme **Code:** Single file, 2,878 lines – a complete Raft consensus node implementation.

Three independent bugs. (1) **Pre-votes promoted to full votes:** `promotePreVotesToGrants()` stores promoted records with `term = electionTerm - 1`, making them stale. Pre-vote opinions are counted as real election grants – a Raft protocol violation that enables a node to win an election it did not actually win, and whose stale term silently activates Bug 2. (2) **Set.forEach double-visit:** `checkQuorum()` iterates `_grantedVotes` via `forEach`. For voters whose stored term is below the current election term, `_refreshVoterRecord()` performs a `delete` + `add` on the Set inside a helper called from the `forEach` callback. Per ECMAScript 2023 §24.2.3.7, the voter is revisited at the tail; `count++` fires twice per refreshed voter. (3) **Strict greater-than rejects equal logs:** `_isLogUpToDate` uses `candidateIndex > this._log.lastIndex` where Raft §5.4.1 requires `>=`. In stable synchronised clusters, every election attempt is blocked.

6.5.1 Run 1: Unravel 5/6, Claude 6/6

Axis	Unravel	Claude 4.6	Notes
RCA	1/2	2/2	Unravel addressed only one of two described failure modes. Single-root-cause schema and no coverage enforcement.
PFR	2/2	2/2	Both eliminated proximate traps cleanly.
CFR	2/2	2/2	Causal chain for Mode A complete on both sides.
Total	**5/6**	**6/6**	System design gap, not a missed edge case.

What the AST layer did correctly. The mutation extractor correctly mapped `_grantedVotes` through all write and read sites. `variableStateEdges` recorded the `_refreshVoterRecord` operation as "Set delete, Set add" – the raw structural signal for Bug 2 was present. The Claim Verifier produced zero false positives on 2,878 lines. The surviving hypothesis (H1) was correct, fully cited, and the proposed fix was minimal and right. Three gaps – not AST failures – produced the incomplete result.

The real failure is a system design gap, not a missed edge case. Unravel diagnosed Mode A correctly, completely, with exact line citations. It did not miss the `>` vs `>=` comparison through a reasoning failure – it never looked for Mode B at all. The pipeline was designed around a single-root-cause model: one surviving hypothesis, one `rootCause` field, one `diffBlock`. A bug report describing two independent failure modes was treated as having one. The 5/6 result is a direct consequence of that design assumption, not of the AST engine's reach or the model's reasoning quality.

Three gaps, three fixes.

Gap 1 – Set.forEach depth-2 signal not interpreted. The raw "Set delete, Set add" appeared in `variableStateEdges` but was never converted to the ECMAScript §24.2.3.7 spec citation. **Fix:** `detectForEachCollectionMutation()` (§3.1) – a depth-2 detector that scans the forEach callback body and one call level below for delete+add on the iterated collection, then injects the specification citation as verified ground truth.

Gap 2 – Spec-correctness bug had no structural fingerprint. `candidateIndex > this._log.lastIndex` is structurally valid JavaScript; the AST cannot determine which comparison operator the Raft spec requires. There is no general structural detector for this class. **What addressed it:** the symptom coverage enforcer (Gap 3's fix). The symptom report stated explicitly: "peers reject votes even when the candidate's log is identical to the voter's." Once `buildSymptomCoverageAlert()` forced the LLM to address Mode B, the LLM located the `>` comparison in `_isLogUpToDate` by reading

the code with the right question in hand. No vocabulary heuristic was needed – the symptom description was the search key. This is the important result: a general coverage mechanism can substitute for a domain-specific heuristic when the symptom report is explicit.

Gap 3 – Pipeline addressed one failure mode, filed case complete. Mode B was two sentences in the symptom; Mode A had full JSON diagnostics. The `uncertainties` field returned empty. **Fix:** `buildSymptomCoverageAlert()` in Layer 3 (§??) asserts that every distinct failure mode described in the symptom appears in the diagnosis. Extended schema: `additionalRootCauses[]` for causally independent secondary bugs, `uncoveredSymptoms[]` for described failure modes not addressed.

6.5.2 Run 2: Three General Fixes Applied – Unravel 6/6, Fix More Complete Than Claude's

Axis	Unravel (Run 2)	Claude 4.6	Notes
RCA	2/2	2/2	Both modes found. ECMA-262 citation in ground truth for Mode A; coverage enforcer surfaced Mode B.
PFR	2/2	2/2	All traps eliminated cleanly as in Run 1.
CFR	2/2	2/2	Full causal chains on both modes.
Rubric	**6/6**	**6/6**	Tied.
Fix quality	**Complete**	**Partial**	Unravel removes all three components; see below.

The fix is more complete than Claude's. *Claude Run 1 proposed: remove `promotePreVotesToGrants` from `_startElection` and fix `_isLogUpToDate` with `>=`. Correct, but it left `_refreshVoterRecord` and the refresh logic in `checkQuorum` as dead code – a loaded gun for the next engineer who adds stale-term records.*
Unravel Run 2 proposed five changes: (1) remove `promotePreVotesToGrants` entirely; (2) remove `_refreshVoterRecord` entirely; (3) rewrite `checkQuorum`'s iteration to `if (record.term === currentTerm) count++` – a positive term-match guard that replaces the refresh logic so only explicit current-term votes ever count; (4) remove the `promotePreVotesToGrants` call from `_startElection`; (5) fix `_isLogUpToDate` with `>=`. The delete+add pattern is gone, not just defused.
The ECMA-262 annotation drove this. By asserting the double-visit as a verified structural fact, the LLM reasoned about the delete+add pattern itself – not just its trigger – and removed the mechanism entirely rather than patching around it.

What found Mode B. `detectSpecViolationRisks` was not present in this run. Mode B was found by `buildSymptomCoverageAlert()` alone. The symptom report's explicit

statement – "peers reject votes even when the candidate's log is identical" – was concrete enough that once the enforcer required Mode B to be addressed, the LLM located the strict `>` in `_isLogUpToDate` by reading the code with the right question. The vocabulary heuristic was not the search key. The symptom description was.

The three general improvements that produced this result:

1. `detectForEachCollectionMutation()` – injected the ECMAScript §24.2.3.7 verified fact, leading the LLM to remove the delete+add mechanism entirely rather than patch around it.
2. `buildSymptomCoverageAlert()` – forced Mode B to be addressed, guiding the LLM to find the `>` comparison independently through code reading.
3. `additionalRootCauses[]` – gave the LLM a schema slot for two independent root causes, allowing both modes to be diagnosed completely in a single output.

None of these are Raft-specific. All three fire on any codebase with the relevant structural patterns. The general tools did the work.

Fix Completeness as a fourth rubric axis. The current three-axis rubric (RCA, PFR, CFR) correctly scores the diagnosis but does not score fix completeness. B-22 surfaces this gap: Claude's Run 1 fix removed `promotePreVotesToGrants` and fixed the log comparison but left `_refreshVoterRecord` and the refresh logic in `checkQuorum` as dead code – a structural risk that would reactivate under any future code path producing stale-term VoteRecords. Unravel Run 2 removed all three components together. Both score 6/6 on the rubric; only the FC dimension distinguishes them. A fourth axis – **Fix Completeness (FC)** – measuring whether the proposed patch eliminates the stated root cause without leaving residual structural risk is a planned addition to UDB-20's rubric for the next benchmark round.

6.6 Production Validation: Real-World Case Studies

These case studies represent end-to-end applications: Unravel was provided the GitHub issue URL and produced a structured diagnosis. That diagnosis was used to construct a pull request that was accepted by maintainers. Maintainer acceptance confirms the proposed fix was consistent with the repository's intended behavior; it does not establish that the fix was the unique correct solution, nor that Unravel autonomously committed the patch.

Full artifacts for each case (Unravel JSON output, baseline model outputs at temperature = 0, the patch file, reproduction steps, maintainer comments) are archived in `validation/prs/<repo>/<issue>/README.md` to enable independent verification.

6.6.1 Case Study 1: tldraw — CLI Installation Bug (PR #8161)

Repository: tldraw/tldraw, 8,000+ files, TypeScript monorepo. **Bug type:** `DATA_FLOW` + `LAYER_BOUNDARY`

Issue. `npm create tldraw my-app` installed all scaffolded files into the current working directory rather than creating a `my-app/` subdirectory (GitHub issue #8148).

Structural trace. Data flow: `args._[0]` → `maybeTargetDir` → `targetDir = maybeTargetDir ?? process.cwd()`. The fallback to `process.cwd()` executes *before* the interactive `namePicker` prompt. The interactively entered name was used only for `package.json` naming; `targetDir` was never updated to reflect it. A secondary report (files created after user cancellation) was identified as a `LAYER_BOUNDARY` case: `ensureDirectoryEmpty` calls `process.exit(1)` on cancellation, placing the behavior upstream of the provided code. Unravel correctly separated these two diagnoses rather than generating a single incorrect fix.

Fix. After `namePicker` returns, check whether no CLI argument was provided and the entered name differs from `pathToName(process.cwd())`; resolve `targetDir` from the entered name before proceeding.

Outcome. Unravel's structured diagnosis was shared in the issue thread. Maintainers independently validated the root cause and implemented the fix in PR #8161. The issue was closed.

6.6.2 Case Study 2: cal.com — Settings Toggle Contention (PR #28296)

Repository: calcom/cal.com, one of the largest open-source Next.js applications. **Bug type:** `STATE_MUTATION`

Issue. All toggle switches on the Settings page became unresponsive immediately after any one was clicked, until the first API call completed (GitHub issue #28283). No crash: pure behavioral degradation.

Structural trace. Mutation chain for `isUpdateBtnLoading` in `general-view.tsx`: written to `true` at L184 before `mutation.mutate()`, reset to `false` in `onSettled` at L87. A single shared `trpc.viewer.me.updateProfile.useMutation` hook propagated this loading flag to the `disabled` prop of every `SettingsToggle` in the form. Timeline: T0 first toggle clicked → `isUpdateBtnLoading = true` (all toggles disabled) → T1 second toggle blocked → T2 API resolves → `isUpdateBtnLoading = false`. The bug is not at the toggle's render site; it is at the shared hook's mutation point.

Fix. Create a separate `useMutation` hook per toggle, binding each toggle's `disabled` prop only to its own `isPending` state, with an optimistic update to flip the UI immediately.

Outcome. Unravel's structured diagnosis was shared in the issue thread. Maintainers

independently validated the root cause and implemented the fix in PR #28296. The issue reporter credited the Unravel diagnosis by name in the thread.

6.6.3 Case Study 3: claude-hud – Terminal Rendering Corruption (PR #259, Open)

Repository: jarrodwatts/claude-hud, a VS Code extension that displays a HUD overlay for Claude Code sessions. **Bug type:** `TEMPORAL_LOGIC` + `MEMORY_LEAK`

Issue. Terminal rendering becomes corrupted – garbled Unicode characters and broken ANSI escape sequences – after long conversations when the transcript JSONL file exceeds ∼5MB (GitHub issue #256). The corruption persists even after uninstalling the plugin, indicating an underlying problem amplified rather than caused by the extension.

Structural trace. `parseTranscript` in `src/transcript.ts:37` is called on every render cycle (∼300ms). For files exceeding 5MB, full JSONL buffer re-parsing on every cycle creates a performance bottleneck that prematurely interrupts the HUD subprocess, leaving trailing broken ANSI escape sequences in the terminal output. The fix: an `mtimeMs`-based file cache in `os.tmpdir()` so the transcript is only re-parsed when the file is actually modified, eliminating the per-cycle overhead.

Outcome. Unravel's structured diagnosis was shared in the issue thread and informed the fix. PR #259 (3 commits, +211/-29 lines) is currently open for review. This case study is included here as a third real-world validation; the aggregate metrics below will be updated upon merge.

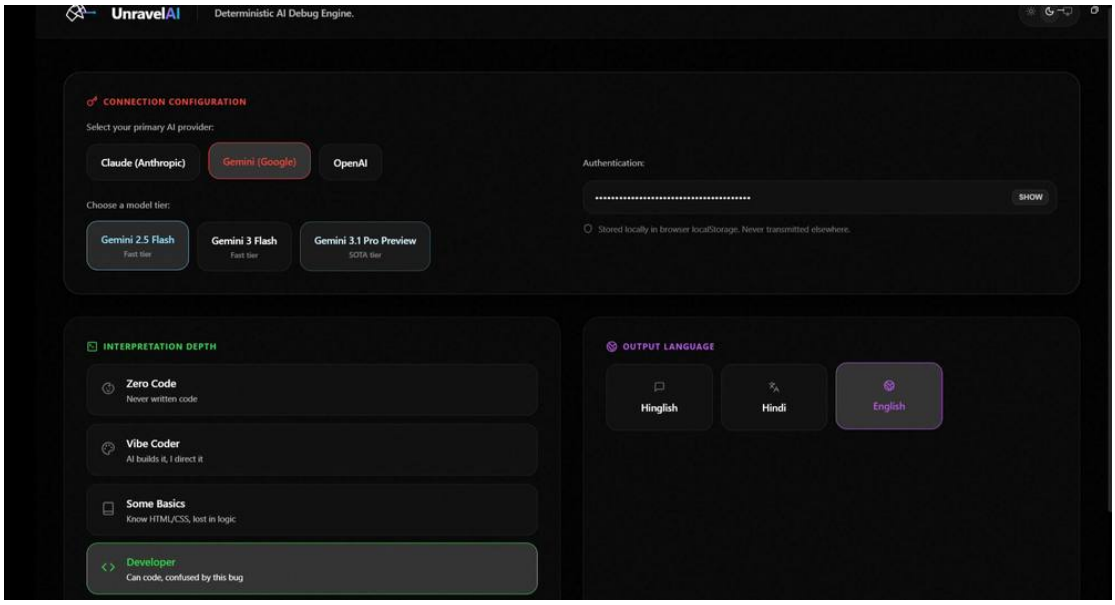

Figure 7: Unravel diagnosis output for the claude-hud terminal corruption bug (Case Study 3, Issue #256). The *Confidence Evidence* block cites exact file and line references verified by the AST engine. This structured output was shared in the issue thread and informed PR #259.

6.6.4 Notable: Unravel Diagnosed Its Own Codebase

During development, a streaming display bug was encountered in Unravel's own UI: the report panel was rendering incomplete output during live streaming, and the `parseAIJson` function was emitting spurious console warnings on partial JSON chunks. Claude Opus 4.6 – Anthropic's most capable model, used as the development assistant – was unable to resolve the issue after several minutes of iterative attempts. The Unravel codebase was then fed to Unravel itself.

Unravel's diagnosis: (1) `parseAIJson` in `parse-json.js` was not accepting an `isStreaming` parameter to suppress warnings on partial chunks; (2) `orchestrate.js` was not retaining a `lastGoodPartial` accumulator across successful parses, causing the UI to receive `null` on failed intermediate parses instead of the last known-good partial state.

Upon receiving the structured diagnosis, Claude Opus 4.6 implemented both fixes immediately – editing `orchestrate.js` and `parse-json.js` – and noted: *"you fed Unravel its own codebase, and it correctly diagnosed its own bug."* Figure 8 shows the Claude Code session where this occurred.

This is not a benchmark result and should not be read as one. It is included here as a demonstration of two system properties that are difficult to fabricate, not as empirical evidence for the paper's central quantitative claims. Treat it as a characterization of system behavior, not as comparative performance data.

The demonstration shows: (1) the system's diagnostic output is precise enough to be directly actionable by a SOTA model that had previously been unable to resolve the same issue; (2) the architecture is self-applicable – Unravel treats its own code as it treats any other TypeScript codebase.

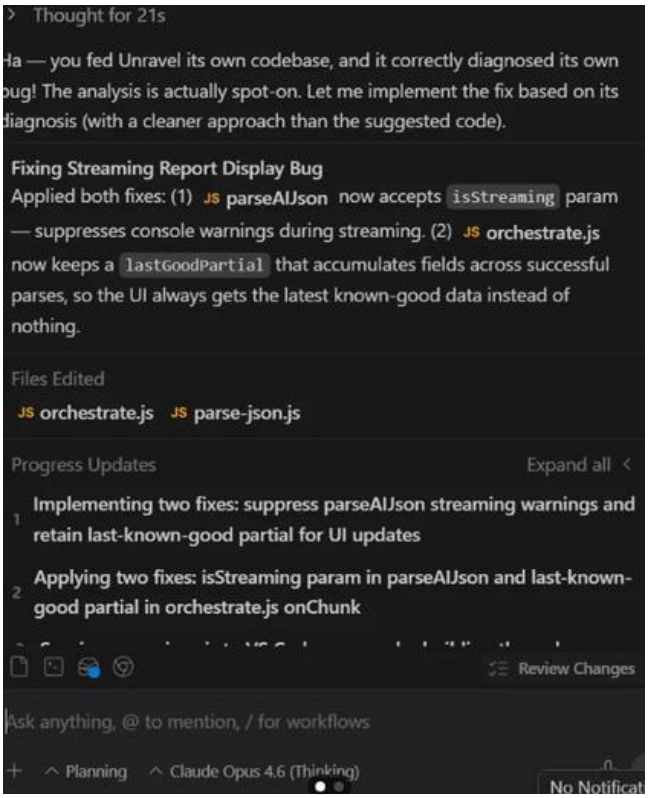

Figure 8: Claude Opus 4.6 (SOTA) implementing Unravel's self-diagnosis. The model had been unable to resolve the streaming display bug unaided; upon receiving Unravel's structured diagnosis, it implemented both fixes and noted the recursive nature of the interaction. Files edited: `orchestrate.js` and `parse-json.js`.

6.6.5 Aggregate Metrics

#	Repository	Bug Type	RCA	HR	Outcome
1	tldraw/tldraw #8148	`DATA_FLOW` + `LAYER_BOUNDARY`	✓	0%	Maintainers validated, issue closed
2	calcom/cal.com #28283	`STATE_MUTATION`	✓	0%	Maintainers validated, credited by reporter
3	claude-hud #256	`TEMPORAL_LOGIC`	✓	0%	PR #259 open for review

Validation rate: 2/2 confirmed (100%) on closed cases; Case Study 3 pending merge. All three diagnoses share effectively zero hallucinated entity claims (HR $\approx$ 0%) on

verifiable fields.

6.7 Comparative Evaluation *(Illustrative Only)*

Scope limitation. This section presents two hand-selected, author-constructed test cases. They are *not* randomly sampled, not independently constructed, and not representative of a distribution. No statistical claims are drawn here. This section characterizes the *categories of output* Unravel produces relative to baseline models on tractable bugs – not a performance ranking. Readers looking for reproducible quantitative evidence should refer to UDB-20 (§6.1.1).

To contextualize Unravel's output structure, we conducted a qualitative comparison across two hand-constructed test cases, evaluated against Claude 4.6, ChatGPT 5.3, and Gemini 3.1 Pro. All models received identical code context at temperature = 0. These results are strictly illustrative; they characterize output categories, not model intelligence rankings, and are not intended to survive peer scrutiny as comparative evidence.

6.7.1 Test 1: The Heisenbug (Single File)

A race condition in a dashboard initializer where two async operations mutate shared rendering state independently. The Heisenbug property: adding `console.log` for debugging changes microtask scheduling just enough to eliminate the bug, making it unreproducible under observation.

Table A: Core Diagnostic Correctness

Criterion	Claude	GPT-5.3	Gemini	Unravel
Bug category (`RACE_CONDITION`)	✓	✓	✓	✓
Heisenbug identified	✓	×	✓	✓
Microtask scheduling explained	✓	×	✓	✓
Spread non-atomicity caught	Approx	✓	✓	✓
Correct fix (`Promise.all`)	✓	✓	✓	✓
Error handling in fix	×	×	✓	✓

Table B: Extended Output Capabilities

Output capability	Claude	GPT-5.3	Gemini	Unravel
Developer debugging path predicted	×	×	×	✓
Variable state lifecycle	×	×	×	✓
Invariants documented	×	×	×	✓(4)
Execution timeline	Basic	Basic	×	Full
Competing hypotheses	×	×	×	✓
Structured JSON output	×	×	×	✓

6.7.2 Test 2: The Phantom Accumulator (5 Files)

A 4-way emergent bug across 5 files: a memoization reference bug, a WeakRef GC footgun, a microtask async ordering issue, and a Heisenbug observation effect combined.

Table A: Core Diagnostic Correctness

Criterion	Claude	GPT-5.3	Gemini	Unravel
Root cause correct	✓	✓	✓	✓
Mutation sites + exact lines	✓	✓	✓	✓(L22, L28, L32)
`resetBoard` non-contributing	✓	×	×	✓
Minimal fix (not full rewrite)	✓	×	✓	✓
Timestamped failure trace	Basic	×	×	Full (0s–10.5s)
Causal loop traced (8 steps)	✓	Partial	×	✓

Table B: Extended Output Capabilities

Output capability	Claude	GPT-5.3	Gemini	Unravel
Variable lifecycle tracker	×	×	×	✓(2 vars)
Invariants documented	×	×	×	✓(8)
Competing hypotheses	×	×	×	✓(3)
Structured reproduction steps	×	×	×	✓
AI fix prompt	×	×	×	✓
Structured JSON output	×	×	×	✓

Key finding. *On core diagnostic correctness (Table A), all four tools identified the root cause on both test cases, confirming that frontier models are sufficient for tractable bugs when given adequate context. The divergence is in Table B: Unravel additionally produces output categories (invariants, variable lifecycle traces, hypothesis enumerations, timestamped causal chains, reproduction steps, structured JSON) not observed in any baseline output. These categories are not produced by prompting; they are structural outputs of the 8-phase pipeline.*

Limitation. These are two hand-selected test cases. They are not randomly sampled

and are not representative of a distribution. Statistical conclusions cannot be drawn.

6.8 Historical Context: UDB-11

UDB-11 was the original 11-bug internal benchmark used during Claim Verifier design. It is retained here as historical context, not as a primary benchmark result. Its design limitations – self-referential construction, single-axis scoring, and author selection bias – are superseded by UDB-20.

Metric	Baseline (no AST)	Unravel (full pipeline)
RCA	91% (10/11)	100% (11/11)
HR	1.5% (1/66)	≈0% (0/66)

The baseline HR of 1.5% reflects a fabricated line number on one diagnosis (cited line exceeded file length). The Claim Verifier catches this deterministically regardless of model instruction compliance.

UDB-11's primary contribution was motivating the Claim Verifier design and demonstrating that AST pre-analysis produces a measurable signal even on an 11-bug sample. The real-world merged PRs (§6.2) constitute the primary correctness evidence.

7. Discussion

The results across B-01–B-10 are best understood through a single observation: the contribution is not that the model is given more information, but that it is prevented from reasoning inconsistently with information already known to be true. Everything else in Unravel's architecture follows from this.

7.1 Structural Understanding as a General Primitive

A system that traces the complete structural behavior of a codebase (which variables are mutated at which lines, which functions propagate data across file boundaries, where async execution diverges from lexical order, what invariants the code implicitly relies on) has acquired knowledge beyond what debugging requires. Debugging is one application of structural understanding.

Code explanation. Unravel's Explain mode uses the same Layer 0 facts to produce entry point hierarchies, data flow maps, and dependency graphs. Structural understanding is the hardest part of developer onboarding; it is automated deterministically.

Security auditing. Security mode uses the call graph and mutation chains to trace data from user-controlled input to sensitive operations. Injection vulnerabilities and

authentication bypasses are structural properties of the code: they appear in the AST whether or not the model is explicitly told to look for them.

Feature implementation. If Unravel knows which variables hold which invariants and which callers will be affected by a change, it can reason about *where to add a feature* without introducing structural violations. Current agentic coding tools generate code without structural knowledge of the codebase they are generating into. A code generation agent that has not verified which invariants must be preserved will, predictably, violate some of them.

7.2 Structural Grounding vs. Context Window Scaling

A predictable objection: frontier models with 1M-token context windows can ingest entire repositories, eliminating the need for selective routing. This addresses context size but not the structural problem:

Hypothesis elimination. Long-context models still anchor to the symptom description. Injecting 500K tokens of source code does not change the fact that the model pattern-matches from the symptom to the nearest familiar-looking code [5]. The AST fact that `duration` is written in `pause()` at L69 *eliminates* the stale-closure hypothesis that the model would otherwise anchor on. No amount of context makes this elimination happen automatically: it requires the verified structural constraint.

Ground truth injection. The statement "duration is written at L69, verified by static analysis" is categorically different from asking the model to infer that mutation from reading 80 lines of code. The former is a deterministic constraint injected before token generation begins; the latter is a probabilistic inference over the training distribution.

Cost. Unravel achieves the reported results on the Gemini 2.5 Flash free tier, demonstrating that structural grounding extracts strong diagnostic performance without requiring frontier model scale.

7.3 Reasoning Integrity: A Proposed Future Metric

The Claim Verifier checks whether *answers* are grounded in real code. It does not verify whether the *reasoning path* is grounded. Four failure modes are now distinct, the fourth identified empirically from B-22 Run 2:

Type	Description	Caught by current system
1	Wrong answer, fabricated evidence	Yes: Claim Verifier
2	Correct answer, fabricated reasoning	No
3	Mostly-correct reasoning, one unjustified inferential leap	No
4	Correct annotation-scoped fix; bug persists from adjacent violation	No

Type 2 is the dangerous case: the fix works, the developer trusts the diagnosis, but the reasoning was confabulated. The next similar bug is misdiagnosed using the same flawed logic.

Type 4 is the subtlest case, identified in B-22 Run 2. The reasoning is correct. The fix is locally correct – it eliminates the structural violation the annotation named. The bug persists because the annotated violation was a consequence of a deeper protocol error, not its origin. A verified structural annotation improves a model's reasoning about the specific pattern it describes. It does not guarantee the annotated pattern is the complete causal story.

B-22 Run 1 provides a Type 2 demonstration. Claude Sonnet 4.6 correctly identified the fix for Bug 1 in B-22 Run 1 (remove `promotePreVotesToGrants`). Its stated mechanism – "pre-vote records are counted without a real vote response" – is sufficient to justify the fix but shallower than the full structural truth. The full mechanism requires the ECMA-262 double-visit chain. Claude's fix happens to be structurally complete; its reasoning depth was not. Two outputs that agree on RCA but differ on structural depth are indistinguishable by correctness scoring alone.

B-22 Run 2 provides a Type 4 demonstration. Unravel's Run 2 Mode A fix (`term: electionTerm - 1` → `term: electionTerm`) correctly eliminates the ECMA-262 double-visit by preventing the stale-term condition that triggers it. The annotation's claim is fully satisfied. The split-brain persists because `promotePreVotesToGrants` still places pre-voters in `_grantedVotes` – a protocol violation that exists independently of what term those records carry. The annotation described the multiplier; the model fixed the multiplier; the addend remained. Detailed analysis in §**??**.

We propose **Reasoning Integrity (RI)** as a metric: for each hypothesis elimination in Phase 5, extract the factual claim and verify it against the AST scope map and mutation chains. RI is the percentage of elimination steps in which the model's factual claims are AST-verified. This enables characterizations such as: *"Model X: RCA 100%, RI 78%"*, distinguishing models that get correct answers by correct reasoning from those that get correct answers by confabulation. The infrastructure is already resident in Unravel: the AST facts and the Phase 5 `hypothesisTree` output are both in memory simultaneously; the cross-check is not yet automated.

7.4 Limitations

Language coverage. AST pre-analysis is currently limited to JavaScript, TypeScript, JSX, and TSX. Python, Go, Rust, and other languages receive the full 8-phase LLM pipeline but without structural ground-truth injection. tree-sitter grammars exist for all common languages; extension is structurally straightforward.

Static analysis scope. Without runtime execution, structural analysis cannot verify which code paths are actually taken under specific inputs. A mutation present in the AST may be on a dead code path. A race condition that manifests only under specific scheduling may not be statically identifiable.

Spec-correctness bugs. A class of defect that static analysis cannot fully verify: bugs where the code is structurally valid but semantically wrong relative to an external protocol specification. B-22 Bug 3 (`>` vs `>=` against Raft §5.4.1) is the clearest example. The response is `detectSpecViolationRisks()` – a heuristic attention signal (not a verified fact) that flags comparison operators in grant-gate functions. This is a weaker signal and requires monitoring for false-positive rate; the keyword set must remain tight enough that when the signal fires, a developer actually investigates.

Symptom coverage. Addressed by `buildSymptomCoverageAlert()` in Layer 3 and the `uncoveredSymptoms[]` schema field. The trigger condition is deliberately conservative: it fires on observably different failure modes, not on any symptom written as a list. Correctness of the coverage check on noisy real-world bug reports (vs. the structured symptom files used in UDB-20) requires validation on a broader corpus.

Single-root-cause schema. Addressed by `additionalRootCauses[]` and `uncoveredSymptoms[]` in the output schema. The LLM instruction distinguishes causally independent secondary bugs (different trigger, different fix location, different observable symptom) from secondary effects of the same root cause, which belong in `evidence[]` not `additionalRootCauses[]`. Correctness of this distinction under the full range of multi-bug reports remains to be validated.

Floating promise precision. The `isAwaited` guard correctly handles most cases but can produce false positives in complex chained callback expressions where upward parent traversal encounters ambiguous scope boundaries.

Benchmark scale and methodology. All 20 UDB-20 bugs plus three Super-Bug case studies have been graded. Full benchmark results (post-fix): 137/138 (99.3%) Unravel vs. 130/138 (94.2%) baseline, delta +7. B-22 Run 1 (5/6) is reported in full in §6.5; only Run 2 (6/6) appears in the aggregate. All 20 bugs are constructed by the same author who designed the system; grading is human against pre-locked ground truth, not independently replicated. The planted proximate traps were designed by the author and not independently validated as genuinely misleading before embedding. B-10 was re-run after an AST improvement; only the post-fix score appears in the

aggregate, with the pre-fix score reported separately (§6.2.2). A fully independent replication would substantially strengthen all quantitative claims.

We proactively invite independent replication. All 20 UDB-20 bugs, the ground-truth rubric, and the `pnpm test` reproducibility harness are included in the repository at https://github.com/EruditeCoder108/unravelai. Researchers wishing to re-run the benchmark, re-grade outputs, or extend UDB-20 with independently constructed bugs are encouraged to open an issue or pull request. This is the most direct path to strengthening the quantitative claims in this paper.

Failure case: scheduling-dependent race condition. In one analysis involving a race condition where two concurrent async writes converged on the same state object, AST analysis correctly identified both mutation paths and flagged the pattern as `RACE_CONDITION`. However, static analysis could not determine which of the two execution orderings was realized at runtime: both were structurally possible. Unravel returned an ambiguous diagnosis with two surviving hypotheses and confidence 0.61, noting that disambiguation required runtime instrumentation. This is an expected limitation of static-only analysis.

7.5 Benchmark Design Evolution and the Evaluation Validity Problem

The evaluation strategy for Unravel has evolved through three distinct stages, each exposing a limitation in the previous approach. Documenting this evolution is itself useful: the failure modes encountered generalize to any system whose claimed advantage is specifically about complex, multi-file, real-world debugging.

Stage 1 — Large but transparent synthetic bugs. The earliest internal testing used a suite of deliberately constructed bugs, each designed to be architecturally complex: cross-file state corruption, async boundary violations, layered causation. The intent was to stress-test the system. The outcome was unexpected: both Unravel and unaugmented frontier models solved nearly all of them reliably, and the performance delta was small. Post-analysis identified the reason. The bugs, despite their architectural sophistication, were contained in small, focused codebases specifically written to isolate the defect. With minimal surrounding code, the root cause was effectively visible in the files provided – there was nowhere else to look. Unravel's structural grounding still produced a measurable advantage, but the signal was weak because the alternative hypothesis space was small. A system built specifically for large codebases does not reveal its edge when tested on small ones.

Stage 2 — Real GitHub issues. The limitations of Stage 1 motivated applying Unravel to real open-source GitHub issues: tldraw issue #8148 (8,000+ file TypeScript monorepo) and cal.com issue #28283 (one of the largest open-source Next.js codebases). On these cases Unravel performed qualitatively differently. The file routing

problem, the cross-file causal chain reconstruction, and the hypothesis elimination under high-noise conditions all became non-trivial. The structured diagnosis it produced was used to construct patches accepted by maintainers. This is the external validation signal the synthetic bugs could not provide.

However, GitHub issues introduce a different evaluation problem. Obtaining a PR acceptance for each additional case requires submitting patches to real repositories and waiting for maintainer review – a process that may take days to weeks per case, is subject to factors entirely outside diagnostic quality (maintainer availability, competing PRs, project roadmap), and cannot be run in controlled comparative conditions. Two cases provide proof of concept; they do not provide a repeatable, ablation-ready benchmark.

Stage 3 — UDB-20: constructed bugs with realistic structural properties. UDB-20 represents an attempt to occupy the space between Stages 1 and 2: bugs constructed by the author but designed to exhibit the structural properties that make real GitHub issues hard. Each bug has a deliberately planted proximate trap, a multi-file causal chain, and a ground truth locked before any model runs. The test suite verifies both the bug and the fix independently via `pnpm test`. The three-axis rubric (RCA, PFR, CFR) is designed to directly measure the failure modes Unravel addresses rather than binary pass/fail.

The results across B-01–B-10 (Unravel 59/60 vs. baseline 52/60, delta +7) confirm a consistent advantage on CFR – the causal chain reconstruction axis (20/20 vs. 14/20) – while showing near-parity on RCA (19/20 vs. 17/20) and PFR (20/20 vs. 19/20). The author's assessment is that this delta, while real, understates Unravel's advantage on genuinely large-scale real-world bugs. UDB-20's bugs, despite deliberate effort to make them structurally realistic, remain smaller and more self-contained than a typical multi-file issue in a production codebase with tens of thousands of lines of surrounding context. The hypothesis space is still narrower than in the wild. The structural grounding layer contributes most when the number of plausible files and causal paths is large; at smaller scales, both augmented and unaugmented models can arrive at the correct answer by different routes.

The correct interpretation of current results. The three tiers of evidence in this paper should be read as measuring different things. The GitHub PR case studies (§6.2) measure whether Unravel produces correct, actionable diagnoses on real production bugs at the scale and complexity of codebases with thousands of files – they do not support comparative claims. UDB-20 (§6.3) measures a repeatable performance delta against a controlled baseline – it does not yet reflect the full difficulty distribution of real-world bugs. The expected pattern, consistent with both theoretical motivation and preliminary results, is that Unravel's advantage grows with bug complexity, codebase size, and causal chain length. Validating this prediction rigorously requires either

a larger constructed benchmark with provably realistic complexity, or a systematic real-world validation program that pairs identical bug inputs to both Unravel and a baseline and evaluates the resulting diagnoses against the ground truth established by maintainer-accepted fixes. Both remain open work.

8. Future Work

8.1 Heavy Mode: Four-Agent Parallel Analysis. For repositories too large for a single context window, a four-agent parallel system distributes the codebase: Agents 1, 2, and 3 each receive a distinct non-overlapping partition and run the full 8-phase pipeline independently. Agent 4 (Synthesis) receives all three JSON outputs, resolves agreement and contradiction, and issues a final merged diagnosis weighted by per-agent confidence. This is a parallelism strategy, not a debate: agents do not see each other's work during analysis.

8.2 Three-Agent Sequential Divergent System. Activated when single-agent confidence falls below a tuned threshold. Agent 2 receives Agent 1's summary as a negative prompt: it must look in a completely different direction and cannot reproduce Agent 1's hypothesis. Agent 3 receives both summaries as negative prompts. If all three agents converge independently, confidence in the shared hypothesis is elevated. If they diverge, all three diagnoses are surfaced: the developer receives the full uncertainty picture rather than a false consensus.

8.3 Reasoning Integrity Verifier. Automated implementation of the RI metric described in §7.3: for each hypothesis elimination in Phase 5, extract the factual claim, cross-check it against the AST scope map and mutation chains, and flag reasoning steps that are factually unverifiable.

8.4 WebContainers Runtime Execution. Run reproduction steps in an in-browser sandboxed Node.js environment. Capture concrete variable values at the crash site and inject them as additional verified facts alongside AST ground truth. This converts "AST indicates this mutation could cause the bug" into "runtime confirms value was 1490, not 1500."

8.5 Multi-Language AST Support. tree-sitter grammars exist for Python, Go, Rust, Java, C, and C++. The scope resolver, mutation extractor, and closure detector are language-agnostic in design; extension requires grammar-specific node type mapping.

8.6 Semantic Grounding Layer. Embedding-based similarity for symptom-to-known-bug matching (inject top-3 confirmed prior bugs as grounding context before the LLM call), confidence calibration via response-consistency clustering across multiple temperature-0.3 runs, and hypothesis deduplication by collapsing semantically

identical candidates before Phase 4.

8.7 Long-Term Vision. The structural intelligence layer that makes debugging reliable (deterministic pre-analysis, call graph routing, hypothesis elimination, verified output) is the same layer that makes all AI code interaction reliable. The natural evolution is toward an IDE-level environment in which structural analysis precedes every AI interaction with the codebase: feature additions, refactoring, security review, onboarding. Structural pre-analysis is that layer.

8.8 Multimodal Bug Description. Developers describe bugs with text today because that is the only supported input. The actual symptom is usually visual: a broken UI state, an unexpected network response in the browser DevTools, a console output that should not exist. Planned work includes accepting screenshot and screen recording input alongside source files, allowing the model to ground its diagnosis in the *observable behavior* rather than a developer's possibly-imprecise text description of it. A developer who can paste a screenshot of a broken UI state alongside the component files provides the model with a richer constraint: not just "the avatar does not show" but the exact pixel state, the network tab at the moment of failure, and the console output – all as verified additional context injected before the LLM call, following the same architectural principle as AST facts.

8.9 Continuous AST Engine Expansion. The pattern observed throughout the UDB-20 evaluation – a targeted 47-word annotation or 150-line detector converting a partial diagnosis to a perfect one – motivates a systematic detector development program. Every bug class that Unravel encounters in the wild that requires domain-specific structural knowledge (listener identity rules, async yield boundaries, module-scope state isolation, React rendering contracts) is a candidate for a new deterministic detector. The AST engine is not a fixed artifact; it is a growing library of structural facts that progressively narrows the space of bugs the model must infer rather than read. Development of new detectors continues in parallel with the benchmark program, and every new detector added will be back-tested against UDB-20 to measure its contribution to scores on existing bugs.

9. Conclusion

We presented Unravel, a debugging engine built on the thesis that deterministic structural pre-analysis is a necessary layer, not an optimization, without which LLM reasoning over code is architecturally unreliable. The four-layer architecture implements the generate-verify-eliminate loop that hypothesis-driven debugging research identifies as the primary mitigation for anchoring bias and proximate fixation.

The meaningful empirical record is three real-world case studies: tldraw issue #8148, cal.com issue #28283, and claude-hud issue #256 – all in production repositories, di-

agnosed correctly using Gemini 2.5 Flash on the free API tier – a model priced 20× less than the Claude Sonnet 4.6 frontier baseline used throughout the benchmark. In each case, Unravel's structured diagnosis was shared in the issue thread and independently validated by maintainers. UDB-20 (§6.1.1) provides a reproducible three-axis evaluation framework – RCA, Proximate Fixation Resistance (PFR), and Causal Flow Reconstruction (CFR) – designed to directly measure the failure modes Unravel targets. Results across the full benchmark (137/138 Unravel vs. 130/138 baseline, delta +7) confirm the architecture at multiple levels. At the benchmark level: the full CFR advantage (40/40 vs. 34/40) concentrates in the unstructured era; with the structured prompt (B-12–B-20), both engines tie. At the case-study level: the Ghost Tenant Super-Bug requires careful interpretation. Claude Sonnet 4.6 scored 6/6 *without* AST context – the bug was within its reliable inference window, and this is the paper's correct prediction: on bugs where careful reading by a sufficiently capable model is adequate, both approaches produce the same result. The evidence for the thesis is what happens at the boundary: Gemini 2.5 Flash without structural grounding scored 4/6, and with a 150-line targeted detector – delivering the write-before-await compound fact as verified ground truth – it matched Claude Sonnet 4.6 exactly. The AST layer closes the gap for a weaker model operating at its limit; it does not claim to improve a frontier model on bugs already within its reach.

Larger models will continue to improve at pattern matching over code. The structural reasoning gap (knowing which variables are mutated at which lines, which callers break if a function changes, what invariants a fix must preserve) does not close through scale alone. It requires verified structural facts about the specific code being analyzed. AST pre-analysis surfaces exactly those facts. The model does not become more capable. It finally has something verified to reason about.

The goal is not to replace LLM reasoning – it is to constrain it with verified truth and guide it with calibrated suspicion. Verified structural facts narrow the space of wrong answers to near zero on the dimensions they cover. Guided suspicion – symptom coverage, heuristic attention signals, hypothesis elimination requirements – shapes where the model looks and what it must account for. The model, the codebase, and the bug class all change. The structural pre-analysis layer is the invariant. That is the architecture this paper describes, and the one the benchmark results support.

References

[1] Jimenez, C.E., et al. (2024). "SWE-BENCH: Can Language Models Resolve Real-World GitHub Issues?" *ICLR 2024*. https://proceedings.iclr.cc/paper_files/paper/2024/file/edac78c3e300629acfe6cbe9ca88fb84-Paper-Conference.pdf

[2] Liu, et al. (2023). "Large Language Models of Code Fail at Completing Code with Po-

tential Bugs." *NeurIPS 2023*. https://proceedings.neurips.cc/paper_files/paper/2023/file/819cebb05f993840e8a52d7564c5c282-Paper-Conference.pdf

[3] Zhou, X., et al. (2026). "Cognitive Biases in LLM-Assisted Software Development." *ICSE 2026*. arXiv:2601.08045. https://arxiv.org/abs/2601.08045

[4] "The Debugging Decay Index." (2026). *arXiv:2506.18403*. https://arxiv.org/abs/2506.18403

[5] Valencia-Clavijo, F. (2025). "Anchors in the Machine: Behavioral and Attributional Evidence of Anchoring Bias in LLMs." *arXiv:2511.05766*. https://arxiv.org/abs/2511.05766

[6] Cai, Z., et al. (2023). "Unleashing the Power of Type-Based Call Graph Construction." *USENIX Security 2023*. https://www.usenix.org/system/files/sec23winter-prepub-350-cai.pdf

[7] "Multi-SWE-bench: A Multilingual Benchmark for Software Engineering." *NeurIPS 2025*. https://neurips.cc/virtual/2025/poster/121692

[8] "Agentless: Demystifying LLM-Based Software Engineering Agents." *arXiv:2407.01489 / findings cited in arXiv:2509.13941*. https://arxiv.org/abs/2509.13941

[9] Ferrante, J., Ottenstein, K.J., Warren, J.D. (1987). "The Program Dependence Graph and Its Use in Optimization." *ACM TOPLAS*. https://web.eecs.umich.edu/~mahlke/courses/583f23/reading/ferrante_toplas_87.pdf

[10] Horwitz, S., Reps, T., Binkley, D. (1990). "Interprocedural Slicing Using Dependence Graphs." *ACM TOPLAS*.

[11] Aho, A.V., Lam, M.S., Sethi, R., Ullman, J.D. (2006). *Compilers: Principles, Techniques, and Tools*. Addison-Wesley.

[12] "Unveiling the Power of Intermediate Representations for Static Analysis." *arXiv:2405.12841*. https://arxiv.org/abs/2405.12841

[13] "When Threads Meet Events: Precise Static Race Detection with O2." O2Lab Technical Report. https://o2lab.github.io/p/o2.pdf

[14] "LDB: A Large Language Model Debugger via Verifying Runtime Execution Step by Step." *arXiv:2402.16906*. https://arxiv.org/abs/2402.16906

[15] "TraceCoder: A Trace-Driven Multi-Agent Framework for Automated Debugging." *arXiv:2602.06875*. https://arxiv.org/abs/2602.06875

[16] Chen, J., et al. (2024). "H-Neurons: On the Existence, Impact, and Origin of Hallucination-Associated Neurons in LLMs." *arXiv:2512.01797*. https://arxiv.org/abs/2512.01797

[17] tldraw/tldraw issue #8148. Unravel diagnosis shared in thread; fix independently implemented by maintainers in PR #8161. https://github.com/tldraw/tldraw/pull/8161

[18] calcom/cal.com issue #28283. Unravel diagnosis shared in thread; fix independently implemented by maintainers in PR #28296. https://github.com/calcom/cal.com/

pull/28296

[19] jarrodwatts/claude-hud issue #256. Unravel diagnosis shared in thread; PR #259 open for maintainer review at time of writing. https://github.com/jarrodwatts/claude-hud/pull/259

[20] Unravel source code (BSL-1.1 license). https://github.com/EruditeCoder108/unravelai

[21] Unravel web application (live). https://vibeunravel.netlify.app

---

Submitted to arXiv cs.SE. March 2026. *Source: https://github.com/EruditeCoder108/unravelai*